The Linux Kernel API

The Linux Kernel API

This documentation is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.
This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.
You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place, Suite 330, Boston, MA 02111-1307 USA
For more details see the file COPYING in the source distribution of Linux.

Table of Contents

1. Driver Basics 1
1.1. Driver Entry and Exit points 1
module_init 1
module_exit 1
1.2. Atomic and pointer manipulation 2
atomic_read 3
atomic_set 3
atomic_add 4
atomic_sub 5
atomic_sub_and_test 6
atomic_inc 7
atomic_dec 8
atomic_dec_and_test 9
atomic_inc_and_test 10
atomic_add_negative 11
get_unaligned 12
put_unaligned 12
1.3. Delaying, scheduling, and timer routines 13
schedule_timeout 14
2. Data Types 16
2.1. Doubly Linked Lists 16
list_add 16
list_add_tail 16
list_del 17
list_del_init 18
list_empty 19
list_splice 19
list_entry 20
list_for_each 21
list_for_each_safe 22
3. Basic C Library Functions 24
3.1. String Conversions 24
simple_strtol 24
simple_strtoll 25
simple_strtoul 25
simple_strtoull 26
vsnprintf 27

snprintf 28
vsprintf 29
sprintf 30
3.2. String Manipulation 31
strcpy 31
strncpy 31
strcat 32
strncat 33
strcmp 34
strncmp 35
strchr 36
strrchr 36
strlen 37
strnlen 38
strpbrk 38
strtok 39
memset 40
bcopy 41
memcpy 42
memmove 43
memcmp 43
memscan 44
strstr 45
memchr 46
3.3. Bit Operations 47
set_bit 47
__set_bit 48
clear_bit 49
__change_bit 49
change_bit 50
test_and_set_bit 51
__test_and_set_bit 52
test_and_clear_bit 53
__test_and_clear_bit 54
test_and_change_bit 55
test_bit 56
find_first_zero_bit 56
find_next_zero_bit 57
ffz 58
ffs 59

hweight32 60
4. Memory Management in Linux 61
4.1. The Slab Cache 61
kmem_cache_create 61
kmem_cache_shrink 62
kmem_cache_destroy 63
kmem_cache_alloc 64
kmalloc 65
kmem_cache_free 66
kfree 67
5. The proc filesystem 69
5.1. sysctl interface 69
register_sysctl_table 69
unregister_sysctl_table 71
proc_dostring 71
proc_dointvec 73
proc_dointvec_minmax 74
proc_doulongvec_minmax 75
proc_doulongvec_ms_jiffies_minmax 76
proc_dointvec_jiffies 77
6. The Linux VFS 80
6.1. The Directory Cache 80
d_invalidate 80
d_find_alias 80
prune_dcache 81
shrink_dcache_sb 82
have_submounts 83
shrink_dcache_parent 84
d_alloc 84
d_instantiate 85
d_alloc_root 86
d_lookup 87
d_validate 88
d_delete 89
d_rehash 90
d_move 91
__d_path 91
is_subdir 93
find_inode_number 94

d_drop 94
d_add 95
dget 96
d_unhashed 97
6.2. Inode Handling 98
__mark_inode_dirty 98
write_inode_now 99
clear_inode 100
invalidate_inodes 100
get_empty_inode 101
iunique 102
insert_inode_hash 103
remove_inode_hash 104
iput 105
bmap 105
update_atime 106
make_bad_inode 107
is_bad_inode 108
6.3. Registration and Superblocks 109
register_filesystem 109
unregister_filesystem 110
get_super 111
6.4. File Locks 111
posix_lock_file 112
__get_lease 113
lease_get_mtime 114
posix_block_lock 114
posix_unblock_lock 115
lock_may_read 116
lock_may_write 117
fcntl_getlease 118
fcntl_setlease 119
sys_flock 120
get_locks_status 121
7. Linux Networking 123
7.1. Socket Buffer Functions 123
skb_queue_empty 123
skb_get 123
kfree_skb 124
skb_cloned 125

skb_shared....126
skb_share_check....127
skb_unshare....128
skb_peek....129
skb_peek_tail....129
skb_queue_len....130
__skb_queue_head....131
skb_queue_head....132
__skb_queue_tail....133
skb_queue_tail....134
__skb_dequeue....135
skb_dequeue....136
skb_insert....137
skb_append....138
skb_unlink....138
__skb_dequeue_tail....139
skb_dequeue_tail....140
skb_put....141
skb_push....142
skb_pull....143
skb_headroom....144
skb_tailroom....144
skb_reserve....145
skb_trim....146
skb_orphan....147
skb_queue_purge....148
__skb_queue_purge....148
__dev_alloc_skb....149
dev_alloc_skb....150
skb_cow....151
skb_over_panic....152
skb_under_panic....153
alloc_skb....154
__kfree_skb....155
skb_clone....156
skb_copy....157
pskb_copy....158
pskb_expand_head....159
skb_copy_expand....160
__pskb_pull_tail....161

7.2. Socket Filter 162
sk_run_filter 162
sk_chk_filter 163
8. Network device support 165
8.1. Driver Support 165
init_etherdev 165
alloc_etherdev 166
init_fddidev 167
alloc_fddidev 168
init_hippi_dev 168
alloc_hippi_dev 169
init_trdev 170
alloc_trdev 171
init_fcdev 172
alloc_fcdev 173
dev_add_pack 174
dev_remove_pack 175
__dev_get_by_name 175
dev_get_by_name 176
dev_get 177
__dev_get_by_index 178
dev_get_by_index 179
dev_alloc_name 179
dev_alloc 180
netdev_state_change 181
dev_load 182
dev_open 183
dev_close 184
register_netdevice_notifier 185
unregister_netdevice_notifier 185
dev_queue_xmit 186
netif_rx 187
net_call_rx_atomic 188
register_gifconf 189
netdev_set_master 190
dev_set_promiscuity 191
dev_set_allmulti 192
dev_ioctl 192
dev_new_index 193
netdev_finish_unregister 194

unregister_netdevice 195
8.2. 8390 Based Network Cards 196
ei_open 196
ei_close 197
ei_tx_timeout 198
ei_interrupt 198
ethdev_init 199
NS8390_init 200
8.3. Synchronous PPP 201
sppp_input 201
sppp_close 202
sppp_open 203
sppp_reopen 204
sppp_change_mtu 205
sppp_do_ioctl 206
sppp_attach 207
sppp_detach 207
9. Module Support 209
9.1. Module Loading 209
request_module 209
call_usermodehelper 209
9.2. Inter Module support 210
inter_module_register 211
inter_module_unregister 211
inter_module_get 212
inter_module_get_request 213
inter_module_put 214
10. Hardware Interfaces 216
10.1. Interrupt Handling 216
disable_irq_nosync 216
disable_irq 216
enable_irq 217
probe_irq_mask 218
10.2. MTRR Handling 219
mtrr_add 219
mtrr_del 221
10.3. PCI Support Library 222
pci_find_slot 222
pci_find_subsys 223

pci_find_device 224
pci_find_class 225
pci_find_capability 226
pci_find_parent_resource 227
pci_set_power_state 228
pci_save_state 229
pci_restore_state 230
pci_enable_device 231
pci_disable_device 231
pci_enable_wake 232
pci_release_regions 233
pci_request_regions 234
pci_match_device 235
pci_register_driver 236
pci_unregister_driver 237
pci_insert_device 237
pci_remove_device 238
pci_dev_driver 239
pci_set_master 240
pci_setup_device 241
pci_pool_create 241
pci_pool_destroy 243
pci_pool_alloc 244
pci_pool_free 245
10.4. MCA Architecture 246
10.4.1. MCA Device Functions 246
mca_find_adapter 246
mca_find_unused_adapter 247
mca_read_stored_pos 248
mca_read_pos 249
mca_write_pos 249
mca_set_adapter_name 251
mca_set_adapter_procfn 251
mca_is_adapter_used 253
mca_mark_as_used 253
mca_mark_as_unused 254
mca_get_adapter_name 255
mca_isadapter 256
mca_isenabled 256
10.4.2. MCA Bus DMA 257

mca_enable_dma....258
mca_disable_dma....258
mca_set_dma_addr....259
mca_get_dma_addr....260
mca_set_dma_count....261
mca_get_dma_residue....262
mca_set_dma_io....262
mca_set_dma_mode....263
11. The Device File System....265
devfs_register....265
devfs_unregister....266
devfs_mk_symlink....267
devfs_mk_dir....268
devfs_find_handle....269
devfs_get_flags....270
devfs_get_maj_min....271
devfs_get_handle_from_inode....272
devfs_generate_path....273
devfs_get_ops....274
devfs_set_file_size....275
devfs_get_info....275
devfs_set_info....276
devfs_get_parent....277
devfs_get_first_child....278
devfs_get_next_sibling....279
devfs_auto_unregister....279
devfs_get_unregister_slave....280
devfs_get_name....281
devfs_register_chrdev....282
devfs_register_blkdev....283
devfs_unregister_chrdev....284
devfs_unregister_blkdev....285
12. Power Management....286
pm_register....286
pm_unregister....287
pm_unregister_all....287
pm_send....288
pm_send_all....290
pm_find....291

13. Block Devices 293
blk_cleanup_queue 293
blk_queue_headactive 294
blk_queue_make_request 295
blk_init_queue 296
generic_make_request 297
submit_bh 298
ll_rw_block 299
end_that_request_first 301
14. Miscellaneous Devices 303
misc_register 303
misc_deregister 303
15. Video4Linux 305
video_unregister_device 305
16. Sound Devices 306
register_sound_special 306
register_sound_mixer 306
register_sound_midi 307
register_sound_dsp 308
register_sound_synth 309
unregister_sound_special 310
unregister_sound_mixer 311
unregister_sound_midi 312
unregister_sound_dsp 313
unregister_sound_synth 314
17. USB Devices 315
usb_register 315
usb_scan_devices 315
usb_deregister 316
usb_alloc_bus 317
usb_free_bus 318
usb_register_bus 319
usb_deregister_bus 319
usb_match_id 320
usb_alloc_urb 322
usb_free_urb 323
usb_control_msg 324
usb_bulk_msg 325

18. 16x50 UART Driver 328
register_serial 328
unregister_serial 328
19. Z85230 Support Library 330
z8530_interrupt 330
z8530_sync_open 331
z8530_sync_close 331
z8530_sync_dma_open 332
z8530_sync_dma_close 333
z8530_sync_txdma_open 334
z8530_sync_txdma_close 335
z8530_describe 336
z8530_init 337
z8530_shutdown 338
z8530_channel_load 339
z8530_null_rx 339
z8530_queue_xmit 340
z8530_get_stats 341
20. Frame Buffer Library 343
20.1. Frame Buffer Memory 343
register_framebuffer 343
unregister_framebuffer 344
20.2. Frame Buffer Console 345
fbcon_redraw_clear 345
fbcon_redraw_bmove 346
20.3. Frame Buffer Colormap 347
fb_alloc_cmap 348
fb_copy_cmap 348
fb_get_cmap 350
fb_set_cmap 351
fb_default_cmap 352
fb_invert_cmaps 353
20.4. Frame Buffer Generic Functions 353
fbgen_get_fix 354
fbgen_get_var 355
fbgen_set_var 356
fbgen_get_cmap 356
fbgen_set_cmap 357
fbgen_pan_display 358

fbgen_do_set_var 359
fbgen_set_disp 360
fbgen_install_cmap 361
fbgen_update_var 362
fbgen_switch 363
fbgen_blank 364
20.5. Frame Buffer Video Mode Database 365
fb_find_mode 365
__fb_try_mode 366
20.6. Frame Buffer Macintosh Video Mode Database 368
console_getmode 368
console_setmode 368
console_setcmap 369
console_powermode 371
mac_vmode_to_var 371
mac_var_to_vmode 372
mac_map_monitor_sense 373
mac_find_mode 374
20.7. Frame Buffer Fonts 375
fbcon_find_font 376
fbcon_get_default_font 376

Chapter 1. Driver Basics

1.1. Driver Entry and Exit points

module_init

Name

`module_init` — driver initialization entry point

Synopsis

```
module_init ( x);
```

Arguments

x

function to be run at kernel boot time or module insertion

Description

`module_init` will add the driver initialization routine in the "__initcall.int" code segment if the driver is checked as "y" or static, or else it will wrap the driver initialization routine with `init_module` which is used by insmod and modprobe when the driver is used as a module.

module_exit

Name

`module_exit` — driver exit entry point

Synopsis

```
module_exit ( x);
```

Arguments

x

function to be run when driver is removed

Description

`module_exit` will wrap the driver clean-up code with `cleanup_module` when used with rmmod when the driver is a module. If the driver is statically compiled into the kernel, `module_exit` has no effect.

1.2. Atomic and pointer manipulation

atomic_read

Name

`atomic_read` — read atomic variable

Synopsis

```
atomic_read ( v);
```

Arguments

v

pointer of type atomic_t

Description

Atomically reads the value of *v*. Note that the guaranteed useful range of an atomic_t is only 24 bits.

atomic_set

Name

`atomic_set` — set atomic variable

Synopsis

```
atomic_set ( v,  i);
```

Arguments

v

pointer of type atomic_t

i

required value

Description

Atomically sets the value of *v* to *i*. Note that the guaranteed useful range of an atomic_t is only 24 bits.

atomic_add

Name

`atomic_add` — add integer to atomic variable

Synopsis

```
void atomic_add (int i, atomic_t * v);
```

Arguments

`i`

integer value to add

`v`

pointer of type atomic_t

Description

Atomically adds *`i`* to *`v`*. Note that the guaranteed useful range of an atomic_t is only 24 bits.

atomic_sub

Name

`atomic_sub` — subtract the atomic variable

Synopsis

```
void atomic_sub (int i, atomic_t * v);
```

Arguments

`i`

integer value to subtract

`v`

pointer of type atomic_t

Description

Atomically subtracts *`i`* from *`v`*. Note that the guaranteed useful range of an atomic_t is only 24 bits.

atomic_sub_and_test

Name

`atomic_sub_and_test` — subtract value from variable and test result

Synopsis

```
int atomic_sub_and_test (int i, atomic_t * v);
```

Arguments

`i`

integer value to subtract

`v`

pointer of type atomic_t

Description

Atomically subtracts *`i`* from *`v`* and returns true if the result is zero, or false for all other cases. Note that the guaranteed useful range of an atomic_t is only 24 bits.

atomic_inc

Name

`atomic_inc` — increment atomic variable

Synopsis

```
void atomic_inc (atomic_t * v);
```

Arguments

v

pointer of type atomic_t

Description

Atomically increments *v* by 1. Note that the guaranteed useful range of an atomic_t is only 24 bits.

atomic_dec

Name

`atomic_dec` — decrement atomic variable

Synopsis

```
void atomic_dec (atomic_t * v);
```

Arguments

v

pointer of type atomic_t

Description

Atomically decrements *v* by 1. Note that the guaranteed useful range of an atomic_t is only 24 bits.

atomic_dec_and_test

Name

`atomic_dec_and_test` — decrement and test

Synopsis

```
int atomic_dec_and_test (atomic_t * v);
```

Arguments

v

pointer of type atomic_t

Description

Atomically decrements *v* by 1 and returns true if the result is 0, or false for all other cases. Note that the guaranteed useful range of an atomic_t is only 24 bits.

atomic_inc_and_test

Name

`atomic_inc_and_test` — increment and test

Synopsis

```
int atomic_inc_and_test (atomic_t * v);
```

Arguments

v

pointer of type atomic_t

Description

Atomically increments *v* by 1 and returns true if the result is zero, or false for all other cases. Note that the guaranteed useful range of an atomic_t is only 24 bits.

atomic_add_negative

Name

`atomic_add_negative` — add and test if negative

Synopsis

```
int atomic_add_negative (int i, atomic_t * v);
```

Arguments

i

integer value to add

v

pointer of type atomic_t

Description

Atomically adds *i* to *v* and returns true if the result is negative, or false when result is greater than or equal to zero. Note that the guaranteed useful range of an atomic_t is only 24 bits.

get_unaligned

Name

`get_unaligned` — get value from possibly mis-aligned location

Synopsis

```
get_unaligned ( ptr);
```

Arguments

`ptr`

pointer to value

Description

This macro should be used for accessing values larger in size than single bytes at locations that are expected to be improperly aligned, e.g. retrieving a u16 value from a location not u16-aligned.

Note that unaligned accesses can be very expensive on some architectures.

put_unaligned

Name

`put_unaligned` — put value to a possibly mis-aligned location

Synopsis

```
put_unaligned ( val,  ptr);
```

Arguments

`val`

value to place

`ptr`

pointer to location

Description

This macro should be used for placing values larger in size than single bytes at locations that are expected to be improperly aligned, e.g. writing a u16 value to a location not u16-aligned.

Note that unaligned accesses can be very expensive on some architectures.

1.3. Delaying, scheduling, and timer routines

schedule_timeout

Name

`schedule_timeout` — sleep until timeout

Synopsis

```
signed long schedule_timeout (signed long timeout);
```

Arguments

`timeout`

timeout value in jiffies

Description

Make the current task sleep until *`timeout`* jiffies have elapsed. The routine will return immediately unless the current task state has been set (see `set_current_state`).

You can set the task state as follows -

`TASK_UNINTERRUPTIBLE` - at least *`timeout`* jiffies are guaranteed to pass before the routine returns. The routine will return 0

`TASK_INTERRUPTIBLE` - the routine may return early if a signal is delivered to the current task. In this case the remaining time in jiffies will be returned, or 0 if the timer expired in time

The current task state is guaranteed to be TASK_RUNNING when this routine returns.

Specifying a *`timeout`* value of `MAX_SCHEDULE_TIMEOUT` will schedule the CPU away without a bound on the timeout. In this case the return value will be `MAX_SCHEDULE_TIMEOUT`.

In all cases the return value is guaranteed to be non-negative.

Chapter 2. Data Types

2.1. Doubly Linked Lists

list_add

Name

`list_add` — add a new entry

Synopsis

```
void list_add (struct list_head * new, struct list_head * head);
```

Arguments

new

new entry to be added

head

list head to add it after

Description

Insert a new entry after the specified head. This is good for implementing stacks.

list_add_tail

Name

`list_add_tail` — add a new entry

Synopsis

```
void list_add_tail (struct list_head * new, struct list_head *
head);
```

Arguments

new

new entry to be added

head

list head to add it before

Description

Insert a new entry before the specified head. This is useful for implementing queues.

list_del

Name

`list_del` — deletes entry from list.

Synopsis

```
void list_del (struct list_head * entry);
```

Arguments

entry

the element to delete from the list.

Note

list_empty on entry does not return true after this, the entry is in an undefined state.

list_del_init

Name

`list_del_init` — deletes entry from list and reinitialize it.

Synopsis

```
void list_del_init (struct list_head * entry);
```

Arguments

entry

the element to delete from the list.

list_empty

Name

`list_empty` — tests whether a list is empty

Synopsis

```
int list_empty (struct list_head * head);
```

Arguments

head

the list to test.

list_splice

Name

`list_splice` — join two lists

Synopsis

```
void list_splice (struct list_head * list, struct list_head * head);
```

Arguments

`list`

the new list to add.

`head`

the place to add it in the first list.

list_entry

Name

`list_entry` — get the struct for this entry

Synopsis

```
list_entry ( ptr,  type,  member);
```

Arguments

ptr

the &struct list_head pointer.

type

the type of the struct this is embedded in.

member

the name of the list_struct within the struct.

list_for_each

Name

`list_for_each` — iterate over a list

Synopsis

```
list_for_each ( pos,  head);
```

Arguments

pos

the &struct list_head to use as a loop counter.

head

the head for your list.

list_for_each_safe

Name

`list_for_each_safe` — iterate over a list safe against removal of list entry

Synopsis

```
list_for_each_safe ( pos,  n,  head);
```

Arguments

pos

the &struct list_head to use as a loop counter.

n

another &struct list_head to use as temporary storage

head

the head for your list.

Chapter 3. Basic C Library Functions

When writing drivers, you cannot in general use routines which are from the C Library. Some of the functions have been found generally useful and they are listed below. The behaviour of these functions may vary slightly from those defined by ANSI, and these deviations are noted in the text.

3.1. String Conversions

simple_strtol

Name

`simple_strtol` — convert a string to a signed long

Synopsis

```
long simple_strtol (const char * cp, char ** endp, unsigned int base);
```

Arguments

cp

The start of the string

endp

A pointer to the end of the parsed string will be placed here

base

The number base to use

simple_strtoll

Name

`simple_strtoll` — convert a string to a signed long long

Synopsis

```
long long simple_strtoll (const char * cp, char ** endp,
unsigned int base);
```

Arguments

cp

The start of the string

endp

A pointer to the end of the parsed string will be placed here

base

The number base to use

simple_strtoul

Name

`simple_strtoul` — convert a string to an unsigned long

Synopsis

```
unsigned long simple_strtoul (const char * cp, char ** endp,
unsigned int base);
```

Arguments

cp

The start of the string

endp

A pointer to the end of the parsed string will be placed here

base

The number base to use

simple_strtoull

Name

`simple_strtoull` — convert a string to an unsigned long long

Synopsis

```
unsigned long long simple_strtoull (const char * cp, char **
endp, unsigned int base);
```

Arguments

cp

The start of the string

endp

A pointer to the end of the parsed string will be placed here

base

The number base to use

vsnprintf

Name

`vsnprintf` — Format a string and place it in a buffer

Synopsis

```
int vsnprintf (char * buf, size_t size, const char * fmt,
va_list args);
```

Arguments

buf

The buffer to place the result into

size

The size of the buffer, including the trailing null space

fmt

The format string to use

args

Arguments for the format string

Description

Call this function if you are already dealing with a va_list. You probably want snprintf instead.

snprintf

Name

snprintf — Format a string and place it in a buffer

Synopsis

```
int snprintf (char * buf, size_t size, const char * fmt, ...
...);
```

Arguments

buf

The buffer to place the result into

size

The size of the buffer, including the trailing null space

fmt

The format string to use @...: Arguments for the format string

...

variable arguments

vsprintf

Name

vsprintf — Format a string and place it in a buffer

Synopsis

```
int vsprintf (char * buf, const char * fmt, va_list args);
```

Arguments

buf

The buffer to place the result into

fmt

The format string to use

args

Arguments for the format string

Description

Call this function if you are already dealing with a va_list. You probably want sprintf instead.

sprintf

Name

`sprintf` — Format a string and place it in a buffer

Synopsis

```
int sprintf (char * buf, const char * fmt, ... ...);
```

Arguments

buf

The buffer to place the result into

fmt

The format string to use @...: Arguments for the format string

`...`

variable arguments

3.2. String Manipulation

strcpy

Name

`strcpy` — Copy a `NUL` terminated string

Synopsis

```
char * strcpy (char * dest, const char * src);
```

Arguments

`dest`

Where to copy the string to

`src`

Where to copy the string from

strncpy

Name

`strncpy` — Copy a length-limited, `NUL-terminated` string

Synopsis

```
char * strncpy (char * dest, const char * src, size_t count);
```

Arguments

`dest`

Where to copy the string to

`src`

Where to copy the string from

`count`

The maximum number of bytes to copy

Description

Note that unlike userspace strncpy, this does not `NUL-pad` the buffer. However, the result is not `NUL-terminated` if the source exceeds *`count`* bytes.

strcat

Name

`strcat` — Append one `NUL-terminated` string to another

Synopsis

```
char * strcat (char * dest, const char * src);
```

Arguments

`dest`

The string to be appended to

`src`

The string to append to it

strncat

Name

`strncat` — Append a length-limited, `NUL-terminated` string to another

Synopsis

```
char * strncat (char * dest, const char * src, size_t count);
```

Arguments

dest

The string to be appended to

src

The string to append to it

count

The maximum numbers of bytes to copy

Description

Note that in contrast to strncpy, strncat ensures the result is terminated.

strcmp

Name

strcmp — Compare two strings

Synopsis

```
int strcmp (const char * cs, const char * ct);
```

Arguments

cs

One string

ct

Another string

strncmp

Name

strncmp — Compare two length-limited strings

Synopsis

```
int strncmp (const char * cs, const char * ct, size_t count);
```

Arguments

cs

One string

ct

Another string

count

The maximum number of bytes to compare

strchr

Name

`strchr` — Find the first occurrence of a character in a string

Synopsis

```
char * strchr (const char * s, int c);
```

Arguments

s

The string to be searched

c

The character to search for

strrchr

Name

`strrchr` — Find the last occurrence of a character in a string

Synopsis

```
char * strrchr (const char * s, int c);
```

Arguments

s

The string to be searched

c

The character to search for

strlen

Name

`strlen` — Find the length of a string

Synopsis

```
size_t strlen (const char * s);
```

Arguments

s

The string to be sized

strnlen

Name

`strnlen` — Find the length of a length-limited string

Synopsis

```
size_t strnlen (const char * s, size_t count);
```

Arguments

`s`

The string to be sized

`count`

The maximum number of bytes to search

strpbrk

Name

`strpbrk` — Find the first occurrence of a set of characters

Synopsis

```
char * strpbrk (const char * cs, const char * ct);
```

Arguments

cs

The string to be searched

ct

The characters to search for

strtok

Name

`strtok` — Split a string into tokens

Synopsis

```
char * strtok (char * s, const char * ct);
```

Arguments

s

The string to be searched

ct

The characters to search for

WARNING

strtok is deprecated, use strsep instead.

memset

Name

memset — Fill a region of memory with the given value

Synopsis

```
void * memset (void * s, int c, size_t count);
```

Arguments

s

Pointer to the start of the area.

c

The byte to fill the area with

count

The size of the area.

Description

Do not use `memset` to access IO space, use `memset_io` instead.

bcopy

Name

`bcopy` — Copy one area of memory to another

Synopsis

```
char * bcopy (const char * src, char * dest, int count);
```

Arguments

`src`

Where to copy from

`dest`

Where to copy to

`count`

The size of the area.

Description

Note that this is the same as `memcpy`, with the arguments reversed. `memcpy` is the standard, `bcopy` is a legacy BSD function.

You should not use this function to access IO space, use `memcpy_toio` or `memcpy_fromio` instead.

memcpy

Name

`memcpy` — Copy one area of memory to another

Synopsis

```
void * memcpy (void * dest, const void * src, size_t count);
```

Arguments

`dest`

Where to copy to

`src`

Where to copy from

`count`

The size of the area.

Description

You should not use this function to access IO space, use `memcpy_toio` or `memcpy_fromio` instead.

memmove

Name

`memmove` — Copy one area of memory to another

Synopsis

```
void * memmove (void * dest, const void * src, size_t count);
```

Arguments

`dest`

Where to copy to

`src`

Where to copy from

`count`

The size of the area.

Description

Unlike `memcpy`, `memmove` copes with overlapping areas.

memcmp

Name

memcmp — Compare two areas of memory

Synopsis

```
int memcmp (const void * cs, const void * ct, size_t count);
```

Arguments

cs

One area of memory

ct

Another area of memory

count

The size of the area.

memscan

Name

memscan — Find a character in an area of memory.

Synopsis

```
void * memscan (void * addr, int c, size_t size);
```

Arguments

addr

The memory area

c

The byte to search for

size

The size of the area.

Description

returns the address of the first occurrence of *c*, or 1 byte past the area if *c* is not found

strstr

Name

strstr — Find the first substring in a NUL terminated string

Synopsis

```
char * strstr (const char * s1, const char * s2);
```

Arguments

s1

The string to be searched

s2

The string to search for

memchr

Name

`memchr` — Find a character in an area of memory.

Synopsis

```
void * memchr (const void * s, int c, size_t n);
```

Arguments

s

The memory area

c

The byte to search for

n

The size of the area.

Description

returns the address of the first occurrence of *c*, or `NULL` if *c* is not found

3.3. Bit Operations

set_bit

Name

`set_bit` — Atomically set a bit in memory

Synopsis

```
void set_bit (int nr, volatile void * addr);
```

Arguments

nr

the bit to set

addr

the address to start counting from

Description

This function is atomic and may not be reordered. See `__set_bit` if you do not require the atomic guarantees. Note that *`nr`* may be almost arbitrarily large; this function is not restricted to acting on a single-word quantity.

__set_bit

Name

`__set_bit` — Set a bit in memory

Synopsis

```
void __set_bit (int nr, volatile void * addr);
```

Arguments

`nr`

the bit to set

`addr`

the address to start counting from

Description

Unlike `set_bit`, this function is non-atomic and may be reordered. If it's called on the same region of memory simultaneously, the effect may be that only one operation succeeds.

clear_bit

Name

`clear_bit` — Clears a bit in memory

Synopsis

```
void clear_bit (int nr, volatile void * addr);
```

Arguments

`nr`

Bit to clear

`addr`

Address to start counting from

Description

`clear_bit` is atomic and may not be reordered. However, it does not contain a memory barrier, so if it is used for locking purposes, you should call `smp_mb__before_clear_bit` and/or `smp_mb__after_clear_bit` in order to ensure changes are visible on other processors.

__change_bit

Name

__change_bit — Toggle a bit in memory

Synopsis

```
void __change_bit (int nr, volatile void * addr);
```

Arguments

nr

the bit to set

addr

the address to start counting from

Description

Unlike change_bit, this function is non-atomic and may be reordered. If it's called on the same region of memory simultaneously, the effect may be that only one operation succeeds.

change_bit

Name

`change_bit` — Toggle a bit in memory

Synopsis

```
void change_bit (int nr, volatile void * addr);
```

Arguments

`nr`

Bit to clear

`addr`

Address to start counting from

Description

`change_bit` is atomic and may not be reordered. Note that *`nr`* may be almost arbitrarily large; this function is not restricted to acting on a single-word quantity.

test_and_set_bit

Name

`test_and_set_bit` — Set a bit and return its old value

Synopsis

```
int test_and_set_bit (int nr, volatile void * addr);
```

Arguments

nr

Bit to set

addr

Address to count from

Description

This operation is atomic and cannot be reordered. It also implies a memory barrier.

__test_and_set_bit

Name

`__test_and_set_bit` — Set a bit and return its old value

Synopsis

```
int __test_and_set_bit (int nr, volatile void * addr);
```

Arguments

nr

Bit to set

addr

Address to count from

Description

This operation is non-atomic and can be reordered. If two examples of this operation race, one can appear to succeed but actually fail. You must protect multiple accesses with a lock.

test_and_clear_bit

Name

`test_and_clear_bit` — Clear a bit and return its old value

Synopsis

```
int test_and_clear_bit (int nr, volatile void * addr);
```

Arguments

nr

Bit to set

addr

Address to count from

Description

This operation is atomic and cannot be reordered. It also implies a memory barrier.

__test_and_clear_bit

Name

__test_and_clear_bit — Clear a bit and return its old value

Synopsis

```
int __test_and_clear_bit (int nr, volatile void * addr);
```

Arguments

nr

Bit to set

addr

Address to count from

Description

This operation is non-atomic and can be reordered. If two examples of this operation race, one can appear to succeed but actually fail. You must protect multiple accesses with a lock.

test_and_change_bit

Name

`test_and_change_bit` — Change a bit and return its new value

Synopsis

```
int test_and_change_bit (int nr, volatile void * addr);
```

Arguments

nr

Bit to set

addr

Address to count from

Description

This operation is atomic and cannot be reordered. It also implies a memory barrier.

test_bit

Name

`test_bit` — Determine whether a bit is set

Synopsis

```
int test_bit (int nr, const volatile void * addr);
```

Arguments

nr

bit number to test

addr

Address to start counting from

find_first_zero_bit

Name

`find_first_zero_bit` — find the first zero bit in a memory region

Synopsis

```
int find_first_zero_bit (void * addr, unsigned size);
```

Arguments

`addr`

The address to start the search at

`size`

The maximum size to search

Description

Returns the bit-number of the first zero bit, not the number of the byte containing a bit.

find_next_zero_bit

Name

`find_next_zero_bit` — find the first zero bit in a memory region

Synopsis

```
int find_next_zero_bit (void * addr, int size, int offset);
```

Arguments

addr

The address to base the search on

size

The maximum size to search

offset

The bitnumber to start searching at

ffz

Name

ffz — find first zero in word.

Synopsis

```
unsigned long ffz (unsigned long word);
```

Arguments

word

The word to search

Description

Undefined if no zero exists, so code should check against ~0UL first.

ffs

Name

ffs — find first bit set

Synopsis

```
int ffs (int x);
```

Arguments

x

the word to search

Description

This is defined the same way as the libc and compiler builtin ffs routines, therefore differs in spirit from the above ffz (man ffs).

hweight32

Name

`hweight32` — returns the hamming weight of a N-bit word

Synopsis

```
hweight32 ( x);
```

Arguments

x

the word to weigh

Description

The Hamming Weight of a number is the total number of bits set in it.

Chapter 4. Memory Management in Linux

4.1. The Slab Cache

kmem_cache_create

Name

`kmem_cache_create` — Create a cache.

Synopsis

```
kmem_cache_t * kmem_cache_create (const char * name, size_t
size, size_t offset, unsigned long flags, void (*ctor) (void*,
kmem_cache_t *, unsigned long), void (*dtor) (void*,
kmem_cache_t *, unsigned long));
```

Arguments

`name`

A string which is used in /proc/slabinfo to identify this cache.

`size`

The size of objects to be created in this cache.

`offset`

The offset to use within the page.

flags

SLAB flags

ctor

A constructor for the objects.

dtor

A destructor for the objects.

Description

Returns a ptr to the cache on success, NULL on failure. Cannot be called within a int, but can be interrupted. The *ctor* is run when new pages are allocated by the cache and the *dtor* is run before the pages are handed back. The flags are

SLAB_POISON - Poison the slab with a known test pattern (a5a5a5a5) to catch references to uninitialised memory.

SLAB_RED_ZONE - Insert ‘Red’ zones around the allocated memory to check for buffer overruns.

SLAB_NO_REAP - Don’t automatically reap this cache when we’re under memory pressure.

SLAB_HWCACHE_ALIGN - Align the objects in this cache to a hardware cacheline. This can be beneficial if you’re counting cycles as closely as davem.

kmem_cache_shrink

Name

kmem_cache_shrink — Shrink a cache.

Synopsis

```
int kmem_cache_shrink (kmem_cache_t * cachep);
```

Arguments

cachep

The cache to shrink.

Description

Releases as many slabs as possible for a cache. To help debugging, a zero exit status indicates all slabs were released.

kmem_cache_destroy

Name

kmem_cache_destroy — delete a cache

Synopsis

```
int kmem_cache_destroy (kmem_cache_t * cachep);
```

Arguments

cachep

the cache to destroy

Description

Remove a kmem_cache_t object from the slab cache. Returns 0 on success.

It is expected this function will be called by a module when it is unloaded. This will remove the cache completely, and avoid a duplicate cache being allocated each time a module is loaded and unloaded, if the module doesn't have persistent in-kernel storage across loads and unloads.

The caller must guarantee that noone will allocate memory from the cache during the `kmem_cache_destroy`.

kmem_cache_alloc

Name

`kmem_cache_alloc` — Allocate an object

Synopsis

```
void * kmem_cache_alloc (kmem_cache_t * cachep, int flags);
```

Arguments

cachep

The cache to allocate from.

flags

See kmalloc.

Description

Allocate an object from this cache. The flags are only relevant if the cache has no available objects.

kmalloc

Name

kmalloc — allocate memory

Synopsis

```
void * kmalloc (size_t size, int flags);
```

Arguments

size

how many bytes of memory are required.

flags

the type of memory to allocate.

Description

kmalloc is the normal method of allocating memory in the kernel.

The *flags* argument may be one of:

GFP_USER - Allocate memory on behalf of user. May sleep.

GFP_KERNEL - Allocate normal kernel ram. May sleep.

GFP_ATOMIC - Allocation will not sleep. Use inside interrupt handlers.

Additionally, the GFP_DMA flag may be set to indicate the memory must be suitable for DMA. This can mean different things on different platforms. For example, on i386, it means that the memory must come from the first 16MB.

kmem_cache_free

Name

kmem_cache_free — Deallocate an object

Synopsis

```
void kmem_cache_free (kmem_cache_t * cachep, void * objp);
```

Arguments

cachep

The cache the allocation was from.

objp

The previously allocated object.

Description

Free an object which was previously allocated from this cache.

kfree

Name

kfree — free previously allocated memory

Synopsis

```
void kfree (const void * objp);
```

Arguments

objp

pointer returned by kmalloc.

Description

Don’t free memory not originally allocated by `kmalloc` or you will run into trouble.

Chapter 5. The proc filesystem

5.1. sysctl interface

register_sysctl_table

Name

`register_sysctl_table` — register a sysctl heirarchy

Synopsis

```
struct ctl_table_header * register_sysctl_table (ctl_table *
table, int insert_at_head);
```

Arguments

`table`

the top-level table structure

`insert_at_head`

whether the entry should be inserted in front or at the end

Description

Register a sysctl table heirarchy. *`table`* should be a filled in ctl_table array. An entry with a ctl_name of 0 terminates the table.

The members of the &ctl_table structure are used as follows:

ctl_name - This is the numeric sysctl value used by sysctl(2). The number must be unique within that level of sysctl

procname - the name of the sysctl file under /proc/sys. Set to `NULL` to not enter a sysctl file

data - a pointer to data for use by proc_handler

maxlen - the maximum size in bytes of the data

mode - the file permissions for the /proc/sys file, and for sysctl(2)

child - a pointer to the child sysctl table if this entry is a directory, or `NULL`.

proc_handler - the text handler routine (described below)

strategy - the strategy routine (described below)

de - for internal use by the sysctl routines

extra1, extra2 - extra pointers usable by the proc handler routines

Leaf nodes in the sysctl tree will be represented by a single file under /proc; non-leaf nodes will be represented by directories.

sysctl(2) can automatically manage read and write requests through the sysctl table. The data and maxlen fields of the ctl_table struct enable minimal validation of the values being written to be performed, and the mode field allows minimal authentication.

More sophisticated management can be enabled by the provision of a strategy routine with the table entry. This will be called before any automatic read or write of the data is performed.

The strategy routine may return

< 0 - Error occurred (error is passed to user process)

0 - OK - proceed with automatic read or write.

> 0 - OK - read or write has been done by the strategy routine, so return immediately.

There must be a proc_handler routine for any terminal nodes mirrored under /proc/sys (non-terminals are handled by a built-in directory handler). Several default handlers are available to cover common cases -

`proc_dostring`, `proc_dointvec`, `proc_dointvec_jiffies`, `proc_dointvec_minmax`, `proc_doulongvec_ms_jiffies_minmax`, `proc_doulongvec_minmax`

It is the handler’s job to read the input buffer from user memory and process it. The handler should return 0 on success.

This routine returns `NULL` on a failure to register, and a pointer to the table header on success.

unregister_sysctl_table

Name

`unregister_sysctl_table` — unregister a sysctl table heirarchy

Synopsis

```
void unregister_sysctl_table (struct ctl_table_header * header);
```

Arguments

`header`

the header returned from register_sysctl_table

Description

Unregisters the sysctl table and all children. proc entries may not actually be removed until they are no longer used by anyone.

proc_dostring

Name

`proc_dostring` — read a string sysctl

Synopsis

```
int proc_dostring (ctl_table * table, int write, struct file *
filp, void * buffer, size_t * lenp);
```

Arguments

`table`

the sysctl table

`write`

`TRUE` if this is a write to the sysctl file

`filp`

the file structure

`buffer`

the user buffer

`lenp`

the size of the user buffer

Description

Reads/writes a string from/to the user buffer. If the kernel buffer provided is not large enough to hold the string, the string is truncated. The copied string is

`NULL-terminated`. If the string is being read by the user process, it is copied and a newline '\n' is added. It is truncated if the buffer is not large enough.

Returns 0 on success.

proc_dointvec

Name

`proc_dointvec` — read a vector of integers

Synopsis

```
int proc_dointvec (ctl_table * table, int write, struct file *
filp, void * buffer, size_t * lenp);
```

Arguments

`table`

the sysctl table

`write`

`TRUE` if this is a write to the sysctl file

`filp`

the file structure

`buffer`

the user buffer

lenp

the size of the user buffer

Description

Reads/writes up to table->maxlen/sizeof(unsigned int) integer values from/to the user buffer, treated as an ASCII string.

Returns 0 on success.

proc_dointvec_minmax

Name

proc_dointvec_minmax — read a vector of integers with min/max values

Synopsis

```
int proc_dointvec_minmax (ctl_table * table, int write, struct
file * filp, void * buffer, size_t * lenp);
```

Arguments

table

the sysctl table

write

TRUE if this is a write to the sysctl file

filp

the file structure

buffer

the user buffer

lenp

the size of the user buffer

Description

Reads/writes up to table->maxlen/sizeof(unsigned int) integer values from/to the user buffer, treated as an ASCII string.

This routine will ensure the values are within the range specified by table->extra1 (min) and table->extra2 (max).

Returns 0 on success.

proc_doulongvec_minmax

Name

proc_doulongvec_minmax — read a vector of long integers with min/max values

Synopsis

```
int proc_doulongvec_minmax (ctl_table * table, int write, struct
file * filp, void * buffer, size_t * lenp);
```

Arguments

`table`

the sysctl table

`write`

`TRUE` if this is a write to the sysctl file

`filp`

the file structure

`buffer`

the user buffer

`lenp`

the size of the user buffer

Description

Reads/writes up to table->maxlen/sizeof(unsigned long) unsigned long values from/to the user buffer, treated as an ASCII string.

This routine will ensure the values are within the range specified by table->extra1 (min) and table->extra2 (max).

Returns 0 on success.

proc_doulongvec_ms_jiffies_minmax

Name

`proc_doulongvec_ms_jiffies_minmax` — read a vector of millisecond values with min/max values

Synopsis

```
int proc_doulongvec_ms_jiffies_minmax (ctl_table * table, int
write, struct file * filp, void * buffer, size_t * lenp);
```

Arguments

table

the sysctl table

write

TRUE if this is a write to the sysctl file

filp

the file structure

buffer

the user buffer

lenp

the size of the user buffer

Description

Reads/writes up to table->maxlen/sizeof(unsigned long) unsigned long values from/to the user buffer, treated as an ASCII string. The values are treated as milliseconds, and converted to jiffies when they are stored.

This routine will ensure the values are within the range specified by table->extra1 (min) and table->extra2 (max).

Returns 0 on success.

proc_dointvec_jiffies

Name

`proc_dointvec_jiffies` — read a vector of integers as seconds

Synopsis

```
int proc_dointvec_jiffies (ctl_table * table, int write, struct
file * filp, void * buffer, size_t * lenp);
```

Arguments

`table`

the sysctl table

`write`

`TRUE` if this is a write to the sysctl file

`filp`

the file structure

`buffer`

the user buffer

`lenp`

the size of the user buffer

Description

Reads/writes up to table->maxlen/sizeof(unsigned int) integer values from/to the user buffer, treated as an ASCII string. The values read are assumed to be in seconds, and

are converted into jiffies.

Returns 0 on success.

Chapter 6. The Linux VFS

6.1. The Directory Cache

d_invalidate

Name

`d_invalidate` — invalidate a dentry

Synopsis

```
int d_invalidate (struct dentry * dentry);
```

Arguments

`dentry`

dentry to invalidate

Description

Try to invalidate the dentry if it turns out to be possible. If there are other dentries that can be reached through this one we can’t delete it and we return -EBUSY. On success we return 0.

no dcache lock.

d_find_alias

Name

d_find_alias — grab a hashed alias of inode

Synopsis

```
struct dentry * d_find_alias (struct inode * inode);
```

Arguments

inode

inode in question

Description

If inode has a hashed alias - acquire the reference to alias and return it. Otherwise return NULL. Notice that if inode is a directory there can be only one alias and it can be unhashed only if it has no children.

prune_dcache

Name

prune_dcache — shrink the dcache

Synopsis

```
void prune_dcache (int count);
```

Arguments

count

number of entries to try and free

Description

Shrink the dcache. This is done when we need more memory, or simply when we need to unmount something (at which point we need to unuse all dentries).

This function may fail to free any resources if all the dentries are in use.

shrink_dcache_sb

Name

shrink_dcache_sb — shrink dcache for a superblock

Synopsis

```
void shrink_dcache_sb (struct super_block * sb);
```

Arguments

sb

superblock

Description

Shrink the dcache for the specified super block. This is used to free the dcache before unmounting a file system

have_submounts

Name

have_submounts — check for mounts over a dentry

Synopsis

```
int have_submounts (struct dentry * parent);
```

Arguments

parent

dentry to check.

Description

Return true if the parent or its subdirectories contain a mount point

shrink_dcache_parent

Name

`shrink_dcache_parent` — prune dcache

Synopsis

```
void shrink_dcache_parent (struct dentry * parent);
```

Arguments

`parent`

parent of entries to prune

Description

Prune the dcache to remove unused children of the parent dentry.

d_alloc

Name

d_alloc — allocate a dcache entry

Synopsis

```
struct dentry * d_alloc (struct dentry * parent, const struct
qstr * name);
```

Arguments

parent

parent of entry to allocate

name

qstr of the name

Description

Allocates a dentry. It returns NULL if there is insufficient memory available. On a success the dentry is returned. The name passed in is copied and the copy passed in may be reused after this call.

d_instantiate

Name

`d_instantiate` — fill in inode information for a dentry

Synopsis

```
void d_instantiate (struct dentry * entry, struct inode *
inode);
```

Arguments

`entry`

dentry to complete

`inode`

inode to attach to this dentry

Description

Fill in inode information in the entry.

This turns negative dentries into productive full members of society.

NOTE! This assumes that the inode count has been incremented (or otherwise set) by the caller to indicate that it is now in use by the dcache.

d_alloc_root

Name

d_alloc_root — allocate root dentry

Synopsis

```
struct dentry * d_alloc_root (struct inode * root_inode);
```

Arguments

root_inode

inode to allocate the root for

Description

Allocate a root ("/") dentry for the inode given. The inode is instantiated and returned. NULL is returned if there is insufficient memory or the inode passed is NULL.

d_lookup

Name

d_lookup — search for a dentry

Synopsis

```
struct dentry * d_lookup (struct dentry * parent, struct qstr *
name);
```

Arguments

parent

parent dentry

name

qstr of name we wish to find

Description

Searches the children of the parent dentry for the name in question. If the dentry is found its reference count is incremented and the dentry is returned. The caller must use d_put to free the entry when it has finished using it. NULL is returned on failure.

d_validate

Name

d_validate — verify dentry provided from insecure source

Synopsis

```
int d_validate (struct dentry * dentry, struct dentry *
dparent);
```

Arguments

dentry

The dentry alleged to be valid child of *dparent*

dparent

The parent dentry (known to be valid)

Description

An insecure source has sent us a dentry, here we verify it and dget it. This is used by ncpfs in its readdir implementation. Zero is returned in the dentry is invalid.

d_delete

Name

d_delete — delete a dentry

Synopsis

```
void d_delete (struct dentry * dentry);
```

Arguments

dentry

The dentry to delete

Description

Turn the dentry into a negative dentry if possible, otherwise remove it from the hash queues so it can be deleted later

d_rehash

Name

d_rehash — add an entry back to the hash

Synopsis

```
void d_rehash (struct dentry * entry);
```

Arguments

entry

dentry to add to the hash

Description

Adds a dentry to the hash according to its name.

d_move

Name

d_move — move a dentry

Synopsis

```
void d_move (struct dentry * dentry, struct dentry * target);
```

Arguments

dentry

entry to move

target

new dentry

Description

Update the dcache to reflect the move of a file name. Negative dcache entries should not be moved in this way.

__d_path

Name

__d_path — return the path of a dentry

Synopsis

```
char * __d_path (struct dentry * dentry, struct vfsmount *
vfsmnt, struct dentry * root, struct vfsmount * rootmnt, char *
buffer, int buflen);
```

Arguments

dentry

dentry to report

vfsmnt

vfsmnt to which the dentry belongs

root

root dentry

rootmnt

vfsmnt to which the root dentry belongs

buffer

buffer to return value in

buflen

buffer length

Description

Convert a dentry into an ASCII path name. If the entry has been deleted the string " (deleted)" is appended. Note that this is ambiguous. Returns the buffer.

"buflen" should be `PAGE_SIZE` or more. Caller holds the dcache_lock.

is_subdir

Name

`is_subdir` — is new dentry a subdirectory of old_dentry

Synopsis

```
int is_subdir (struct dentry * new_dentry, struct dentry *
old_dentry);
```

Arguments

`new_dentry`

new dentry

`old_dentry`

old dentry

Description

Returns 1 if new_dentry is a subdirectory of the parent (at any depth). Returns 0 otherwise.

find_inode_number

Name

`find_inode_number` — check for dentry with name

Synopsis

```
ino_t find_inode_number (struct dentry * dir, struct qstr *
name);
```

Arguments

`dir`

directory to check

`name`

Name to find.

Description

Check whether a dentry already exists for the given name, and return the inode number if it has an inode. Otherwise 0 is returned.

This routine is used to post-process directory listings for filesystems using synthetic inode numbers, and is necessary to keep `getcwd` working.

d_drop

Name

d_drop — drop a dentry

Synopsis

```
void d_drop (struct dentry * dentry);
```

Arguments

dentry

dentry to drop

Description

d_drop unhashes the entry from the parent dentry hashes, so that it won't be found through a VFS lookup any more. Note that this is different from deleting the dentry - d_delete will try to mark the dentry negative if possible, giving a successful _negative_ lookup, while d_drop will just make the cache lookup fail.

d_drop is used mainly for stuff that wants to invalidate a dentry for some reason (NFS timeouts or autofs deletes).

d_add

Name

d_add — add dentry to hash queues

Synopsis

```
void d_add (struct dentry * entry, struct inode * inode);
```

Arguments

entry

dentry to add

inode

The inode to attach to this dentry

Description

This adds the entry to the hash queues and initializes *inode*. The entry was actually filled in earlier during d_alloc.

dget

Name

dget — get a reference to a dentry

Synopsis

```
struct dentry * dget (struct dentry * dentry);
```

Arguments

dentry

dentry to get a reference to

Description

Given a dentry or NULL pointer increment the reference count if appropriate and return the dentry. A dentry will not be destroyed when it has references. dget should never be called for dentries with zero reference counter. For these cases (preferably none, functions in dcache.c are sufficient for normal needs and they take necessary precautions) you should hold dcache_lock and call dget_locked instead of dget.

d_unhashed

Name

d_unhashed — is dentry hashed

Synopsis

```
int d_unhashed (struct dentry * dentry);
```

Arguments

dentry

entry to check

Description

Returns true if the dentry passed is not currently hashed.

6.2. Inode Handling

__mark_inode_dirty

Name

__mark_inode_dirty — internal function

Synopsis

```
void __mark_inode_dirty (struct inode * inode, int flags);
```

Arguments

inode

inode to mark

flags

what kind of dirty (i.e. I_DIRTY_SYNC) Mark an inode as dirty. Callers should use mark_inode_dirty or mark_inode_dirty_sync.

write_inode_now

Name

`write_inode_now` — write an inode to disk

Synopsis

```
void write_inode_now (struct inode * inode, int sync);
```

Arguments

inode

inode to write to disk

sync

whether the write should be synchronous or not

Description

This function commits an inode to disk immediately if it is dirty. This is primarily needed by knfsd.

clear_inode

Name

`clear_inode` — clear an inode

Synopsis

```
void clear_inode (struct inode * inode);
```

Arguments

`inode`

inode to clear

Description

This is called by the filesystem to tell us that the inode is no longer useful. We just terminate it with extreme prejudice.

invalidate_inodes

Name

`invalidate_inodes` — discard the inodes on a device

Synopsis

```
int invalidate_inodes (struct super_block * sb);
```

Arguments

sb

superblock

Description

Discard all of the inodes for a given superblock. If the discard fails because there are busy inodes then a non zero value is returned. If the discard is successful all the inodes have been discarded.

get_empty_inode

Name

`get_empty_inode` — obtain an inode

Synopsis

```
struct inode * get_empty_inode ( void);
```

Arguments

void

no arguments

Description

This is called by things like the networking layer etc that want to get an inode without any inode number, or filesystems that allocate new inodes with no pre-existing information.

On a successful return the inode pointer is returned. On a failure a `NULL` pointer is returned. The returned inode is not on any superblock lists.

iunique

Name

`iunique` — get a unique inode number

Synopsis

```
ino_t iunique (struct super_block * sb, ino_t max_reserved);
```

Arguments

sb

superblock

max_reserved

highest reserved inode number

Description

Obtain an inode number that is unique on the system for a given superblock. This is used by file systems that have no natural permanent inode numbering system. An inode number is returned that is higher than the reserved limit but unique.

BUGS

With a large number of inodes live on the file system this function currently becomes quite slow.

insert_inode_hash

Name

`insert_inode_hash` — hash an inode

Synopsis

```
void insert_inode_hash (struct inode * inode);
```

Arguments

inode

unhashed inode

Description

Add an inode to the inode hash for this superblock. If the inode has no superblock it is added to a separate anonymous chain.

remove_inode_hash

Name

remove_inode_hash — remove an inode from the hash

Synopsis

```
void remove_inode_hash (struct inode * inode);
```

Arguments

inode

inode to unhash

Description

Remove an inode from the superblock or anonymous hash.

iput

Name

`iput` — put an inode

Synopsis

```
void iput (struct inode * inode);
```

Arguments

`inode`

inode to put

Description

Puts an inode, dropping its usage count. If the inode use count hits zero the inode is also then freed and may be destroyed.

bmap

Name

`bmap` — find a block number in a file

Synopsis

```
int bmap (struct inode * inode, int block);
```

Arguments

`inode`

inode of file

`block`

block to find

Description

Returns the block number on the device holding the inode that is the disk block number for the block of the file requested. That is, asked for block 4 of inode 1 the function will return the disk block relative to the disk start that holds that block of the file.

update_atime

Name

update_atime — update the access time

Synopsis

```
void update_atime (struct inode * inode);
```

Arguments

inode

inode accessed

Description

Update the accessed time on an inode and mark it for writeback. This function automatically handles read only file systems and media, as well as the “noatime” flag and inode specific “noatime” markers.

make_bad_inode

Name

make_bad_inode — mark an inode bad due to an I/O error

Synopsis

```
void make_bad_inode (struct inode * inode);
```

Arguments

inode

Inode to mark bad

Description

When an inode cannot be read due to a media or remote network failure this function makes the inode “bad” and causes I/O operations on it to fail from this point on.

is_bad_inode

Name

`is_bad_inode` — is an inode errored

Synopsis

```
int is_bad_inode (struct inode * inode);
```

Arguments

inode

inode to test

Description

Returns true if the inode in question has been marked as bad.

6.3. Registration and Superblocks

register_filesystem

Name

`register_filesystem` — register a new filesystem

Synopsis

```
int register_filesystem (struct file_system_type * fs);
```

Arguments

fs

the file system structure

Description

Adds the file system passed to the list of file systems the kernel is aware of for mount and other syscalls. Returns 0 on success, or a negative errno code on an error.

The &struct file_system_type that is passed is linked into the kernel structures and must not be freed until the file system has been unregistered.

unregister_filesystem

Name

`unregister_filesystem` — unregister a file system

Synopsis

```
int unregister_filesystem (struct file_system_type * fs);
```

Arguments

`fs`

filesystem to unregister

Description

Remove a file system that was previously successfully registered with the kernel. An error is returned if the file system is not found. Zero is returned on a success.

Once this function has returned the &struct file_system_type structure may be freed or reused.

get_super

Name

get_super — get the superblock of a device

Synopsis

```
struct super_block * get_super (kdev_t dev);
```

Arguments

dev

device to get the superblock for

Description

Scans the superblock list and finds the superblock of the file system mounted on the device given. NULL is returned if no match is found.

6.4. File Locks

posix_lock_file

Name

`posix_lock_file` —

Synopsis

```
int posix_lock_file (struct file * filp, struct file_lock *
caller, unsigned int wait);
```

Arguments

`filp`

The file to apply the lock to

`caller`

The lock to be applied

`wait`

1 to retry automatically, 0 to return -EAGAIN

Description

Add a POSIX style lock to a file. We merge adjacent locks whenever possible. POSIX locks are sorted by owner task, then by starting address

Kai Petzke writes

To make freeing a lock much faster, we keep a pointer to the lock before the actual one. But the real gain of the new coding was, that `lock_it` and `unlock_it` became one function.

To all purists

Yes, I use a few goto's. Just pass on to the next function.

__get_lease

Name

`__get_lease` — revoke all outstanding leases on file

Synopsis

```
int __get_lease (struct inode * inode, unsigned int mode);
```

Arguments

`inode`

the inode of the file to return

`mode`

the open mode (read or write)

Description

get_lease (inlined for speed) has checked there already is a lease on this file. Leases are broken on a call to `open` or `truncate`. This function can sleep unless you specified `O_NONBLOCK` to your `open`.

lease_get_mtime

Name

`lease_get_mtime` —

Synopsis

```
time_t lease_get_mtime (struct inode * inode);
```

Arguments

`inode`

the inode

Description

This is to force NFS clients to flush their caches for files with exclusive leases. The justification is that if someone has an exclusive lease, then they could be modifiying it.

posix_block_lock

Name

posix_block_lock — blocks waiting for a file lock

Synopsis

```
void posix_block_lock (struct file_lock * blocker, struct
file_lock * waiter);
```

Arguments

blocker

the lock which is blocking

waiter

the lock which conflicts and has to wait

Description

lockd needs to block waiting for locks.

posix_unblock_lock

Name

posix_unblock_lock — stop waiting for a file lock

Synopsis

```
void posix_unblock_lock (struct file_lock * waiter);
```

Arguments

waiter

the lock which was waiting

Description

lockd needs to block waiting for locks.

lock_may_read

Name

`lock_may_read` — checks that the region is free of locks

Synopsis

```
int lock_may_read (struct inode * inode, loff_t start, unsigned
long len);
```

Arguments

inode

the inode that is being read

start

the first byte to read

len

the number of bytes to read

Description

Emulates Windows locking requirements. Whole-file mandatory locks (share modes) can prohibit a read and byte-range POSIX locks can prohibit a read if they overlap.

N.B. this function is only ever called from knfsd and ownership of locks is never checked.

lock_may_write

Name

`lock_may_write` — checks that the region is free of locks

Synopsis

```
int lock_may_write (struct inode * inode, loff_t start, unsigned
long len);
```

Arguments

inode

the inode that is being written

start

the first byte to write

len

the number of bytes to write

Description

Emulates Windows locking requirements. Whole-file mandatory locks (share modes) can prohibit a write and byte-range POSIX locks can prohibit a write if they overlap.

N.B. this function is only ever called from knfsd and ownership of locks is never checked.

fcntl_getlease

Name

`fcntl_getlease` — Enquire what lease is currently active

Synopsis

```
int fcntl_getlease (struct file * filp);
```

Arguments

filp

the file

Description

The value returned by this function will be one of

F_RDLCK to indicate a read-only (type II) lease is held.

F_WRLCK to indicate an exclusive lease is held.

XXX

sfr & i disagree over whether F_INPROGRESS should be returned to userspace.

fcntl_setlease

Name

fcntl_setlease — sets a lease on an open file

Synopsis

```
int fcntl_setlease (unsigned int fd, struct file * filp, long
arg);
```

Arguments

fd

open file descriptor

filp

file pointer

arg

type of lease to obtain

Description

Call this fcntl to establish a lease on the file. Note that you also need to call F_SETSIG to receive a signal when the lease is broken.

sys_flock

Name

sys_flock — flock system call.

Synopsis

```
asmlinkage long sys_flock (unsigned int fd, unsigned int cmd);
```

Arguments

fd

the file descriptor to lock.

cmd

the type of lock to apply.

Description

Apply a FL_FLOCK style lock to an open file descriptor. The *cmd* can be one of

LOCK_SH -- a shared lock.

LOCK_EX -- an exclusive lock.

LOCK_UN -- remove an existing lock.

LOCK_MAND -- a ‘mandatory’ flock. This exists to emulate Windows Share Modes.

LOCK_MAND can be combined with LOCK_READ or LOCK_WRITE to allow other processes read and write access respectively.

get_locks_status

Name

get_locks_status — reports lock usage in /proc/locks

Synopsis

```
int get_locks_status (char * buffer, char ** start, off_t
offset, int length);
```

Arguments

buffer

address in userspace to write into

start

?

offset

how far we are through the buffer

length

how much to read

Chapter 7. Linux Networking

7.1. Socket Buffer Functions

skb_queue_empty

Name

`skb_queue_empty` — check if a queue is empty

Synopsis

```
int skb_queue_empty (struct sk_buff_head * list);
```

Arguments

`list`

queue head

Description

Returns true if the queue is empty, false otherwise.

skb_get

Name

skb_get — reference buffer

Synopsis

```
struct sk_buff * skb_get (struct sk_buff * skb);
```

Arguments

skb

buffer to reference

Description

Makes another reference to a socket buffer and returns a pointer to the buffer.

kfree_skb

Name

kfree_skb — free an sk_buff

Synopsis

```
void kfree_skb (struct sk_buff * skb);
```

Arguments

skb

buffer to free

Description

Drop a reference to the buffer and free it if the usage count has hit zero.

skb_cloned

Name

skb_cloned — is the buffer a clone

Synopsis

```
int skb_cloned (struct sk_buff * skb);
```

Arguments

skb

buffer to check

Description

Returns true if the buffer was generated with `skb_clone` and is one of multiple shared copies of the buffer. Cloned buffers are shared data so must not be written to under normal circumstances.

skb_shared

Name

`skb_shared` — is the buffer shared

Synopsis

```
int skb_shared (struct sk_buff * skb);
```

Arguments

skb

buffer to check

Description

Returns true if more than one person has a reference to this buffer.

skb_share_check

Name

skb_share_check — check if buffer is shared and if so clone it

Synopsis

```
struct sk_buff * skb_share_check (struct sk_buff * skb, int
pri);
```

Arguments

skb

buffer to check

pri

priority for memory allocation

Description

If the buffer is shared the buffer is cloned and the old copy drops a reference. A new clone with a single reference is returned. If the buffer is not shared the original buffer is returned. When being called from interrupt status or with spinlocks held pri must be GFP_ATOMIC.

NULL is returned on a memory allocation failure.

skb_unshare

Name

skb_unshare — make a copy of a shared buffer

Synopsis

```
struct sk_buff * skb_unshare (struct sk_buff * skb, int pri);
```

Arguments

skb

buffer to check

pri

priority for memory allocation

Description

If the socket buffer is a clone then this function creates a new copy of the data, drops a reference count on the old copy and returns the new copy with the reference count at 1. If the buffer is not a clone the original buffer is returned. When called with a spinlock held or from interrupt state *pri* must be GFP_ATOMIC

NULL is returned on a memory allocation failure.

skb_peek

Name

skb_peek —

Synopsis

```
struct sk_buff * skb_peek (struct sk_buff_head * list_);
```

Arguments

*list_*

list to peek at

Description

Peek an &sk_buff. Unlike most other operations you _MUST_ be careful with this one. A peek leaves the buffer on the list and someone else may run off with it. You must hold the appropriate locks or have a private queue to do this.

Returns NULL for an empty list or a pointer to the head element. The reference count is not incremented and the reference is therefore volatile. Use with caution.

skb_peek_tail

Name

`skb_peek_tail` —

Synopsis

```
struct sk_buff * skb_peek_tail (struct sk_buff_head * list_);
```

Arguments

`list_`

list to peek at

Description

Peek an &sk_buff. Unlike most other operations you _MUST_ be careful with this one. A peek leaves the buffer on the list and someone else may run off with it. You must hold the appropriate locks or have a private queue to do this.

Returns `NULL` for an empty list or a pointer to the tail element. The reference count is not incremented and the reference is therefore volatile. Use with caution.

skb_queue_len

Name

skb_queue_len — get queue length

Synopsis

```
__u32 skb_queue_len (struct sk_buff_head * list_);
```

Arguments

*list_*

list to measure

Description

Return the length of an &sk_buff queue.

__skb_queue_head

Name

__skb_queue_head — queue a buffer at the list head

Synopsis

```
void __skb_queue_head (struct sk_buff_head * list, struct
sk_buff * newsk);
```

Arguments

list

list to use

newsk

buffer to queue

Description

Queue a buffer at the start of a list. This function takes no locks and you must therefore hold required locks before calling it.

A buffer cannot be placed on two lists at the same time.

skb_queue_head

Name

skb_queue_head — queue a buffer at the list head

Synopsis

```
void skb_queue_head (struct sk_buff_head * list, struct sk_buff
* newsk);
```

Arguments

list

list to use

newsk

buffer to queue

Description

Queue a buffer at the start of the list. This function takes the list lock and can be used safely with other locking &sk_buff functions safely.

A buffer cannot be placed on two lists at the same time.

__skb_queue_tail

Name

__skb_queue_tail — queue a buffer at the list tail

Synopsis

```
void __skb_queue_tail (struct sk_buff_head * list, struct
sk_buff * newsk);
```

Arguments

list

list to use

newsk

buffer to queue

Description

Queue a buffer at the end of a list. This function takes no locks and you must therefore hold required locks before calling it.

A buffer cannot be placed on two lists at the same time.

skb_queue_tail

Name

skb_queue_tail — queue a buffer at the list tail

Synopsis

```
void skb_queue_tail (struct sk_buff_head * list, struct sk_buff
* newsk);
```

Arguments

list

list to use

newsk

buffer to queue

Description

Queue a buffer at the tail of the list. This function takes the list lock and can be used safely with other locking &sk_buff functions safely.

A buffer cannot be placed on two lists at the same time.

__skb_dequeue

Name

__skb_dequeue — remove from the head of the queue

Synopsis

```
struct sk_buff * __skb_dequeue (struct sk_buff_head * list);
```

Arguments

list

list to dequeue from

Description

Remove the head of the list. This function does not take any locks so must be used with appropriate locks held only. The head item is returned or NULL if the list is empty.

skb_dequeue

Name

skb_dequeue — remove from the head of the queue

Synopsis

```
struct sk_buff * skb_dequeue (struct sk_buff_head * list);
```

Arguments

list

list to dequeue from

Description

Remove the head of the list. The list lock is taken so the function may be used safely with other locking list functions. The head item is returned or `NULL` if the list is empty.

skb_insert

Name

`skb_insert` — insert a buffer

Synopsis

```
void skb_insert (struct sk_buff * old, struct sk_buff * newsk);
```

Arguments

`old`

buffer to insert before

`newsk`

buffer to insert

Description

Place a packet before a given packet in a list. The list locks are taken and this function is atomic with respect to other list locked calls A buffer cannot be placed on two lists at the same time.

skb_append

Name

skb_append — append a buffer

Synopsis

```
void skb_append (struct sk_buff * old, struct sk_buff * newsk);
```

Arguments

old

buffer to insert after

newsk

buffer to insert

Description

Place a packet after a given packet in a list. The list locks are taken and this function is atomic with respect to other list locked calls. A buffer cannot be placed on two lists at the same time.

skb_unlink

Name

skb_unlink — remove a buffer from a list

Synopsis

```
void skb_unlink (struct sk_buff * skb);
```

Arguments

skb

buffer to remove

Description

Place a packet after a given packet in a list. The list locks are taken and this function is atomic with respect to other list locked calls

Works even without knowing the list it is sitting on, which can be handy at times. It also means that THE LIST MUST EXIST when you unlink. Thus a list must have its contents unlinked before it is destroyed.

__skb_dequeue_tail

Name

__skb_dequeue_tail — remove from the tail of the queue

Synopsis

```
struct sk_buff * __skb_dequeue_tail (struct sk_buff_head *
list);
```

Arguments

list

list to dequeue from

Description

Remove the tail of the list. This function does not take any locks so must be used with appropriate locks held only. The tail item is returned or NULL if the list is empty.

skb_dequeue_tail

Name

skb_dequeue_tail — remove from the head of the queue

Synopsis

```
struct sk_buff * skb_dequeue_tail (struct sk_buff_head * list);
```

Arguments

list

list to dequeue from

Description

Remove the head of the list. The list lock is taken so the function may be used safely with other locking list functions. The tail item is returned or NULL if the list is empty.

skb_put

Name

skb_put — add data to a buffer

Synopsis

```
unsigned char * skb_put (struct sk_buff * skb, unsigned int
len);
```

Arguments

skb

buffer to use

len

amount of data to add

Description

This function extends the used data area of the buffer. If this would exceed the total buffer size the kernel will panic. A pointer to the first byte of the extra data is returned.

skb_push

Name

skb_push — add data to the start of a buffer

Synopsis

```
unsigned char * skb_push (struct sk_buff * skb, unsigned int
len);
```

Arguments

skb

buffer to use

len

amount of data to add

Description

This function extends the used data area of the buffer at the buffer start. If this would exceed the total buffer headroom the kernel will panic. A pointer to the first byte of the extra data is returned.

skb_pull

Name

skb_pull — remove data from the start of a buffer

Synopsis

```
unsigned char * skb_pull (struct sk_buff * skb, unsigned int
len);
```

Arguments

skb

buffer to use

len

amount of data to remove

Description

This function removes data from the start of a buffer, returning the memory to the headroom. A pointer to the next data in the buffer is returned. Once the data has been pulled future pushes will overwrite the old data.

skb_headroom

Name

`skb_headroom` — bytes at buffer head

Synopsis

```
int skb_headroom (const struct sk_buff * skb);
```

Arguments

`skb`

buffer to check

Description

Return the number of bytes of free space at the head of an &sk_buff.

skb_tailroom

Name

`skb_tailroom` — bytes at buffer end

Synopsis

```
int skb_tailroom (const struct sk_buff * skb);
```

Arguments

`skb`

buffer to check

Description

Return the number of bytes of free space at the tail of an sk_buff

skb_reserve

Name

`skb_reserve` — adjust headroom

Synopsis

```
void skb_reserve (struct sk_buff * skb, unsigned int len);
```

Arguments

skb

buffer to alter

len

bytes to move

Description

Increase the headroom of an empty &sk_buff by reducing the tail room. This is only allowed for an empty buffer.

skb_trim

Name

skb_trim — remove end from a buffer

Synopsis

```
void skb_trim (struct sk_buff * skb, unsigned int len);
```

Arguments

skb

buffer to alter

len

new length

Description

Cut the length of a buffer down by removing data from the tail. If the buffer is already under the length specified it is not modified.

skb_orphan

Name

skb_orphan — orphan a buffer

Synopsis

```
void skb_orphan (struct sk_buff * skb);
```

Arguments

skb

buffer to orphan

Description

If a buffer currently has an owner then we call the owner's destructor function and make the *skb* unowned. The buffer continues to exist but is no longer charged to its former owner.

skb_queue_purge

Name

skb_queue_purge — empty a list

Synopsis

```
void skb_queue_purge (struct sk_buff_head * list);
```

Arguments

list

list to empty

Description

Delete all buffers on an &sk_buff list. Each buffer is removed from the list and one reference dropped. This function takes the list lock and is atomic with respect to other list locking functions.

__skb_queue_purge

Name

__skb_queue_purge — empty a list

Synopsis

```
void __skb_queue_purge (struct sk_buff_head * list);
```

Arguments

list

list to empty

Description

Delete all buffers on an &sk_buff list. Each buffer is removed from the list and one reference dropped. This function does not take the list lock and the caller must hold the relevant locks to use it.

__dev_alloc_skb

Name

__dev_alloc_skb — allocate an skbuff for sending

Synopsis

```
struct sk_buff * __dev_alloc_skb (unsigned int length, int
gfp_mask);
```

Arguments

length

length to allocate

gfp_mask

get_free_pages mask, passed to alloc_skb

Description

Allocate a new &sk_buff and assign it a usage count of one. The buffer has unspecified headroom built in. Users should allocate the headroom they think they need without accounting for the built in space. The built in space is used for optimisations.

`NULL` is returned in there is no free memory.

dev_alloc_skb

Name

`dev_alloc_skb` — allocate an skbuff for sending

Synopsis

```
struct sk_buff * dev_alloc_skb (unsigned int length);
```

Arguments

length

length to allocate

Description

Allocate a new &sk_buff and assign it a usage count of one. The buffer has unspecified headroom built in. Users should allocate the headroom they think they need without accounting for the built in space. The built in space is used for optimisations.

`NULL` is returned in there is no free memory. Although this function allocates memory it can be called from an interrupt.

skb_cow

Name

`skb_cow` — copy header of skb when it is required

Synopsis

```
int skb_cow (struct sk_buff * skb, unsigned int headroom);
```

Arguments

skb

buffer to cow

headroom

needed headroom

Description

If the skb passed lacks sufficient headroom or its data part is shared, data is reallocated. If reallocation fails, an error is returned and original skb is not changed.

The result is skb with writable area skb->head...skb->tail and at least *headroom* of space at head.

skb_over_panic

Name

skb_over_panic — private function

Synopsis

```
void skb_over_panic (struct sk_buff * skb, int sz, void * here);
```

Arguments

skb

buffer

sz

size

here

address

Description

Out of line support code for skb_put. Not user callable.

skb_under_panic

Name

skb_under_panic — private function

Synopsis

```
void skb_under_panic (struct sk_buff * skb, int sz, void *
here);
```

Arguments

skb

buffer

sz

size

here

address

Description

Out of line support code for skb_push. Not user callable.

alloc_skb

Name

alloc_skb — allocate a network buffer

Synopsis

```
struct sk_buff * alloc_skb (unsigned int size, int gfp_mask);
```

Arguments

size

size to allocate

gfp_mask

allocation mask

Description

Allocate a new &sk_buff. The returned buffer has no headroom and a tail room of size bytes. The object has a reference count of one. The return is the buffer. On a failure the return is NULL.

Buffers may only be allocated from interrupts using a *gfp_mask* of GFP_ATOMIC.

__kfree_skb

Name

__kfree_skb — private function

Synopsis

```
void __kfree_skb (struct sk_buff * skb);
```

Arguments

skb

buffer

Description

Free an sk_buff. Release anything attached to the buffer. Clean the state. This is an internal helper function. Users should always call kfree_skb

skb_clone

Name

skb_clone — duplicate an sk_buff

Synopsis

```
struct sk_buff * skb_clone (struct sk_buff * skb, int gfp_mask);
```

Arguments

skb

buffer to clone

gfp_mask

allocation priority

Description

Duplicate an &sk_buff. The new one is not owned by a socket. Both copies share the same packet data but not structure. The new buffer has a reference count of 1. If the allocation fails the function returns `NULL` otherwise the new buffer is returned.

If this function is called from an interrupt `gfp_mask` must be `GFP_ATOMIC`.

skb_copy

Name

`skb_copy` — create private copy of an sk_buff

Synopsis

```
struct sk_buff * skb_copy (const struct sk_buff * skb, int
gfp_mask);
```

Arguments

`skb`

buffer to copy

`gfp_mask`

allocation priority

Description

Make a copy of both an &sk_buff and its data. This is used when the caller wishes to modify the data and needs a private copy of the data to alter. Returns `NULL` on failure or the pointer to the buffer on success. The returned buffer has a reference count of 1.

As by-product this function converts non-linear &sk_buff to linear one, so that &sk_buff becomes completely private and caller is allowed to modify all the data of returned buffer. This means that this function is not recommended for use in circumstances when only header is going to be modified. Use `pskb_copy` instead.

pskb_copy

Name

`pskb_copy` — create copy of an sk_buff with private head.

Synopsis

```
struct sk_buff * pskb_copy (struct sk_buff * skb, int gfp_mask);
```

Arguments

`skb`

buffer to copy

`gfp_mask`

allocation priority

Description

Make a copy of both an &sk_buff and part of its data, located in header. Fragmented data remain shared. This is used when the caller wishes to modify only header of &sk_buff and needs private copy of the header to alter. Returns `NULL` on failure or the pointer to the buffer on success. The returned buffer has a reference count of 1.

pskb_expand_head

Name

`pskb_expand_head` — reallocate header of sk_buff

Synopsis

```
int pskb_expand_head (struct sk_buff * skb, int nhead, int
ntail, int gfp_mask);
```

Arguments

`skb`

buffer to reallocate

`nhead`

room to add at head

`ntail`

room to add at tail

gfp_mask

allocation priority

Description

Expands (or creates identical copy, if &nhead and &ntail are zero) header of skb. &sk_buff itself is not changed. &sk_buff MUST have reference count of 1. Returns zero in the case of success or error, if expansion failed. In the last case, &sk_buff is not changed.

All the pointers pointing into skb header may change and must be reloaded after call to this function.

skb_copy_expand

Name

skb_copy_expand — copy and expand sk_buff

Synopsis

```
struct sk_buff * skb_copy_expand (const struct sk_buff * skb,
int newheadroom, int newtailroom, int gfp_mask);
```

Arguments

skb

buffer to copy

newheadroom

new free bytes at head

newtailroom

new free bytes at tail

gfp_mask

allocation priority

Description

Make a copy of both an &sk_buff and its data and while doing so allocate additional space.

This is used when the caller wishes to modify the data and needs a private copy of the data to alter as well as more space for new fields. Returns `NULL` on failure or the pointer to the buffer on success. The returned buffer has a reference count of 1.

You must pass `GFP_ATOMIC` as the allocation priority if this function is called from an interrupt.

__pskb_pull_tail

Name

`__pskb_pull_tail` — advance tail of skb header

Synopsis

```
unsigned char * __pskb_pull_tail (struct sk_buff * skb, int
delta);
```

Arguments

skb

buffer to reallocate

delta

number of bytes to advance tail

Description

The function makes a sense only on a fragmented &sk_buff, it expands header moving its tail forward and copying necessary data from fragmented part.

&sk_buff MUST have reference count of 1.

Returns NULL (and &sk_buff does not change) if pull failed or value of new tail of skb in the case of success.

All the pointers pointing into skb header may change and must be reloaded after call to this function.

7.2. Socket Filter

sk_run_filter

Name

sk_run_filter — run a filter on a socket

Synopsis

```
int sk_run_filter (struct sk_buff * skb, struct sock_filter *
filter, int flen);
```

Arguments

skb

buffer to run the filter on

filter

filter to apply

flen

length of filter

Description

Decode and apply filter instructions to the skb->data. Return length to keep, 0 for none. skb is the data we are filtering, filter is the array of filter instructions, and len is the number of filter blocks in the array.

sk_chk_filter

Name

sk_chk_filter — verify socket filter code

Synopsis

```
int sk_chk_filter (struct sock_filter * filter, int flen);
```

Arguments

filter

filter to verify

flen

length of filter

Description

Check the user’s filter code. If we let some ugly filter code slip through kaboom! The filter must contain no references or jumps that are out of range, no illegal instructions and no backward jumps. It must end with a RET instruction

Returns 0 if the rule set is legal or a negative errno code if not.

Chapter 8. Network device support

8.1. Driver Support

init_etherdev

Name

`init_etherdev` — Register ethernet device

Synopsis

```
struct net_device * init_etherdev (struct net_device * dev, int
sizeof_priv);
```

Arguments

`dev`

An ethernet device structure to be filled in, or `NULL` if a new struct should be allocated.

`sizeof_priv`

Size of additional driver-private structure to be allocated for this ethernet device

Description

Fill in the fields of the device structure with ethernet-generic values.

If no device structure is passed, a new one is constructed, complete with a private data area of size *`sizeof_priv`*. A 32-byte (not bit) alignment is enforced for this private data area.

If an empty string area is passed as dev->name, or a new structure is made, a new name string is constructed.

alloc_etherdev

Name

`alloc_etherdev` — Allocates and sets up an ethernet device

Synopsis

```
struct net_device * alloc_etherdev (int sizeof_priv);
```

Arguments

`sizeof_priv`

Size of additional driver-private structure to be allocated for this ethernet device

Description

Fill in the fields of the device structure with ethernet-generic values. Basically does everything except registering the device.

Constructs a new net device, complete with a private data area of size *`sizeof_priv`*. A 32-byte (not bit) alignment is enforced for this private data area.

init_fddidev

Name

`init_fddidev` — Register FDDI device

Synopsis

```
struct net_device * init_fddidev (struct net_device * dev, int
sizeof_priv);
```

Arguments

`dev`

A FDDI device structure to be filled in, or `NULL` if a new struct should be allocated.

`sizeof_priv`

Size of additional driver-private structure to be allocated for this ethernet device

Description

Fill in the fields of the device structure with FDDI-generic values.

If no device structure is passed, a new one is constructed, complete with a private data area of size *`sizeof_priv`*. A 32-byte (not bit) alignment is enforced for this private data area.

If an empty string area is passed as dev->name, or a new structure is made, a new name string is constructed.

alloc_fddidev

Name

`alloc_fddidev` — Register FDDI device

Synopsis

```
struct net_device * alloc_fddidev (int sizeof_priv);
```

Arguments

`sizeof_priv`

Size of additional driver-private structure to be allocated for this FDDI device

Description

Fill in the fields of the device structure with FDDI-generic values.

Constructs a new net device, complete with a private data area of size *`sizeof_priv`*. A 32-byte (not bit) alignment is enforced for this private data area.

init_hippi_dev

Name

`init_hippi_dev` — Register HIPPI device

Synopsis

```
struct net_device * init_hippi_dev (struct net_device * dev, int
sizeof_priv);
```

Arguments

`dev`

A HIPPI device structure to be filled in, or `NULL` if a new struct should be allocated.

`sizeof_priv`

Size of additional driver-private structure to be allocated for this ethernet device

Description

Fill in the fields of the device structure with HIPPI-generic values.

If no device structure is passed, a new one is constructed, complete with a private data area of size *`sizeof_priv`*. A 32-byte (not bit) alignment is enforced for this private data area.

If an empty string area is passed as dev->name, or a new structure is made, a new name string is constructed.

alloc_hippi_dev

Name

`alloc_hippi_dev` — Register HIPPI device

Synopsis

```
struct net_device * alloc_hippi_dev (int sizeof_priv);
```

Arguments

`sizeof_priv`

Size of additional driver-private structure to be allocated for this HIPPI device

Description

Fill in the fields of the device structure with HIPPI-generic values.

Constructs a new net device, complete with a private data area of size *`sizeof_priv`*. A 32-byte (not bit) alignment is enforced for this private data area.

init_trdev

Name

`init_trdev` — Register token ring device

Synopsis

```
struct net_device * init_trdev (struct net_device * dev, int
sizeof_priv);
```

Arguments

dev

A token ring device structure to be filled in, or NULL if a new struct should be allocated.

sizeof_priv

Size of additional driver-private structure to be allocated for this ethernet device

Description

Fill in the fields of the device structure with token ring-generic values.

If no device structure is passed, a new one is constructed, complete with a private data area of size *sizeof_priv*. A 32-byte (not bit) alignment is enforced for this private data area.

If an empty string area is passed as dev->name, or a new structure is made, a new name string is constructed.

alloc_trdev

Name

alloc_trdev — Register token ring device

Synopsis

```
struct net_device * alloc_trdev (int sizeof_priv);
```

Arguments

sizeof_priv

Size of additional driver-private structure to be allocated for this token ring device

Description

Fill in the fields of the device structure with token ring-generic values.

Constructs a new net device, complete with a private data area of size *sizeof_priv*. A 32-byte (not bit) alignment is enforced for this private data area.

init_fcdev

Name

init_fcdev — Register fibre channel device

Synopsis

```
struct net_device * init_fcdev (struct net_device * dev, int
sizeof_priv);
```

Arguments

dev

A fibre channel device structure to be filled in, or NULL if a new struct should be allocated.

sizeof_priv

Size of additional driver-private structure to be allocated for this ethernet device

Description

Fill in the fields of the device structure with fibre channel-generic values.

If no device structure is passed, a new one is constructed, complete with a private data area of size *sizeof_priv*. A 32-byte (not bit) alignment is enforced for this private data area.

If an empty string area is passed as dev->name, or a new structure is made, a new name string is constructed.

alloc_fcdev

Name

alloc_fcdev — Register fibre channel device

Synopsis

```
struct net_device * alloc_fcdev (int sizeof_priv);
```

Arguments

sizeof_priv

Size of additional driver-private structure to be allocated for this fibre channel device

Description

Fill in the fields of the device structure with fibre channel-generic values.

Constructs a new net device, complete with a private data area of size *sizeof_priv*. A 32-byte (not bit) alignment is enforced for this private data area.

dev_add_pack

Name

dev_add_pack — add packet handler

Synopsis

```
void dev_add_pack (struct packet_type * pt);
```

Arguments

pt

packet type declaration

Description

Add a protocol handler to the networking stack. The passed &packet_type is linked into kernel lists and may not be freed until it has been removed from the kernel lists.

dev_remove_pack

Name

`dev_remove_pack` — remove packet handler

Synopsis

```
void dev_remove_pack (struct packet_type * pt);
```

Arguments

`pt`

packet type declaration

Description

Remove a protocol handler that was previously added to the kernel protocol handlers by `dev_add_pack`. The passed &packet_type is removed from the kernel lists and can be freed or reused once this function returns.

__dev_get_by_name

Name

__dev_get_by_name — find a device by its name

Synopsis

```
struct net_device * __dev_get_by_name (const char * name);
```

Arguments

name

name to find

Description

Find an interface by name. Must be called under RTNL semaphore or *dev_base_lock*. If the name is found a pointer to the device is returned. If the name is not found then NULL is returned. The reference counters are not incremented so the caller must be careful with locks.

dev_get_by_name

Name

dev_get_by_name — find a device by its name

Synopsis

```
struct net_device * dev_get_by_name (const char * name);
```

Arguments

name

name to find

Description

Find an interface by name. This can be called from any context and does its own locking. The returned handle has the usage count incremented and the caller must use `dev_put` to release it when it is no longer needed. `NULL` is returned if no matching device is found.

dev_get

Name

`dev_get` — test if a device exists

Synopsis

```
int dev_get (const char * name);
```

Arguments

name

name to test for

Description

Test if a name exists. Returns true if the name is found. In order to be sure the name is not allocated or removed during the test the caller must hold the rtnl semaphore.

This function primarily exists for back compatibility with older drivers.

__dev_get_by_index

Name

__dev_get_by_index — find a device by its ifindex

Synopsis

```
struct net_device * __dev_get_by_index (int ifindex);
```

Arguments

ifindex

index of device

Description

Search for an interface by index. Returns `NULL` if the device is not found or a pointer to the device. The device has not had its reference counter increased so the caller must be careful about locking. The caller must hold either the RTNL semaphore or *`dev_base_lock`*.

dev_get_by_index

Name

`dev_get_by_index` — find a device by its ifindex

Synopsis

```
struct net_device * dev_get_by_index (int ifindex);
```

Arguments

`ifindex`

index of device

Description

Search for an interface by index. Returns NULL if the device is not found or a pointer to the device. The device returned has had a reference added and the pointer is safe until the user calls dev_put to indicate they have finished with it.

dev_alloc_name

Name

dev_alloc_name — allocate a name for a device

Synopsis

```
int dev_alloc_name (struct net_device * dev, const char * name);
```

Arguments

dev

device

name

name format string

Description

Passed a format string - eg "ltd" it will try and find a suitable id. Not efficient for many devices, not called a lot. The caller must hold the dev_base or rtnl lock while allocating the name and adding the device in order to avoid duplicates. Returns the number of the unit assigned or a negative errno code.

dev_alloc

Name

dev_alloc — allocate a network device and name

Synopsis

```
struct net_device * dev_alloc (const char * name, int * err);
```

Arguments

name

name format string

err

error return pointer

Description

Passed a format string, eg. "ltd", it will allocate a network device and space for the name. NULL is returned if no memory is available. If the allocation succeeds then the name is assigned and the device pointer returned. NULL is returned if the name allocation failed. The cause of an error is returned as a negative errno code in the variable *err* points to.

The caller must hold the *dev_base* or RTNL locks when doing this in order to avoid duplicate name allocations.

netdev_state_change

Name

`netdev_state_change` — device changes state

Synopsis

```
void netdev_state_change (struct net_device * dev);
```

Arguments

dev

device to cause notification

Description

Called to indicate a device has changed state. This function calls the notifier chains for netdev_chain and sends a NEWLINK message to the routing socket.

dev_load

Name

`dev_load` — load a network module

Synopsis

```
void dev_load (const char * name);
```

Arguments

name

name of interface

Description

If a network interface is not present and the process has suitable privileges this function loads the module. If module loading is not available in this kernel then it becomes a nop.

dev_open

Name

dev_open — prepare an interface for use.

Synopsis

```
int dev_open (struct net_device * dev);
```

Arguments

dev

device to open

Description

Takes a device from down to up state. The device's private open function is invoked and then the multicast lists are loaded. Finally the device is moved into the up state and a NETDEV_UP message is sent to the netdev notifier chain.

Calling this function on an active interface is a nop. On a failure a negative errno code is returned.

dev_close

Name

dev_close — shutdown an interface.

Synopsis

```
int dev_close (struct net_device * dev);
```

Arguments

dev

device to shutdown

Description

This function moves an active device into down state. A `NETDEV_GOING_DOWN` is sent to the netdev notifier chain. The device is then deactivated and finally a `NETDEV_DOWN` is sent to the notifier chain.

register_netdevice_notifier

Name

`register_netdevice_notifier` — register a network notifier block

Synopsis

```
int register_netdevice_notifier (struct notifier_block * nb);
```

Arguments

nb

notifier

Description

Register a notifier to be called when network device events occur. The notifier passed is linked into the kernel structures and must not be reused until it has been unregistered. A negative errno code is returned on a failure.

unregister_netdevice_notifier

Name

`unregister_netdevice_notifier` — unregister a network notifier block

Synopsis

```
int unregister_netdevice_notifier (struct notifier_block * nb);
```

Arguments

nb

notifier

Description

Unregister a notifier previously registered by `register_netdevice_notifier`. The notifier is unlinked into the kernel structures and may then be reused. A negative errno code is returned on a failure.

dev_queue_xmit

Name

`dev_queue_xmit` — transmit a buffer

Synopsis

```
int dev_queue_xmit (struct sk_buff * skb);
```

Arguments

skb

buffer to transmit

Description

Queue a buffer for transmission to a network device. The caller must have set the device and priority and built the buffer before calling this function. The function can be called from an interrupt.

A negative errno code is returned on a failure. A success does not guarantee the frame will be transmitted as it may be dropped due to congestion or traffic shaping.

netif_rx

Name

`netif_rx` — post buffer to the network code

Synopsis

```
int netif_rx (struct sk_buff * skb);
```

Arguments

skb

buffer to post

Description

This function receives a packet from a device driver and queues it for the upper (protocol) levels to process. It always succeeds. The buffer may be dropped during processing for congestion control or by the protocol layers.

return values

NET_RX_SUCCESS (no congestion) NET_RX_CN_LOW (low congestion) NET_RX_CN_MOD (moderate congestion) NET_RX_CN_HIGH (high congestion) NET_RX_DROP (packet was dropped)

net_call_rx_atomic

Name

net_call_rx_atomic —

Synopsis

```
void net_call_rx_atomic (void (*fn) (void));
```

Arguments

fn

function to call

Description

Make a function call that is atomic with respect to the protocol layers.

register_gifconf

Name

register_gifconf — register a SIOCGIF handler

Synopsis

```
int register_gifconf (unsigned int family, gifconf_func_t *
gifconf);
```

Arguments

family

Address family

gifconf

Function handler

Description

Register protocol dependent address dumping routines. The handler that is passed must not be freed or reused until it has been replaced by another handler.

netdev_set_master

Name

`netdev_set_master` — set up master/slave pair

Synopsis

```
int netdev_set_master (struct net_device * slave, struct
net_device * master);
```

Arguments

`slave`

slave device

`master`

new master device

Description

Changes the master device of the slave. Pass `NULL` to break the bonding. The caller must hold the RTNL semaphore. On a failure a negative errno code is returned. On

success the reference counts are adjusted, RTM_NEWLINK is sent to the routing socket and the function returns zero.

dev_set_promiscuity

Name

dev_set_promiscuity — update promiscuity count on a device

Synopsis

```
void dev_set_promiscuity (struct net_device * dev, int inc);
```

Arguments

dev

device

inc

modifier

Description

Add or remove promsicuity from a device. While the count in the device remains above zero the interface remains promiscuous. Once it hits zero the device reverts back to normal filtering operation. A negative inc value is used to drop promiscuity on the device.

dev_set_allmulti

Name

`dev_set_allmulti` — update allmulti count on a device

Synopsis

```
void dev_set_allmulti (struct net_device * dev, int inc);
```

Arguments

`dev`

device

`inc`

modifier

Description

Add or remove reception of all multicast frames to a device. While the count in the device remains above zero the interface remains listening to all interfaces. Once it hits zero the device reverts back to normal filtering operation. A negative *`inc`* value is used to drop the counter when releasing a resource needing all multicasts.

dev_ioctl

Name

`dev_ioctl` — network device ioctl

Synopsis

```
int dev_ioctl (unsigned int cmd, void * arg);
```

Arguments

`cmd`

command to issue

`arg`

pointer to a struct ifreq in user space

Description

Issue ioctl functions to devices. This is normally called by the user space syscall interfaces but can sometimes be useful for other purposes. The return value is the return from the syscall if positive or a negative errno code on error.

dev_new_index

Name

dev_new_index — allocate an ifindex

Synopsis

```
int dev_new_index ( void);
```

Arguments

void

no arguments

Description

Returns a suitable unique value for a new device interface number. The caller must hold the rtnl semaphore or the dev_base_lock to be sure it remains unique.

netdev_finish_unregister

Name

netdev_finish_unregister — complete unregistration

Synopsis

```
int netdev_finish_unregister (struct net_device * dev);
```

Arguments

dev

device

Description

Destroy and free a dead device. A value of zero is returned on success.

unregister_netdevice

Name

unregister_netdevice — remove device from the kernel

Synopsis

```
int unregister_netdevice (struct net_device * dev);
```

Arguments

dev

device

Description

This function shuts down a device interface and removes it from the kernel tables. On success 0 is returned, on a failure a negative errno code is returned.

Callers must hold the rtnl semaphore. See the comment at the end of Space.c for details about the locking. You may want `unregister_netdev` instead of this.

8.2. 8390 Based Network Cards

ei_open

Name

`ei_open` — Open/initialize the board.

Synopsis

```
int ei_open (struct net_device * dev);
```

Arguments

dev

network device to initialize

Description

This routine goes all-out, setting everything up anew at each open, even though many of these registers should only need to be set once at boot.

ei_close

Name

`ei_close` — shut down network device

Synopsis

```
int ei_close (struct net_device * dev);
```

Arguments

dev

network device to close

Description

Opposite of `ei_open`. Only used when “ifconfig <devname> down” is done.

ei_tx_timeout

Name

`ei_tx_timeout` — handle transmit time out condition

Synopsis

```
void ei_tx_timeout (struct net_device * dev);
```

Arguments

dev

network device which has apparently fallen asleep

Description

Called by kernel when device never acknowledges a transmit has completed (or failed) - i.e. never posted a Tx related interrupt.

ei_interrupt

Name

`ei_interrupt` — handle the interrupts from an 8390

Synopsis

```
void ei_interrupt (int irq, void * dev_id, struct pt_regs *
regs);
```

Arguments

`irq`

interrupt number

`dev_id`

a pointer to the net_device

`regs`

unused

Description

Handle the ether interface interrupts. We pull packets from the 8390 via the card specific functions and fire them at the networking stack. We also handle transmit completions and wake the transmit path if neccessary. We also update the counters and do other housekeeping as needed.

ethdev_init

Name

`ethdev_init` — init rest of 8390 device struct

Synopsis

```
int ethdev_init (struct net_device * dev);
```

Arguments

dev

network device structure to init

Description

Initialize the rest of the 8390 device structure. Do NOT __init this, as it is used by 8390 based modular drivers too.

NS8390_init

Name

`NS8390_init` — initialize 8390 hardware

Synopsis

```
void NS8390_init (struct net_device * dev, int startp);
```

Arguments

dev

network device to initialize

startp

boolean. non-zero value to initiate chip processing

Description

Must be called with lock held.

8.3. Synchronous PPP

sppp_input

Name

`sppp_input` — receive and process a WAN PPP frame

Synopsis

```
void sppp_input (struct net_device * dev, struct sk_buff * skb);
```

Arguments

dev

The device it arrived on

skb

The buffer to process

Description

This can be called directly by cards that do not have timing constraints but is normally called from the network layer after interrupt servicing to process frames queued via `netif_rx`.

We process the options in the card. If the frame is destined for the protocol stacks then it requeues the frame for the upper level protocol. If it is a control from it is processed and discarded here.

sppp_close

Name

`sppp_close` — close down a synchronous PPP or Cisco HDLC link

Synopsis

```
int sppp_close (struct net_device * dev);
```

Arguments

dev

The network device to drop the link of

Description

This drops the logical interface to the channel. It is not done politely as we assume we will also be dropping DTR. Any timeouts are killed.

sppp_open

Name

sppp_open — open a synchronous PPP or Cisco HDLC link

Synopsis

```
int sppp_open (struct net_device * dev);
```

Arguments

dev

Network device to activate

Description

Close down any existing synchronous session and commence from scratch. In the PPP case this means negotiating LCP/IPCP and friends, while for Cisco HDLC we simply need to start sending keepalives

sppp_reopen

Name

sppp_reopen — notify of physical link loss

Synopsis

```
int sppp_reopen (struct net_device * dev);
```

Arguments

dev

Device that lost the link

Description

This function informs the synchronous protocol code that the underlying link died (for example a carrier drop on X.21)

We increment the magic numbers to ensure that if the other end failed to notice we will correctly start a new session. It happens do to the nature of telco circuits is that you can lose carrier on one endonly.

Having done this we go back to negotiating. This function may be called from an interrupt context.

sppp_change_mtu

Name

`sppp_change_mtu` — Change the link MTU

Synopsis

```
int sppp_change_mtu (struct net_device * dev, int new_mtu);
```

Arguments

`dev`

Device to change MTU on

`new_mtu`

New MTU

Description

Change the MTU on the link. This can only be called with the link down. It returns an error if the link is up or the mtu is out of range.

sppp_do_ioctl

Name

`sppp_do_ioctl` — Ioctl handler for ppp/hdlc

Synopsis

```
int sppp_do_ioctl (struct net_device * dev, struct ifreq * ifr,
int cmd);
```

Arguments

`dev`

Device subject to ioctl

`ifr`

Interface request block from the user

`cmd`

Command that is being issued

Description

This function handles the ioctls that may be issued by the user to control the settings of a PPP/HDLC link. It does both busy and security checks. This function is intended to be wrapped by callers who wish to add additional ioctl calls of their own.

sppp_attach

Name

`sppp_attach` — attach synchronous PPP/HDLC to a device

Synopsis

```
void sppp_attach (struct ppp_device * pd);
```

Arguments

pd

PPP device to initialise

Description

This initialises the PPP/HDLC support on an interface. At the time of calling the dev element must point to the network device that this interface is attached to. The interface should not yet be registered.

sppp_detach

Name

sppp_detach — release PPP resources from a device

Synopsis

```
void sppp_detach (struct net_device * dev);
```

Arguments

dev

Network device to release

Description

Stop and free up any PPP/HDLC resources used by this interface. This must be called before the device is freed.

Chapter 9. Module Support

9.1. Module Loading

request_module

Name

`request_module` — try to load a kernel module

Synopsis

```
int request_module (const char * module_name);
```

Arguments

`module_name`

Name of module

Description

Load a module using the user mode module loader. The function returns zero on success or a negative errno code on failure. Note that a successful module load does not mean the module did not then unload and exit on an error of its own. Callers must check that the service they requested is now available not blindly invoke it.

If module auto-loading support is disabled then this function becomes a no-operation.

call_usermodehelper

Name

`call_usermodehelper` — start a usermode application

Synopsis

```
int call_usermodehelper (char * path, char ** argv, char **
envp);
```

Arguments

`path`

pathname for the application

`argv`

null-terminated argument list

`envp`

null-terminated environment list

Description

Runs a user-space application. The application is started asynchronously. It runs as a child of keventd. It runs with full root capabilities. keventd silently reaps the child when it exits.

Must be called from process context. Returns zero on success, else a negative error code.

9.2. Inter Module support

inter_module_register

Name

`inter_module_register` — register a new set of inter module data.

Synopsis

```
void inter_module_register (const char * im_name, struct module
* owner, const void * userdata);
```

Arguments

`im_name`

an arbitrary string to identify the data, must be unique

`owner`

module that is registering the data, always use THIS_MODULE

`userdata`

pointer to arbitrary userdata to be registered

Description

Check that the im_name has not already been registered, complain if it has. For new data, add it to the inter_module_entry list.

inter_module_unregister

Name

`inter_module_unregister` — unregister a set of inter module data.

Synopsis

```
void inter_module_unregister (const char * im_name);
```

Arguments

`im_name`

an arbitrary string to identify the data, must be unique

Description

Check that the im_name has been registered, complain if it has not. For existing data, remove it from the inter_module_entry list.

inter_module_get

Name

`inter_module_get` — return arbitrary userdata from another module.

Synopsis

```
const void * inter_module_get (const char * im_name);
```

Arguments

im_name

an arbitrary string to identify the data, must be unique

Description

If the im_name has not been registered, return NULL. Try to increment the use count on the owning module, if that fails then return NULL. Otherwise return the userdata.

inter_module_get_request

Name

inter_module_get_request — im get with automatic request_module.

Synopsis

```
const void * inter_module_get_request (const char * im_name,
const char * modname);
```

Arguments

im_name

an arbitrary string to identify the data, must be unique

modname

module that is expected to register im_name

Description

If inter_module_get fails, do request_module then retry.

inter_module_put

Name

`inter_module_put` — release use of data from another module.

Synopsis

```
void inter_module_put (const char * im_name);
```

Arguments

im_name

an arbitrary string to identify the data, must be unique

Description

If the im_name has not been registered, complain, otherwise decrement the use count on the owning module.

Chapter 10. Hardware Interfaces

10.1. Interrupt Handling

disable_irq_nosync

Name

`disable_irq_nosync` — disable an irq without waiting

Synopsis

```
void disable_irq_nosync (unsigned int irq);
```

Arguments

`irq`

Interrupt to disable

Description

Disable the selected interrupt line. Disables and Enables are nested. Unlike `disable_irq`, this function does not ensure existing instances of the IRQ handler have completed before returning.

This function may be called from IRQ context.

disable_irq

Name

`disable_irq` — disable an irq and wait for completion

Synopsis

```
void disable_irq (unsigned int irq);
```

Arguments

`irq`

Interrupt to disable

Description

Disable the selected interrupt line. Enables and Disables are nested. This function waits for any pending IRQ handlers for this interrupt to complete before returning. If you use this function while holding a resource the IRQ handler may need you will deadlock.

This function may be called - with care - from IRQ context.

enable_irq

Name

`enable_irq` — enable handling of an irq

Synopsis

```
void enable_irq (unsigned int irq);
```

Arguments

irq

Interrupt to enable

Description

Undoes the effect of one call to `disable_irq`. If this matches the last disable, processing of interrupts on this IRQ line is re-enabled.

This function may be called from IRQ context.

probe_irq_mask

Name

`probe_irq_mask` — scan a bitmap of interrupt lines

Synopsis

```
unsigned int probe_irq_mask (unsigned long val);
```

Arguments

val

mask of interrupts to consider

Description

Scan the ISA bus interrupt lines and return a bitmap of active interrupts. The interrupt probe logic state is then returned to its previous value.

Note

we need to scan all the irq's even though we will only return ISA irq numbers - just so that we reset them all to a known state.

10.2. MTRR Handling

mtrr_add

Name

mtrr_add — Add a memory type region

Synopsis

```
int mtrr_add (unsigned long base, unsigned long size, unsigned
int type, char increment);
```

Arguments

base

Physical base address of region

size

Physical size of region

type

Type of MTRR desired

increment

If this is true do usage counting on the region

Description

Memory type region registers control the caching on newer Intel and non Intel processors. This function allows drivers to request an MTRR is added. The details and hardware specifics of each processor's implementation are hidden from the caller, but nevertheless the caller should expect to need to provide a power of two size on an equivalent power of two boundary.

If the region cannot be added either because all regions are in use or the CPU cannot support it a negative value is returned. On success the register number for this entry is returned, but should be treated as a cookie only.

On a multiprocessor machine the changes are made to all processors. This is required on x86 by the Intel processors.

The available types are

`MTRR_TYPE_UNCACHABLE` - No caching

`MTRR_TYPE_WRBACK` - Write data back in bursts whenever

`MTRR_TYPE_WRCOMB` - Write data back soon but allow bursts

`MTRR_TYPE_WRTHROUGH` - Cache reads but not writes

BUGS

Needs a quiet flag for the cases where drivers do not mind failures and do not wish system log messages to be sent.

mtrr_del

Name

`mtrr_del` — delete a memory type region

Synopsis

```
int mtrr_del (int reg, unsigned long base, unsigned long size);
```

Arguments

`reg`

Register returned by mtrr_add

`base`

Physical base address

`size`

Size of region

Description

If register is supplied then base and size are ignored. This is how drivers should call it.

Releases an MTRR region. If the usage count drops to zero the register is freed and the region returns to default state. On success the register is returned, on failure a negative error code.

10.3. PCI Support Library

pci_find_slot

Name

`pci_find_slot` — locate PCI device from a given PCI slot

Synopsis

```
struct pci_dev * pci_find_slot (unsigned int bus, unsigned int devfn);
```

Arguments

`bus`

number of PCI bus on which desired PCI device resides

`devfn`

encodes number of PCI slot in which the desired PCI device resides and the logical device number within that slot in case of multi-function devices.

Description

Given a PCI bus and slot/function number, the desired PCI device is located in system global list of PCI devices. If the device is found, a pointer to its data structure is returned. If no device is found, `NULL` is returned.

pci_find_subsys

Name

`pci_find_subsys` — begin or continue searching for a PCI device by vendor/subvendor/device/subdevice id

Synopsis

```
struct pci_dev * pci_find_subsys (unsigned int vendor, unsigned
int device, unsigned int ss_vendor, unsigned int ss_device,
const struct pci_dev * from);
```

Arguments

`vendor`

PCI vendor id to match, or `PCI_ANY_ID` to match all vendor ids

`device`

PCI device id to match, or `PCI_ANY_ID` to match all device ids

`ss_vendor`

PCI subsystem vendor id to match, or `PCI_ANY_ID` to match all vendor ids

ss_device

PCI subsystem device id to match, or PCI_ANY_ID to match all device ids

from

Previous PCI device found in search, or NULL for new search.

Description

Iterates through the list of known PCI devices. If a PCI device is found with a matching *vendor*, *device*, *ss_vendor* and *ss_device*, a pointer to its device structure is returned. Otherwise, NULL is returned. A new search is initiated by passing NULL to the *from* argument. Otherwise if *from* is not NULL, searches continue from next device on the global list.

pci_find_device

Name

pci_find_device — begin or continue searching for a PCI device by vendor/device id

Synopsis

```
struct pci_dev * pci_find_device (unsigned int vendor, unsigned
int device, const struct pci_dev * from);
```

Arguments

vendor

PCI vendor id to match, or PCI_ANY_ID to match all vendor ids

device

PCI device id to match, or PCI_ANY_ID to match all device ids

from

Previous PCI device found in search, or NULL for new search.

Description

Iterates through the list of known PCI devices. If a PCI device is found with a matching *vendor* and *device*, a pointer to its device structure is returned. Otherwise, NULL is returned. A new search is initiated by passing NULL to the *from* argument. Otherwise if *from* is not NULL, searches continue from next device on the global list.

pci_find_class

Name

pci_find_class — begin or continue searching for a PCI device by class

Synopsis

```
struct pci_dev * pci_find_class (unsigned int class, const
struct pci_dev * from);
```

Arguments

class

search for a PCI device with this class designation

from

Previous PCI device found in search, or NULL for new search.

Description

Iterates through the list of known PCI devices. If a PCI device is found with a matching *class*, a pointer to its device structure is returned. Otherwise, NULL is returned. A new search is initiated by passing NULL to the *from* argument. Otherwise if *from* is not NULL, searches continue from next device on the global list.

pci_find_capability

Name

pci_find_capability — query for devices' capabilities

Synopsis

```
int pci_find_capability (struct pci_dev * dev, int cap);
```

Arguments

dev

PCI device to query

cap

capability code

Description

Tell if a device supports a given PCI capability. Returns the address of the requested capability structure within the device's PCI configuration space or 0 in case the device does not support it. Possible values for *cap*:

`PCI_CAP_ID_PM` Power Management

`PCI_CAP_ID_AGP` Accelerated Graphics Port

`PCI_CAP_ID_VPD` Vital Product Data

`PCI_CAP_ID_SLOTID` Slot Identification

`PCI_CAP_ID_MSI` Message Signalled Interrupts

`PCI_CAP_ID_CHSWP` CompactPCI HotSwap

pci_find_parent_resource

Name

`pci_find_parent_resource` — return resource region of parent bus of given region

Synopsis

```
struct resource * pci_find_parent_resource (const struct pci_dev
* dev, struct resource * res);
```

Arguments

dev

PCI device structure contains resources to be searched

res

child resource record for which parent is sought

Description

For given resource region of given device, return the resource region of parent bus the given region is contained in or where it should be allocated from.

pci_set_power_state

Name

pci_set_power_state — Set the power state of a PCI device

Synopsis

```
int pci_set_power_state (struct pci_dev * dev, int state);
```

Arguments

dev

PCI device to be suspended

state

Power state we're entering

Description

Transition a device to a new power state, using the Power Management Capabilities in the device's config space.

RETURN VALUE

-EINVAL if trying to enter a lower state than we're already in. 0 if we're already in the requested state. -EIO if device does not support PCI PM. 0 if we can successfully change the power state.

pci_save_state

Name

pci_save_state — save the PCI configuration space of a device before suspending

Synopsis

```
int pci_save_state (struct pci_dev * dev, u32 * buffer);
```

Arguments

dev

- PCI device that we’re dealing with

buffer

- buffer to hold config space context

Description

buffer must be large enough to hold the entire PCI 2.2 config space (>= 64 bytes).

pci_restore_state

Name

pci_restore_state — Restore the saved state of a PCI device

Synopsis

```
int pci_restore_state (struct pci_dev * dev, u32 * buffer);
```

Arguments

dev

- PCI device that we’re dealing with

buffer

- saved PCI config space

pci_enable_device

Name

pci_enable_device — Initialize device before it's used by a driver.

Synopsis

```
int pci_enable_device (struct pci_dev * dev);
```

Arguments

dev

PCI device to be initialized

Description

Initialize device before it's used by a driver. Ask low-level code to enable I/O and memory. Wake up the device if it was suspended. Beware, this function can fail.

pci_disable_device

Name

pci_disable_device — Disable PCI device after use

Synopsis

```
void pci_disable_device (struct pci_dev * dev);
```

Arguments

dev

PCI device to be disabled

Description

Signal to the system that the PCI device is not in use by the system anymore. This only involves disabling PCI bus-mastering, if active.

pci_enable_wake

Name

pci_enable_wake — enable device to generate PME# when suspended

Synopsis

```
int pci_enable_wake (struct pci_dev * dev, u32 state, int
enable);
```

Arguments

dev

- PCI device to operate on

state

- Current state of device.

enable

- Flag to enable or disable generation

Description

Set the bits in the device's PM Capabilities to generate PME# when the system is suspended.

-EIO is returned if device doesn't have PM Capabilities. -EINVAL is returned if device supports it, but can't generate wake events. 0 if operation is successful.

pci_release_regions

Name

pci_release_regions — Release reserved PCI I/O and memory resources

Synopsis

```
void pci_release_regions (struct pci_dev * pdev);
```

Arguments

pdev

PCI device whose resources were previously reserved by pci_request_regions

Description

Releases all PCI I/O and memory resources previously reserved by a successful call to pci_request_regions. Call this function only after all use of the PCI regions has ceased.

pci_request_regions

Name

pci_request_regions — Reserved PCI I/O and memory resources

Synopsis

```
int pci_request_regions (struct pci_dev * pdev, char *
res_name);
```

Arguments

pdev

PCI device whose resources are to be reserved

res_name

Name to be associated with resource.

Description

Mark all PCI regions associated with PCI device *pdev* as being reserved by owner *res_name*. Do not access any address inside the PCI regions unless this call returns successfully.

Returns 0 on success, or EBUSY on error. A warning message is also printed on failure.

pci_match_device

Name

pci_match_device — Tell if a PCI device structure has a matching PCI device id structure

Synopsis

```
const struct pci_device_id * pci_match_device (const struct
pci_device_id * ids, const struct pci_dev * dev);
```

Arguments

ids

array of PCI device id structures to search in

dev

the PCI device structure to match against

Description

Used by a driver to check whether a PCI device present in the system is in its list of supported devices.Returns the matching pci_device_id structure or NULL if there is no match.

pci_register_driver

Name

pci_register_driver — register a new pci driver

Synopsis

```
int pci_register_driver (struct pci_driver * drv);
```

Arguments

drv

the driver structure to register

Description

Adds the driver structure to the list of registered drivers Returns the number of pci devices which were claimed by the driver during registration. The driver remains registered even if the return value is zero.

pci_unregister_driver

Name

`pci_unregister_driver` — unregister a pci driver

Synopsis

```
void pci_unregister_driver (struct pci_driver * drv);
```

Arguments

`drv`

the driver structure to unregister

Description

Deletes the driver structure from the list of registered PCI drivers, gives it a chance to clean up by calling its `remove` function for each device it was responsible for, and marks those devices as driverless.

pci_insert_device

Name

`pci_insert_device` — insert a hotplug device

Synopsis

```
void pci_insert_device (struct pci_dev * dev, struct pci_bus *
bus);
```

Arguments

dev

the device to insert

bus

where to insert it

Description

Add a new device to the device lists and notify userspace (/sbin/hotplug).

pci_remove_device

Name

`pci_remove_device` — remove a hotplug device

Synopsis

```
void pci_remove_device (struct pci_dev * dev);
```

Arguments

dev

the device to remove

Description

Delete the device structure from the device lists and notify userspace (/sbin/hotplug).

pci_dev_driver

Name

`pci_dev_driver` — get the pci_driver of a device

Synopsis

```
struct pci_driver * pci_dev_driver (const struct pci_dev * dev);
```

Arguments

dev

the device to query

Description

Returns the appropriate pci_driver structure or `NULL` if there is no registered driver for the device.

pci_set_master

Name

`pci_set_master` — enables bus-mastering for device dev

Synopsis

```
void pci_set_master (struct pci_dev * dev);
```

Arguments

dev

the PCI device to enable

Description

Enables bus-mastering on the device and calls `pcibios_set_master` to do the needed arch specific settings.

pci_setup_device

Name

`pci_setup_device` — fill in class and map information of a device

Synopsis

```
int pci_setup_device (struct pci_dev * dev);
```

Arguments

`dev`

the device structure to fill

Description

Initialize the device structure with information about the device's vendor,class,memory and IO-space addresses,IRQ lines etc. Called at initialisation of the PCI subsystem and by CardBus services. Returns 0 on success and -1 if unknown type of device (not normal, bridge or CardBus).

pci_pool_create

Name

`pci_pool_create` — Creates a pool of pci consistent memory blocks, for dma.

Synopsis

```
struct pci_pool * pci_pool_create (const char * name, struct
pci_dev * pdev, size_t size, size_t align, size_t allocation,
int flags);
```

Arguments

`name`

name of pool, for diagnostics

`pdev`

pci device that will be doing the DMA

`size`

size of the blocks in this pool.

`align`

alignment requirement for blocks; must be a power of two

`allocation`

returned blocks won’t cross this boundary (or zero)

`flags`

SLAB_* flags (not all are supported).

Description

Returns a pci allocation pool with the requested characteristics, or null if one can’t be created. Given one of these pools, `pci_pool_alloc` may be used to allocate memory. Such memory will all have “consistent” DMA mappings, accessible by the device and its driver without using cache flushing primitives. The actual size of blocks allocated may be larger than requested because of alignment.

If allocation is nonzero, objects returned from `pci_pool_alloc` won’t cross that size boundary. This is useful for devices which have addressing restrictions on individual DMA transfers, such as not crossing boundaries of 4KBytes.

pci_pool_destroy

Name

`pci_pool_destroy` — destroys a pool of pci memory blocks.

Synopsis

```
void pci_pool_destroy (struct pci_pool * pool);
```

Arguments

`pool`

pci pool that will be destroyed

Description

Caller guarantees that no more memory from the pool is in use, and that nothing will try to use the pool after this call.

pci_pool_alloc

Name

`pci_pool_alloc` — get a block of consistent memory

Synopsis

```
void * pci_pool_alloc (struct pci_pool * pool, int mem_flags,
dma_addr_t * handle);
```

Arguments

`pool`

pci pool that will produce the block

`mem_flags`

SLAB_KERNEL or SLAB_ATOMIC

`handle`

pointer to dma address of block

Description

This returns the kernel virtual address of a currently unused block, and reports its dma address through the handle. If such a memory block can’t be allocated, null is returned.

pci_pool_free

Name

`pci_pool_free` — put block back into pci pool

Synopsis

```
void pci_pool_free (struct pci_pool * pool, void * vaddr,
dma_addr_t dma);
```

Arguments

`pool`

the pci pool holding the block

`vaddr`

virtual address of block

`dma`

dma address of block

Description

Caller promises neither device nor driver will again touch this block unless it is first re-allocated.

10.4. MCA Architecture

10.4.1. MCA Device Functions

mca_find_adapter

Name

`mca_find_adapter` — scan for adapters

Synopsis

```
int mca_find_adapter (int id, int start);
```

Arguments

`id`

MCA identification to search for

`start`

starting slot

Description

Search the MCA configuration for adapters matching the 16bit ID given. The first time it should be called with start as zero and then further calls made passing the return value of the previous call until `MCA_NOTFOUND` is returned.

Disabled adapters are not reported.

mca_find_unused_adapter

Name

`mca_find_unused_adapter` — scan for unused adapters

Synopsis

```
int mca_find_unused_adapter (int id, int start);
```

Arguments

`id`

MCA identification to search for

`start`

starting slot

Description

Search the MCA configuration for adapters matching the 16bit ID given. The first time it should be called with start as zero and then further calls made passing the return

value of the previous call until `MCA_NOTFOUND` is returned.

Adapters that have been claimed by drivers and those that are disabled are not reported. This function thus allows a driver to scan for further cards when some may already be driven.

mca_read_stored_pos

Name

`mca_read_stored_pos` — read POS register from boot data

Synopsis

```
unsigned char mca_read_stored_pos (int slot, int reg);
```

Arguments

`slot`

slot number to read from

`reg`

register to read from

Description

Fetch a POS value that was stored at boot time by the kernel when it scanned the MCA space. The register value is returned. Missing or invalid registers report 0.

mca_read_pos

Name

mca_read_pos — read POS register from card

Synopsis

```
unsigned char mca_read_pos (int slot, int reg);
```

Arguments

slot

slot number to read from

reg

register to read from

Description

Fetch a POS value directly from the hardware to obtain the current value. This is much slower than mca_read_stored_pos and may not be invoked from interrupt context. It handles the deep magic required for onboard devices transparently.

mca_write_pos

Name

`mca_write_pos` — read POS register from card

Synopsis

```
void mca_write_pos (int slot, int reg, unsigned char byte);
```

Arguments

`slot`

slot number to read from

`reg`

register to read from

`byte`

byte to write to the POS registers

Description

Store a POS value directly from the hardware. You should not normally need to use this function and should have a very good knowledge of MCA bus before you do so. Doing this wrongly can damage the hardware.

This function may not be used from interrupt context.

Note that this a technically a Bad Thing, as IBM tech stuff says you should only set POS values through their utilities. However, some devices such as the 3c523 recommend that you write back some data to make sure the configuration is consistent. I’d say that IBM is right, but I like my drivers to work.

This function can’t do checks to see if multiple devices end up with the same resources, so you might see magic smoke if someone screws up.

mca_set_adapter_name

Name

`mca_set_adapter_name` — Set the description of the card

Synopsis

```
void mca_set_adapter_name (int slot, char* name);
```

Arguments

`slot`

slot to name

`name`

text string for the namen

Description

This function sets the name reported via /proc for this adapter slot. This is for user information only. Setting a name deletes any previous name.

mca_set_adapter_procfn

Name

`mca_set_adapter_procfn` — Set the /proc callback

Synopsis

```
void mca_set_adapter_procfn (int slot, MCA_ProcFn procfn, void*
dev);
```

Arguments

`slot`

slot to configure

`procfn`

callback function to call for /proc

`dev`

device information passed to the callback

Description

This sets up an information callback for /proc/mca/slot?. The function is called with the buffer, slot, and device pointer (or some equally informative context information, or nothing, if you prefer), and is expected to put useful information into the buffer. The adapter name, ID, and POS registers get printed before this is called though, so don’t do it again.

This should be called with a `NULL` *`procfn`* when a module unregisters, thus preventing kernel crashes and other such nastiness.

mca_is_adapter_used

Name

mca_is_adapter_used — check if claimed by driver

Synopsis

```
int mca_is_adapter_used (int slot);
```

Arguments

slot

slot to check

Description

Returns 1 if the slot has been claimed by a driver

mca_mark_as_used

Name

mca_mark_as_used — claim an MCA device

Synopsis

```
int mca_mark_as_used (int slot);
```

Arguments

slot

slot to claim

FIXME

should we make this threadsafe

Claim an MCA slot for a device driver. If the slot is already taken the function returns 1, if it is not taken it is claimed and 0 is returned.

mca_mark_as_unused

Name

mca_mark_as_unused — release an MCA device

Synopsis

```
void mca_mark_as_unused (int slot);
```

Arguments

slot

slot to claim

Description

Release the slot for other drives to use.

mca_get_adapter_name

Name

mca_get_adapter_name — get the adapter description

Synopsis

```
char * mca_get_adapter_name (int slot);
```

Arguments

slot

slot to query

Description

Return the adapter description if set. If it has not been set or the slot is out range then return NULL.

mca_isadapter

Name

`mca_isadapter` — check if the slot holds an adapter

Synopsis

```
int mca_isadapter (int slot);
```

Arguments

`slot`

slot to query

Description

Returns zero if the slot does not hold an adapter, non zero if it does.

mca_isenabled

Name

`mca_isenabled` — check if the slot holds an adapter

Synopsis

```
int mca_isenabled (int slot);
```

Arguments

`slot`

slot to query

Description

Returns a non zero value if the slot holds an enabled adapter and zero for any other case.

10.4.2. MCA Bus DMA

mca_enable_dma

Name

`mca_enable_dma` — channel to enable DMA on

Synopsis

```
void mca_enable_dma (unsigned int dmanr);
```

Arguments

`dmanr`

DMA channel

Description

Enable the MCA bus DMA on a channel. This can be called from IRQ context.

mca_disable_dma

Name

`mca_disable_dma` — channel to disable DMA on

Synopsis

```
void mca_disable_dma (unsigned int dmanr);
```

Arguments

dmanr

DMA channel

Description

Enable the MCA bus DMA on a channel. This can be called from IRQ context.

mca_set_dma_addr

Name

mca_set_dma_addr — load a 24bit DMA address

Synopsis

```
void mca_set_dma_addr (unsigned int dmanr, unsigned int a);
```

Arguments

dmanr

DMA channel

a

24bit bus address

Description

Load the address register in the DMA controller. This has a 24bit limitation (16Mb).

mca_get_dma_addr

Name

mca_get_dma_addr — load a 24bit DMA address

Synopsis

```
unsigned int mca_get_dma_addr (unsigned int dmanr);
```

Arguments

dmanr

DMA channel

Description

Read the address register in the DMA controller. This has a 24bit limitation (16Mb). The return is a bus address.

mca_set_dma_count

Name

`mca_set_dma_count` — load a 16bit transfer count

Synopsis

```
void mca_set_dma_count (unsigned int dmanr, unsigned int count);
```

Arguments

`dmanr`

DMA channel

`count`

count

Description

Set the DMA count for this channel. This can be up to 64Kbytes. Setting a count of zero will not do what you expect.

mca_get_dma_residue

Name

`mca_get_dma_residue` — get the remaining bytes to transfer

Synopsis

```
unsigned int mca_get_dma_residue (unsigned int dmanr);
```

Arguments

`dmanr`

DMA channel

Description

This function returns the number of bytes left to transfer on this DMA channel.

mca_set_dma_io

Name

`mca_set_dma_io` — set the port for an I/O transfer

Synopsis

```
void mca_set_dma_io (unsigned int dmanr, unsigned int io_addr);
```

Arguments

dmanr

DMA channel

io_addr

an I/O port number

Description

Unlike the ISA bus DMA controllers the DMA on MCA bus can transfer with an I/O port target.

mca_set_dma_mode

Name

mca_set_dma_mode — set the DMA mode

Synopsis

```
void mca_set_dma_mode (unsigned int dmanr, unsigned int mode);
```

Arguments

dmanr

DMA channel

mode

mode to set

Description

The DMA controller supports several modes. The mode values you can

set are

MCA_DMA_MODE_READ when reading from the DMA device.

MCA_DMA_MODE_WRITE to writing to the DMA device.

MCA_DMA_MODE_IO to do DMA to or from an I/O port.

MCA_DMA_MODE_16 to do 16bit transfers.

Chapter 11. The Device File System

devfs_register

Name

`devfs_register` — Register a device entry.

Synopsis

```
devfs_handle_t devfs_register (devfs_handle_t dir, const char *
name, unsigned int flags, unsigned int major, unsigned int
minor, umode_t mode, void * ops, void * info);
```

Arguments

`dir`

The handle to the parent devfs directory entry. If this is `NULL` the new name is relative to the root of the devfs.

`name`

The name of the entry.

`flags`

A set of bitwise-ORed flags (DEVFS_FL_*).

`major`

The major number. Not needed for regular files.

`minor`

The minor number. Not needed for regular files.

mode

The default file mode.

ops

The &file_operations or &block_device_operations structure. This must not be externally deallocated.

info

An arbitrary pointer which will be written to the *private_data* field of the &file structure passed to the device driver. You can set this to whatever you like, and change it once the file is opened (the next file opened will not see this change).

Description

Returns a handle which may later be used in a call to devfs_unregister. On failure NULL is returned.

devfs_unregister

Name

devfs_unregister — Unregister a device entry.

Synopsis

```
void devfs_unregister (devfs_handle_t de);
```

Arguments

de

A handle previously created by `devfs_register` or returned from `devfs_find_handle`. If this is `NULL` the routine does nothing.

devfs_mk_symlink

Name

`devfs_mk_symlink` —

Synopsis

```
int devfs_mk_symlink (devfs_handle_t dir, const char * name,
unsigned int flags, const char * link, devfs_handle_t * handle,
void * info);
```

Arguments

dir

The handle to the parent devfs directory entry. If this is `NULL` the new name is relative to the root of the devfs.

name

The name of the entry.

flags

A set of bitwise-ORed flags (DEVFS_FL_*).

link

The destination name.

handle

The handle to the symlink entry is written here. This may be NULL.

info

An arbitrary pointer which will be associated with the entry.

Description

Returns 0 on success, else a negative error code is returned.

devfs_mk_dir

Name

devfs_mk_dir — Create a directory in the devfs namespace.

Synopsis

```
devfs_handle_t devfs_mk_dir (devfs_handle_t dir, const char *
name, void * info);
```

Arguments

dir

The handle to the parent devfs directory entry. If this is NULL the new name is relative to the root of the devfs.

name

The name of the entry.

info

An arbitrary pointer which will be associated with the entry.

Description

Use of this function is optional. The `devfs_register` function will automatically create intermediate directories as needed. This function is provided for efficiency reasons, as it provides a handle to a directory. Returns a handle which may later be used in a call to `devfs_unregister`. On failure `NULL` is returned.

devfs_find_handle

Name

`devfs_find_handle` — Find the handle of a devfs entry.

Synopsis

```
devfs_handle_t devfs_find_handle (devfs_handle_t dir, const char
* name, unsigned int major, unsigned int minor, char type, int
traverse_symlinks);
```

Arguments

dir

The handle to the parent devfs directory entry. If this is `NULL` the name is relative

to the root of the devfs.

name

The name of the entry.

major

The major number. This is used if *name* is NULL.

minor

The minor number. This is used if *name* is NULL.

type

The type of special file to search for. This may be either DEVFS_SPECIAL_CHR or DEVFS_SPECIAL_BLK.

traverse_symlinks

If TRUE then symlink entries in the devfs namespace are traversed. Symlinks pointing out of the devfs namespace will cause a failure. Symlink traversal consumes stack space.

Description

Returns a handle which may later be used in a call to devfs_unregister, devfs_get_flags, or devfs_set_flags. On failure NULL is returned.

devfs_get_flags

Name

devfs_get_flags — Get the flags for a devfs entry.

Synopsis

```
int devfs_get_flags (devfs_handle_t de, unsigned int * flags);
```

Arguments

de

The handle to the device entry.

flags

The flags are written here.

Description

Returns 0 on success, else a negative error code.

devfs_get_maj_min

Name

devfs_get_maj_min — Get the major and minor numbers for a devfs entry.

Synopsis

```
int devfs_get_maj_min (devfs_handle_t de, unsigned int * major,
unsigned int * minor);
```

Arguments

de

The handle to the device entry.

major

The major number is written here. This may be NULL.

minor

The minor number is written here. This may be NULL.

Description

Returns 0 on success, else a negative error code.

devfs_get_handle_from_inode

Name

devfs_get_handle_from_inode — Get the devfs handle for a VFS inode.

Synopsis

```
devfs_handle_t devfs_get_handle_from_inode (struct inode *
inode);
```

Arguments

inode

The VFS inode.

Description

Returns the devfs handle on success, else `NULL`.

devfs_generate_path

Name

`devfs_generate_path` — Generate a pathname for an entry, relative to the devfs root.

Synopsis

```
int devfs_generate_path (devfs_handle_t de, char * path, int
buflen);
```

Arguments

de

The devfs entry.

path

The buffer to write the pathname to. The pathname and '\0' terminator will be written at the end of the buffer.

buflen

The length of the buffer.

Description

Returns the offset in the buffer where the pathname starts on success, else a negative error code.

devfs_get_ops

Name

`devfs_get_ops` — Get the device operations for a devfs entry.

Synopsis

```
void * devfs_get_ops (devfs_handle_t de);
```

Arguments

de

The handle to the device entry.

Description

Returns a pointer to the device operations on success, else NULL.

devfs_set_file_size

Name

`devfs_set_file_size` — Set the file size for a devfs regular file.

Synopsis

```
int devfs_set_file_size (devfs_handle_t de, unsigned long size);
```

Arguments

de

The handle to the device entry.

size

The new file size.

Description

Returns 0 on success, else a negative error code.

devfs_get_info

Name

`devfs_get_info` — Get the info pointer written to private_data of *de* upon open.

Synopsis

```
void * devfs_get_info (devfs_handle_t de);
```

Arguments

de

The handle to the device entry.

Description

Returns the info pointer.

devfs_set_info

Name

`devfs_set_info` — Set the info pointer written to private_data upon open.

Synopsis

```
int devfs_set_info (devfs_handle_t de, void * info);
```

Arguments

de

The handle to the device entry.

info

pointer to the data

Description

Returns 0 on success, else a negative error code.

devfs_get_parent

Name

devfs_get_parent — Get the parent device entry.

Synopsis

```
devfs_handle_t devfs_get_parent (devfs_handle_t de);
```

Arguments

de

The handle to the device entry.

Description

Returns the parent device entry if it exists, else `NULL`.

devfs_get_first_child

Name

`devfs_get_first_child` — Get the first leaf node in a directory.

Synopsis

```
devfs_handle_t devfs_get_first_child (devfs_handle_t de);
```

Arguments

de

The handle to the device entry.

Description

Returns the leaf node device entry if it exists, else `NULL`.

devfs_get_next_sibling

Name

`devfs_get_next_sibling` — Get the next sibling leaf node. for a device entry.

Synopsis

```
devfs_handle_t devfs_get_next_sibling (devfs_handle_t de);
```

Arguments

de

The handle to the device entry.

Description

Returns the leaf node device entry if it exists, else `NULL`.

devfs_auto_unregister

Name

`devfs_auto_unregister` — Configure a devfs entry to be automatically

unregistered.

Synopsis

```
void devfs_auto_unregister (devfs_handle_t master,
devfs_handle_t slave);
```

Arguments

master

The master devfs entry. Only one slave may be registered.

slave

The devfs entry which will be automatically unregistered when the master entry is unregistered. It is illegal to call `devfs_unregister` on this entry.

devfs_get_unregister_slave

Name

`devfs_get_unregister_slave` — Get the slave entry which will be automatically unregistered.

Synopsis

```
devfs_handle_t devfs_get_unregister_slave (devfs_handle_t
master);
```

Arguments

master

The master devfs entry.

Description

Returns the slave which will be unregistered when *master* is unregistered.

devfs_get_name

Name

devfs_get_name — Get the name for a device entry in its parent directory.

Synopsis

```
const char * devfs_get_name (devfs_handle_t de, unsigned int *
namelen);
```

Arguments

de

The handle to the device entry.

namelen

The length of the name is written here. This may be NULL.

Description

Returns the name on success, else `NULL`.

devfs_register_chrdev

Name

`devfs_register_chrdev` — Optionally register a conventional character driver.

Synopsis

```
int devfs_register_chrdev (unsigned int major, const char *
name, struct file_operations * fops);
```

Arguments

`major`

The major number for the driver.

`name`

The name of the driver (as seen in /proc/devices).

`fops`

The &file_operations structure pointer.

Description

This function will register a character driver provided the “devfs=only” option was not provided at boot time. Returns 0 on success, else a negative error code on failure.

devfs_register_blkdev

Name

`devfs_register_blkdev` — Optionally register a conventional block driver.

Synopsis

```
int devfs_register_blkdev (unsigned int major, const char *
name, struct block_device_operations * bdops);
```

Arguments

`major`

The major number for the driver.

`name`

The name of the driver (as seen in /proc/devices).

`bdops`

The &block_device_operations structure pointer.

Description

This function will register a block driver provided the “devfs=only” option was not provided at boot time. Returns 0 on success, else a negative error code on failure.

devfs_unregister_chrdev

Name

`devfs_unregister_chrdev` — Optionally unregister a conventional character driver.

Synopsis

```
int devfs_unregister_chrdev (unsigned int major, const char * name);
```

Arguments

`major`

The major number for the driver.

`name`

The name of the driver (as seen in /proc/devices).

Description

This function will unregister a character driver provided the “devfs=only” option was not provided at boot time. Returns 0 on success, else a negative error code on failure.

devfs_unregister_blkdev

Name

devfs_unregister_blkdev — Optionally unregister a conventional block driver.

Synopsis

```
int devfs_unregister_blkdev (unsigned int major, const char *
name);
```

Arguments

major

The major number for the driver.

name

The name of the driver (as seen in /proc/devices).

Description

This function will unregister a block driver provided the “devfs=only” option was not provided at boot time. Returns 0 on success, else a negative error code on failure.

Chapter 12. Power Management

pm_register

Name

pm_register — register a device with power management

Synopsis

```
struct pm_dev * pm_register (pm_dev_t type, unsigned long id,
pm_callback callback);
```

Arguments

type

device type

id

device ID

callback

callback function

Description

Add a device to the list of devices that wish to be notified about power management events. A &pm_dev structure is returned on success, on failure the return is NULL.

The callback function will be called in process context and it may sleep.

pm_unregister

Name

pm_unregister — unregister a device with power management

Synopsis

```
void pm_unregister (struct pm_dev * dev);
```

Arguments

dev

device to unregister

Description

Remove a device from the power management notification lists. The dev passed must be a handle previously returned by pm_register.

pm_unregister_all

Name

pm_unregister_all — unregister all devices with matching callback

Synopsis

```
void pm_unregister_all (pm_callback callback);
```

Arguments

callback

callback function pointer

Description

Unregister every device that would call the callback passed. This is primarily meant as a helper function for loadable modules. It enables a module to give up all its managed devices without keeping its own private list.

pm_send

Name

pm_send — send request to a single device

Synopsis

```
int pm_send (struct pm_dev * dev, pm_request_t rqst, void *
data);
```

Arguments

dev

device to send to

rqst

power management request

data

data for the callback

Description

Issue a power management request to a given device. The PM_SUSPEND and PM_RESUME events are handled specially. The data field must hold the intended next state. No call is made if the state matches.

BUGS

what stops two power management requests occuring in parallel and conflicting.

WARNING

Calling pm_send directly is not generally recommended, in paticular there is no locking against the pm_dev going away. The caller must maintain all needed locking or have 'inside knowledge' on the safety. Also remember that this function is not locked against pm_unregister. This means that you must handle SMP races on callback execution and unload yourself.

pm_send_all

Name

pm_send_all — send request to all managed devices

Synopsis

```
int pm_send_all (pm_request_t rqst, void * data);
```

Arguments

rqst

power management request

data

data for the callback

Description

Issue a power management request to a all devices. The PM_SUSPEND events are handled specially. Any device is permitted to fail a suspend by returning a non zero (error) value from its callback function. If any device vetoes a suspend request then all other devices that have suspended during the processing of this request are restored to their previous state.

WARNING

This function takes the pm_devs_lock. The lock is not dropped until the callbacks have completed. This prevents races against pm locking functions, races against module unload pm_unregister code. It does mean however that you must not issue pm_ functions within the callback or you will deadlock and users will hate you.

Zero is returned on success. If a suspend fails then the status from the device that vetoes the suspend is returned.

BUGS

what stops two power management requests occuring in parallel and conflicting.

pm_find

Name

`pm_find` — find a device

Synopsis

```
struct pm_dev * pm_find (pm_dev_t type, struct pm_dev * from);
```

Arguments

`type`

type of device

from

where to start looking

Description

Scan the power management list for devices of a specific type. The return value for a matching device may be passed to further calls to this function to find further matches. A NULL indicates the end of the list.

To search from the beginning pass NULL as the *from* value.

The caller MUST hold the pm_devs_lock lock when calling this function. The instant that the lock is dropped all pointers returned may become invalid.

Chapter 13. Block Devices

blk_cleanup_queue

Name

`blk_cleanup_queue` — release a request_queue_t when it is no longer needed

Synopsis

```
void blk_cleanup_queue (request_queue_t * q);
```

Arguments

q

the request queue to be released

Description

blk_cleanup_queue is the pair to `blk_init_queue`. It should be called when a request queue is being released; typically when a block device is being de-registered. Currently, its primary task it to free all the &struct request structures that were allocated to the queue.

Caveat

Hopefully the low level driver will have finished any outstanding requests first...

blk_queue_headactive

Name

`blk_queue_headactive` — indicate whether head of request queue may be active

Synopsis

```
void blk_queue_headactive (request_queue_t * q, int active);
```

Arguments

`q`

The queue which this applies to.

`active`

A flag indication where the head of the queue is active.

Description

The driver for a block device may choose to leave the currently active request on the request queue, removing it only when it has completed. The queue handling routines assume this by default for safety reasons and will not involve the head of the request queue in any merging or reordering of requests when the queue is unplugged (and thus may be working on this particular request).

If a driver removes requests from the queue before processing them, then it may indicate that it does so, there by allowing the head of the queue to be involved in merging and reordering. This is done be calling `blk_queue_headactive` with an *`active`* flag of `0`.

If a driver processes several requests at once, it must remove them (or at least all but one of them) from the request queue.

When a queue is plugged the head will be assumed to be inactive.

blk_queue_make_request

Name

`blk_queue_make_request` — define an alternate make_request function for a device

Synopsis

```
void blk_queue_make_request (request_queue_t * q,
make_request_fn * mfn);
```

Arguments

q

the request queue for the device to be affected

mfn

the alternate make_request function

Description

The normal way for &struct buffer_heads to be passed to a device driver is for them to be collected into requests on a request queue, and then to allow the device driver to select requests off that queue when it is ready. This works well for many block devices. However some block devices (typically virtual devices such as md or lvm) do not benefit from the processing on the request queue, and are served best by having the

requests passed directly to them. This can be achieved by providing a function to `blk_queue_make_request`.

Caveat

The driver that does this *must* be able to deal appropriately with buffers in “highmemory”, either by calling `bh_kmap` to get a kernel mapping, to by calling `create_bounce` to create a buffer in normal memory.

blk_init_queue

Name

`blk_init_queue` — prepare a request queue for use with a block device

Synopsis

```
void blk_init_queue (request_queue_t * q, request_fn_proc *
rfn);
```

Arguments

q

The &request_queue_t to be initialised

rfn

The function to be called to process requests that have been placed on the queue.

Description

If a block device wishes to use the standard request handling procedures, which sorts requests and coalesces adjacent requests, then it must call `blk_init_queue`. The function *`rfn`* will be called when there are requests on the queue that need to be processed. If the device supports plugging, then *`rfn`* may not be called immediately when requests are available on the queue, but may be called at some time later instead. Plugged queues are generally unplugged when a buffer belonging to one of the requests on the queue is needed, or due to memory pressure.

`rfn` is not required, or even expected, to remove all requests off the queue, but only as many as it can handle at a time. If it does leave requests on the queue, it is responsible for arranging that the requests get dealt with eventually.

A global spin lock $io_request_lock must be held while manipulating the requests on the request queue.

The request on the head of the queue is by default assumed to be potentially active, and it is not considered for re-ordering or merging whenever the given queue is unplugged. This behaviour can be changed with `blk_queue_headactive`.

Note

`blk_init_queue` must be paired with a `blk_cleanup_queue` call when the block device is deactivated (such as at module unload).

generic_make_request

Name

`generic_make_request` —

Synopsis

```
void generic_make_request (int rw, struct buffer_head * bh);
```

Arguments

rw

READ, WRITE, or READA - what sort of I/O is desired.

bh

The buffer head describing the location in memory and on the device.

Description

`generic_make_request` is used to make I/O requests of block devices. It is passed a &struct buffer_head and a &rw value. The `READ` and `WRITE` options are (hopefully) obvious in meaning. The `READA` value means that a read is required, but that the driver is free to fail the request if, for example, it cannot get needed resources immediately.

`generic_make_request` does not return any status. The success/failure status of the request, along with notification of completion, is delivered asynchronously through the bh->b_end_io function described (one day) else where.

The caller of generic_make_request must make sure that b_page, b_addr, b_size are set to describe the memory buffer, that b_rdev and b_rsector are set to describe the device address, and the b_end_io and optionally b_private are set to describe how completion notification should be signaled. BH_Mapped should also be set (to confirm that b_dev and b_blocknr are valid).

generic_make_request and the drivers it calls may use b_reqnext, and may change b_rdev and b_rsector. So the values of these fields should NOT be depended on after the call to generic_make_request. Because of this, the caller should record the device address information in b_dev and b_blocknr.

Apart from those fields mentioned above, no other fields, and in particular, no other flags, are changed by generic_make_request or any lower level drivers.

submit_bh

Name

submit_bh —

Synopsis

```
void submit_bh (int rw, struct buffer_head * bh);
```

Arguments

rw

whether to READ or WRITE, or maybe to READA (read ahead)

bh

The &struct buffer_head which describes the I/O

Description

submit_bh is very similar in purpose to generic_make_request, and uses that function to do most of the work.

The extra functionality provided by submit_bh is to determine b_rsector from b_blocknr and b_size, and to set b_rdev from b_dev. This is is appropriate for IO requests that come from the buffer cache and page cache which (currently) always use aligned blocks.

ll_rw_block

Name

`ll_rw_block` — level access to block devices

Synopsis

```
void ll_rw_block (int rw, int nr, struct buffer_head * * bhs);
```

Arguments

rw

whether to `READ` or `WRITE` or maybe `READA` (readahead)

nr

number of &struct buffer_heads in the array

bhs

array of pointers to &struct buffer_head

Description

`ll_rw_block` takes an array of pointers to &struct buffer_heads, and requests an I/O operation on them, either a `READ` or a `WRITE`. The third `READA` option is described in the documentation for `generic_make_request` which `ll_rw_block` calls.

This function provides extra functionality that is not in `generic_make_request` that is relevant to buffers in the buffer cache or page cache. In particular it drops any buffer that it cannot get a lock on (with the BH_Lock state bit), any buffer that appears to be clean when doing a write request, and any buffer that appears to be up-to-date when doing read request. Further it marks as clean buffers that are processed for writing (the buffer cache wont assume that they are actually clean until the buffer gets unlocked).

ll_rw_block sets b_end_io to simple completion handler that marks the buffer up-to-date (if approriate), unlocks the buffer and wakes any waiters. As client that needs a more interesting completion routine should call `submit_bh` (or `generic_make_request`) directly.

Caveat

All of the buffers must be for the same device, and must also be

end_that_request_first

Name

`end_that_request_first` — end I/O on one buffer.

Synopsis

```
int end_that_request_first (struct request * req, int uptodate,
char * name);
```

Arguments

`req`

the request being processed

`uptodate`

0 for I/O error

name

the name printed for an I/O error

Description

Ends I/O on the first buffer attached to *req*, and sets it up for the next buffer_head (if any) in the cluster.

Return

0 - we are done with this request, call `end_that_request_last` 1 - still buffers pending for this request

Caveat

Drivers implementing their own end_request handling must call `blk_finished_io` appropriately.

Chapter 14. Miscellaneous Devices

misc_register

Name

`misc_register` — register a miscellaneous device

Synopsis

```
int misc_register (struct miscdevice * misc);
```

Arguments

`misc`

device structure

Description

Register a miscellaneous device with the kernel. If the minor number is set to `MISC_DYNAMIC_MINOR` a minor number is assigned and placed in the minor field of the structure. For other cases the minor number requested is used.

The structure passed is linked into the kernel and may not be destroyed until it has been unregistered.

A zero is returned on success and a negative errno code for failure.

misc_deregister

Name

`misc_deregister` — unregister a miscellaneous device

Synopsis

```
int misc_deregister (struct miscdevice * misc);
```

Arguments

`misc`

device to unregister

Description

Unregister a miscellaneous device that was previously successfully registered with `misc_register`. Success is indicated by a zero return, a negative errno code indicates an error.

Chapter 15. Video4Linux

video_unregister_device

Name

`video_unregister_device` — unregister a video4linux device

Synopsis

```
void video_unregister_device (struct video_device * vfd);
```

Arguments

`vfd`

the device to unregister

Description

This unregisters the passed device and deassigns the minor number. Future open calls will be met with errors.

Chapter 16. Sound Devices

register_sound_special

Name

`register_sound_special` — register a special sound node

Synopsis

```
int register_sound_special (struct file_operations * fops, int
unit);
```

Arguments

`fops`

File operations for the driver

`unit`

Unit number to allocate

Description

Allocate a special sound device by minor number from the sound subsystem. The allocated number is returned on succes. On failure a negative error code is returned.

register_sound_mixer

Name

`register_sound_mixer` — register a mixer device

Synopsis

```
int register_sound_mixer (struct file_operations * fops, int
dev);
```

Arguments

`fops`

File operations for the driver

`dev`

Unit number to allocate

Description

Allocate a mixer device. Unit is the number of the mixer requested. Pass -1 to request the next free mixer unit. On success the allocated number is returned, on failure a negative error code is returned.

register_sound_midi

Name

`register_sound_midi` — register a midi device

Synopsis

```
int register_sound_midi (struct file_operations * fops, int
dev);
```

Arguments

`fops`

File operations for the driver

`dev`

Unit number to allocate

Description

Allocate a midi device. Unit is the number of the midi device requested. Pass -1 to request the next free midi unit. On success the allocated number is returned, on failure a negative error code is returned.

register_sound_dsp

Name

`register_sound_dsp` — register a DSP device

Synopsis

```
int register_sound_dsp (struct file_operations * fops, int dev);
```

Arguments

`fops`

File operations for the driver

`dev`

Unit number to allocate

Description

Allocate a DSP device. Unit is the number of the DSP requested. Pass -1 to request the next free DSP unit. On success the allocated number is returned, on failure a negative error code is returned.

This function allocates both the audio and dsp device entries together and will always allocate them as a matching pair - eg dsp3/audio3

register_sound_synth

Name

`register_sound_synth` — register a synth device

Synopsis

```
int register_sound_synth (struct file_operations * fops, int
dev);
```

Arguments

`fops`

File operations for the driver

`dev`

Unit number to allocate

Description

Allocate a synth device. Unit is the number of the synth device requested. Pass -1 to request the next free synth unit. On success the allocated number is returned, on failure a negative error code is returned.

unregister_sound_special

Name

`unregister_sound_special` — unregister a special sound device

Synopsis

```
void unregister_sound_special (int unit);
```

Arguments

`unit`

unit number to allocate

Description

Release a sound device that was allocated with `register_sound_special`. The unit passed is the return value from the register function.

unregister_sound_mixer

Name

`unregister_sound_mixer` — unregister a mixer

Synopsis

```
void unregister_sound_mixer (int unit);
```

Arguments

unit

unit number to allocate

Description

Release a sound device that was allocated with `register_sound_mixer`. The unit passed is the return value from the register function.

unregister_sound_midi

Name

`unregister_sound_midi` — unregister a midi device

Synopsis

```
void unregister_sound_midi (int unit);
```

Arguments

unit

unit number to allocate

Description

Release a sound device that was allocated with `register_sound_midi`. The unit passed is the return value from the register function.

unregister_sound_dsp

Name

`unregister_sound_dsp` — unregister a DSP device

Synopsis

```
void unregister_sound_dsp (int unit);
```

Arguments

unit

unit number to allocate

Description

Release a sound device that was allocated with `register_sound_dsp`. The unit passed is the return value from the register function.

Both of the allocated units are released together automatically.

unregister_sound_synth

Name

`unregister_sound_synth` — unregister a synth device

Synopsis

```
void unregister_sound_synth (int unit);
```

Arguments

`unit`

unit number to allocate

Description

Release a sound device that was allocated with `register_sound_synth`. The unit passed is the return value from the register function.

Chapter 17. USB Devices

usb_register

Name

`usb_register` — register a USB driver

Synopsis

```
int usb_register (struct usb_driver * new_driver);
```

Arguments

`new_driver`

USB operations for the driver

Description

Registers a USB driver with the USB core. The list of unattached interfaces will be rescanned whenever a new driver is added, allowing the new driver to attach to any recognized devices. Returns a negative error code on failure and 0 on success.

usb_scan_devices

Name

`usb_scan_devices` — scans all unclaimed USB interfaces

Synopsis

```
void usb_scan_devices ( void);
```

Arguments

`void`

no arguments

Description

Goes through all unclaimed USB interfaces, and offers them to all registered USB drivers through the ’probe’ function. This will automatically be called after usb_register is called. It is called by some of the USB subsystems after one of their subdrivers are registered.

usb_deregister

Name

`usb_deregister` — unregister a USB driver

Synopsis

```
void usb_deregister (struct usb_driver * driver);
```

Arguments

driver

USB operations of the driver to unregister

Description

Unlinks the specified driver from the internal USB driver list.

usb_alloc_bus

Name

usb_alloc_bus — creates a new USB host controller structure

Synopsis

```
struct usb_bus * usb_alloc_bus (struct usb_operations * op);
```

Arguments

op

pointer to a struct usb_operations that this bus structure should use

Description

Creates a USB host controller bus structure with the specified usb_operations and initializes all the necessary internal objects. (For use only by USB Host Controller Drivers.)

If no memory is available, NULL is returned.

The caller should call `usb_free_bus` when it is finished with the structure.

usb_free_bus

Name

`usb_free_bus` — frees the memory used by a bus structure

Synopsis

```
void usb_free_bus (struct usb_bus * bus);
```

Arguments

bus

pointer to the bus to free

Description

(For use only by USB Host Controller Drivers.)

usb_register_bus

Name

`usb_register_bus` — registers the USB host controller with the usb core

Synopsis

```
void usb_register_bus (struct usb_bus * bus);
```

Arguments

bus

pointer to the bus to register

Description

(For use only by USB Host Controller Drivers.)

usb_deregister_bus

Name

usb_deregister_bus — deregisters the USB host controller

Synopsis

```
void usb_deregister_bus (struct usb_bus * bus);
```

Arguments

bus

pointer to the bus to deregister

Description

(For use only by USB Host Controller Drivers.)

usb_match_id

Name

usb_match_id — find first usb_device_id matching device or interface

Synopsis

```
const struct usb_device_id * usb_match_id (struct usb_device *
dev, struct usb_interface * interface, const struct
usb_device_id * id);
```

Arguments

dev

the device whose descriptors are considered when matching

interface

the interface of interest

id

array of usb_device_id structures, terminated by zero entry

Description

usb_match_id searches an array of usb_device_id's and returns the first one matching the device or interface, or null. This is used when binding (or rebinding) a driver to an interface. Most USB device drivers will use this indirectly, through the usb core, but some layered driver frameworks use it directly. These device tables are exported with MODULE_DEVICE_TABLE, through modutils and "modules.usbmap", to support the driver loading functionality of USB hotplugging.

What Matches

The "match_flags" element in a usb_device_id controls which members are used. If the corresponding bit is set, the value in the device_id must match its corresponding member in the device or interface descriptor, or else the device_id does not match.

"driver_info" is normally used only by device drivers, but you can create a wildcard "matches anything" usb_device_id as a driver's "modules.usbmap" entry if you provide

an id with only a nonzero “driver_info” field. If you do this, the USB device driver’s `probe` routine should use additional intelligence to decide whether to bind to the specified interface.

What Makes Good usb_device_id Tables

The match algorithm is very simple, so that intelligence in driver selection must come from smart driver id records. Unless you have good reasons to use another selection policy, provide match elements only in related groups, and order match specifiers from specific to general. Use the macros provided for that purpose if you can.

The most specific match specifiers use device descriptor data. These are commonly used with product-specific matches; the USB_DEVICE macro lets you provide vendor and product IDs, and you can also match against ranges of product revisions. These are widely used for devices with application or vendor specific bDeviceClass values.

Matches based on device class/subclass/protocol specifications are slightly more general; use the USB_DEVICE_INFO macro, or its siblings. These are used with single-function devices where bDeviceClass doesn’t specify that each interface has its own class.

Matches based on interface class/subclass/protocol are the most general; they let drivers bind to any interface on a multiple-function device. Use the USB_INTERFACE_INFO macro, or its siblings, to match class-per-interface style devices (as recorded in bDeviceClass).

Within those groups, remember that not all combinations are meaningful. For example, don’t give a product version range without vendor and product IDs; or specify a protocol without its associated class and subclass.

usb_alloc_urb

Name

`usb_alloc_urb` — creates a new urb for a USB driver to use

Synopsis

```
urb_t * usb_alloc_urb (int iso_packets);
```

Arguments

iso_packets

number of iso packets for this urb

Description

Creates an urb for the USB driver to use and returns a pointer to it. If no memory is available, NULL is returned.

If the driver want to use this urb for interrupt, control, or bulk endpoints, pass ’0’ as the number of iso packets.

The driver should call `usb_free_urb` when it is finished with the urb.

usb_free_urb

Name

`usb_free_urb` — frees the memory used by a urb

Synopsis

```
void usb_free_urb (urb_t* urb);
```

Arguments

urb

pointer to the urb to free

Description

If an urb is created with a call to `usb_create_urb` it should be cleaned up with a call to `usb_free_urb` when the driver is finished with it.

usb_control_msg

Name

`usb_control_msg` — Builds a control urb, sends it off and waits for completion

Synopsis

```
int usb_control_msg (struct usb_device * dev, unsigned int pipe,
__u8 request, __u8 requesttype, __u16 value, __u16 index, void *
data, __u16 size, int timeout);
```

Arguments

dev

pointer to the usb device to send the message to

pipe

endpoint “pipe” to send the message to

request

USB message request value

requesttype

USB message request type value

value

USB message value

index

USB message index value

data

pointer to the data to send

size

length in bytes of the data to send

timeout

time to wait for the message to complete before timing out (if 0 the wait is forever)

Description

This function sends a simple control message to a specified endpoint and waits for the message to complete, or timeout.

If successful, it returns 0, othwise a negative error number.

Don’t use this function from within an interrupt context, like a bottom half handler. If you need a asyncronous message, or need to send a message from within interrupt context, use `usb_submit_urb`

usb_bulk_msg

Name

`usb_bulk_msg` — Builds a bulk urb, sends it off and waits for completion

Synopsis

```
int usb_bulk_msg (struct usb_device * usb_dev, unsigned int
pipe, void * data, int len, int * actual_length, int timeout);
```

Arguments

`usb_dev`

pointer to the usb device to send the message to

`pipe`

endpoint “pipe” to send the message to

`data`

pointer to the data to send

`len`

length in bytes of the data to send

`actual_length`

pointer to a location to put the actual length transferred in bytes

`timeout`

time to wait for the message to complete before timing out (if 0 the wait is forever)

Description

This function sends a simple bulk message to a specified endpoint and waits for the message to complete, or timeout.

If successful, it returns 0, othwise a negative error number. The number of actual bytes transferred will be plaed in the actual_timeout paramater.

Don’t use this function from within an interrupt context, like a bottom half handler. If you need a asyncronous message, or need to send a message from within interrupt context, use `usb_submit_urb`

Chapter 18. 16x50 UART Driver

register_serial

Name

`register_serial` — configure a 16x50 serial port at runtime

Synopsis

```
int register_serial (struct serial_struct * req);
```

Arguments

`req`

request structure

Description

Configure the serial port specified by the request. If the port exists and is in use an error is returned. If the port is not currently in the table it is added.

The port is then probed and if neccessary the IRQ is autodetected If this fails an error is returned.

On success the port is ready to use and the line number is returned.

unregister_serial

Name

`unregister_serial` — deconfigure a 16x50 serial port

Synopsis

```
void unregister_serial (int line);
```

Arguments

`line`

line to deconfigure

Description

The port specified is deconfigured and its resources are freed. Any user of the port is disconnected as if carrier was dropped. Line is the port number returned by `register_serial`.

Chapter 19. Z85230 Support Library

z8530_interrupt

Name

`z8530_interrupt` — Handle an interrupt from a Z8530

Synopsis

```
void z8530_interrupt (int irq, void * dev_id, struct pt_regs *
regs);
```

Arguments

`irq`

Interrupt number

`dev_id`

The Z8530 device that is interrupting.

`regs`

unused

Description

A Z85[2]30 device has stuck its hand in the air for attention. We scan both the channels on the chip for events and then call the channel specific call backs for each channel that has events. We have to use callback functions because the two channels can be in different modes.

z8530_sync_open

Name

z8530_sync_open — Open a Z8530 channel for PIO

Synopsis

```
int z8530_sync_open (struct net_device * dev, struct
z8530_channel * c);
```

Arguments

dev

The network interface we are using

c

The Z8530 channel to open in synchronous PIO mode

Description

Switch a Z8530 into synchronous mode without DMA assist. We raise the RTS/DTR and commence network operation.

z8530_sync_close

Name

z8530_sync_close — Close a PIO Z8530 channel

Synopsis

```
int z8530_sync_close (struct net_device * dev, struct
z8530_channel * c);
```

Arguments

dev

Network device to close

c

Z8530 channel to disassociate and move to idle

Description

Close down a Z8530 interface and switch its interrupt handlers to discard future events.

z8530_sync_dma_open

Name

z8530_sync_dma_open — Open a Z8530 for DMA I/O

Synopsis

```
int z8530_sync_dma_open (struct net_device * dev, struct
z8530_channel * c);
```

Arguments

dev

The network device to attach

c

The Z8530 channel to configure in sync DMA mode.

Description

Set up a Z85x30 device for synchronous DMA in both directions. Two ISA DMA channels must be available for this to work. We assume ISA DMA driven I/O and PC limits on access.

z8530_sync_dma_close

Name

`z8530_sync_dma_close` — Close down DMA I/O

Synopsis

```
int z8530_sync_dma_close (struct net_device * dev, struct
z8530_channel * c);
```

Arguments

dev

Network device to detach

c

Z8530 channel to move into discard mode

Description

Shut down a DMA mode synchronous interface. Halt the DMA, and free the buffers.

z8530_sync_txdma_open

Name

`z8530_sync_txdma_open` — Open a Z8530 for TX driven DMA

Synopsis

```
int z8530_sync_txdma_open (struct net_device * dev, struct
z8530_channel * c);
```

Arguments

dev

The network device to attach

c

The Z8530 channel to configure in sync DMA mode.

Description

Set up a Z85x30 device for synchronous DMA tranmission. One ISA DMA channel must be available for this to work. The receive side is run in PIO mode, but then it has the bigger FIFO.

z8530_sync_txdma_close

Name

`z8530_sync_txdma_close` — Close down a TX driven DMA channel

Synopsis

```
int z8530_sync_txdma_close (struct net_device * dev, struct
z8530_channel * c);
```

Arguments

dev

Network device to detach

c

Z8530 channel to move into discard mode

Description

Shut down a DMA/PIO split mode synchronous interface. Halt the DMA, and free the buffers.

z8530_describe

Name

z8530_describe — Uniformly describe a Z8530 port

Synopsis

```
void z8530_describe (struct z8530_dev * dev, char * mapping,
unsigned long io);
```

Arguments

dev

Z8530 device to describe

mapping

string holding mapping type (eg “I/O” or “Mem”)

io

the port value in question

Description

Describe a Z8530 in a standard format. We must pass the I/O as the port offset isnt predictable. The main reason for this function is to try and get a common format of report.

z8530_init

Name

`z8530_init` — Initialise a Z8530 device

Synopsis

```
int z8530_init (struct z8530_dev * dev);
```

Arguments

dev

Z8530 device to initialise.

Description

Configure up a Z8530/Z85C30 or Z85230 chip. We check the device is present, identify the type and then program it to hopefully keep quite and behave. This matters a lot, a Z8530 in the wrong state will sometimes get into stupid modes generating 10Khz interrupt streams and the like.

We set the interrupt handler up to discard any events, in case we get them during reset or setp.

Return 0 for success, or a negative value indicating the problem in errno form.

z8530_shutdown

Name

`z8530_shutdown` — Shutdown a Z8530 device

Synopsis

```
int z8530_shutdown (struct z8530_dev * dev);
```

Arguments

`dev`

The Z8530 chip to shutdown

Description

We set the interrupt handlers to silence any interrupts. We then reset the chip and wait 100uS to be sure the reset completed. Just in case the caller then tries to do stuff.

z8530_channel_load

Name

`z8530_channel_load` — Load channel data

Synopsis

```
int z8530_channel_load (struct z8530_channel * c, u8 * rtable);
```

Arguments

`c`

Z8530 channel to configure

`rtable`

table of register, value pairs

FIXME

ioctl to allow user uploaded tables

Load a Z8530 channel up from the system data. We use +16 to indicate the “prime” registers. The value 255 terminates the table.

z8530_null_rx

Name

z8530_null_rx — Discard a packet

Synopsis

```
void z8530_null_rx (struct z8530_channel * c, struct sk_buff *
skb);
```

Arguments

c

The channel the packet arrived on

skb

The buffer

Description

We point the receive handler at this function when idle. Instead of syncppp processing the frames we get to throw them away.

z8530_queue_xmit

Name

`z8530_queue_xmit` — Queue a packet

Synopsis

```
int z8530_queue_xmit (struct z8530_channel * c, struct sk_buff *
skb);
```

Arguments

c

The channel to use

skb

The packet to kick down the channel

Description

Queue a packet for transmission. Because we have rather hard to hit interrupt latencies for the Z85230 per packet even in DMA mode we do the flip to DMA buffer if needed here not in the IRQ.

z8530_get_stats

Name

`z8530_get_stats` — Get network statistics

Synopsis

```
struct net_device_stats * z8530_get_stats (struct z8530_channel
* c);
```

Arguments

c

The channel to use

Description

Get the statistics block. We keep the statistics in software as the chip doesn’t do it for us.

Chapter 20. Frame Buffer Library

The frame buffer drivers depend heavily on four data structures. These structures are declared in include/linux/fb.h. They are fb_info, fb_var_screeninfo, fb_fix_screeninfo and fb_monospecs. The last three can be made available to and from userland.

fb_info defines the current state of a particular video card. Inside fb_info, there exists a fb_ops structure which is a collection of needed functions to make fbdev and fbcon work. fb_info is only visible to the kernel.

fb_var_screeninfo is used to describe the features of a video card that are user defined. With fb_var_screeninfo, things such as depth and the resolution may be defined.

The next structure is fb_fix_screeninfo. This defines the properties of a card that are created when a mode is set and can’t be changed otherwise. A good example of this is the start of the frame buffer memory. This "locks" the address of the frame buffer memory, so that it cannot be changed or moved.

The last structure is fb_monospecs. In the old API, there was little importance for fb_monospecs. This allowed for forbidden things such as setting a mode of 800x600 on a fix frequency monitor. With the new API, fb_monospecs prevents such things, and if used correctly, can prevent a monitor from being cooked. fb_monospecs will not be useful until kernels 2.5.x.

20.1. Frame Buffer Memory

register_framebuffer

Name

`register_framebuffer` — registers a frame buffer device

Synopsis

```
int register_framebuffer (struct fb_info * fb_info);
```

Arguments

fb_info

frame buffer info structure

Description

Registers a frame buffer device *fb_info*.

Returns negative errno on error, or zero for success.

unregister_framebuffer

Name

unregister_framebuffer — releases a frame buffer device

Synopsis

```
int unregister_framebuffer (struct fb_info * fb_info);
```

Arguments

fb_info

frame buffer info structure

Description

Unregisters a frame buffer device *fb_info*.

Returns negative errno on error, or zero for success.

20.2. Frame Buffer Console

fbcon_redraw_clear

Name

fbcon_redraw_clear — clear area of the screen

Synopsis

```
void fbcon_redraw_clear (struct vc_data * conp, struct display *
p, int sy, int sx, int height, int width);
```

Arguments

conp

stucture pointing to current active virtual console

p

display structure

sy

starting Y coordinate

sx

starting X coordinate

height

height of area to clear

width

width of area to clear

Description

Clears a specified area of the screen. All dimensions are in pixels.

fbcon_redraw_bmove

Name

fbcon_redraw_bmove — copy area of screen to another area

Synopsis

```
void fbcon_redraw_bmove (struct display * p, int sy, int sx, int
dy, int dx, int h, int w);
```

Arguments

p

display structure

sy

origin Y coordinate

sx

origin X coordinate

dy

destination Y coordinate

dx

destination X coordinate

h

height of area to copy

w

width of area to copy

Description

Copies an area of the screen to another area of the same screen. All dimensions are in pixels.

Note that this function cannot be used together with ypan or ywrap.

20.3. Frame Buffer Colormap

fb_alloc_cmap

Name

`fb_alloc_cmap` — allocate a colormap

Synopsis

```
int fb_alloc_cmap (struct fb_cmap * cmap, int len, int transp);
```

Arguments

`cmap`

frame buffer colormap structure

`len`

length of *`cmap`*

`transp`

boolean, 1 if there is transparency, 0 otherwise

Description

Allocates memory for a colormap *`cmap`*. *`len`* is the number of entries in the palette.

Returns -1 errno on error, or zero on success.

fb_copy_cmap

Name

fb_copy_cmap — copy a colormap

Synopsis

```
void fb_copy_cmap (struct fb_cmap * from, struct fb_cmap * to,
int fsfromto);
```

Arguments

from

frame buffer colormap structure

to

frame buffer colormap structure

fsfromto

determine copy method

Description

Copy contents of colormap from *from* to *to*.

0

memcpy function

1

`copy_from_user` function to copy from userspace

2

`copy_to_user` function to copy to userspace

fb_get_cmap

Name

`fb_get_cmap` — get a colormap

Synopsis

```
int fb_get_cmap (struct fb_cmap * cmap, int kspc, int
(*getcolreg) (u_int, u_int *, u_int *, u_int *, u_int *, struct
fb_info *), struct fb_info * info);
```

Arguments

`cmap`

frame buffer colormap

`kspc`

boolean, 0 copy local, 1 `put_user` function

getcolreg

pointer to a function to get a color register

info

frame buffer info structure

Description

Get a colormap *cmap* for a screen of device *info*.

Returns negative errno on error, or zero on success.

fb_set_cmap

Name

fb_set_cmap — set the colormap

Synopsis

```
int fb_set_cmap (struct fb_cmap * cmap, int kspc, int
(*setcolreg) (u_int, u_int, u_int, u_int, u_int, struct fb_info
*), struct fb_info * info);
```

Arguments

cmap

frame buffer colormap structure

kspc

boolean, 0 copy local, 1 `get_user` function

setcolreg

info

frame buffer info structure

Description

Sets the colormap *cmap* for a screen of device *info*.

Returns negative errno on error, or zero on success.

fb_default_cmap

Name

`fb_default_cmap` — get default colormap

Synopsis

```
struct fb_cmap * fb_default_cmap (int len);
```

Arguments

len

size of palette for a depth

Description

Gets the default colormap for a specific screen depth. *len* is the size of the palette for a particular screen depth.

Returns pointer to a frame buffer colormap structure.

fb_invert_cmaps

Name

fb_invert_cmaps — invert all defaults colormaps

Synopsis

```
void fb_invert_cmaps ( void);
```

Arguments

void

no arguments

Description

Invert all default colormaps.

20.4. Frame Buffer Generic Functions

fbgen_get_fix

Name

`fbgen_get_fix` — get fixed part of display

Synopsis

```
int fbgen_get_fix (struct fb_fix_screeninfo * fix, int con,
struct fb_info * info);
```

Arguments

`fix`

fb_fix_screeninfo structure

`con`

virtual console number

`info`

frame buffer info structure

Description

Get the fixed information part of the display and place it into *`fix`* for virtual console *`con`* on device *`info`*.

Returns negative errno on error, or zero on success.

fbgen_get_var

Name

`fbgen_get_var` — get user defined part of display

Synopsis

```
int fbgen_get_var (struct fb_var_screeninfo * var, int con,
struct fb_info * info);
```

Arguments

`var`

fb_var_screeninfo structure

`con`

virtual console number

`info`

frame buffer info structure

Description

Get the user defined part of the display and place it into *`var`* for virtual console *`con`* on device *`info`*.

Returns negative errno on error, or zero for success.

fbgen_set_var

Name

`fbgen_set_var` — set the user defined part of display

Synopsis

```
int fbgen_set_var (struct fb_var_screeninfo * var, int con,
struct fb_info * info);
```

Arguments

`var`

fb_var_screeninfo user defined part of the display

`con`

virtual console number

`info`

frame buffer info structure

Description

Set the user defined part of the display as dictated by *`var`* for virtual console *`con`* on device *`info`*.

Returns negative errno on error, or zero for success.

fbgen_get_cmap

Name

`fbgen_get_cmap` — get the colormap

Synopsis

```
int fbgen_get_cmap (struct fb_cmap * cmap, int kspc, int con,
struct fb_info * info);
```

Arguments

`cmap`

frame buffer colormap structure

`kspc`

boolean, 0 copy local, 1 `put_user` function

`con`

virtual console number

`info`

frame buffer info structure

Description

Gets the colormap for virtual console *`con`* and places it into *`cmap`* for device *`info`*.

Returns negative errno on error, or zero for success.

fbgen_set_cmap

Name

`fbgen_set_cmap` — set the colormap

Synopsis

```
int fbgen_set_cmap (struct fb_cmap * cmap, int kspc, int con,
struct fb_info * info);
```

Arguments

`cmap`

frame buffer colormap structure

`kspc`

boolean, 0 copy local, 1 `get_user` function

`con`

virtual console number

`info`

frame buffer info structure

Description

Sets the colormap *`cmap`* for virtual console *`con`* on device *`info`*.

Returns negative errno on error, or zero for success.

fbgen_pan_display

Name

`fbgen_pan_display` — pan or wrap the display

Synopsis

```
int fbgen_pan_display (struct fb_var_screeninfo * var, int con,
struct fb_info * info);
```

Arguments

`var`

frame buffer user defined part of display

`con`

virtual console number

`info`

frame buffer info structure

Description

Pan or wrap virtual console *`con`* for device *`info`*.

This call looks only at xoffset, yoffset and the FB_VMODE_YWRAP flag in *`var`*.

Returns negative errno on error, or zero for success.

fbgen_do_set_var

Name

`fbgen_do_set_var` — change the video mode

Synopsis

```
int fbgen_do_set_var (struct fb_var_screeninfo * var, int
isactive, struct fb_info_gen * info);
```

Arguments

`var`

frame buffer user defined part of display

`isactive`

boolean, 0 inactive, 1 active

`info`

generic frame buffer info structure

Description

Change the video mode settings for device *`info`*. If *`isactive`* is non-zero, the changes will be activated immediately.

Return negative errno on error, or zero for success.

fbgen_set_disp

Name

`fbgen_set_disp` — set generic display

Synopsis

```
void fbgen_set_disp (int con, struct fb_info_gen * info);
```

Arguments

`con`

virtual console number

`info`

generic frame buffer info structure

Description

Sets a display on virtual console *`con`* for device *`info`*.

fbgen_install_cmap

Name

`fbgen_install_cmap` — install the current colormap

Synopsis

```
void fbgen_install_cmap (int con, struct fb_info_gen * info);
```

Arguments

con

virtual console number

info

generic frame buffer info structure

Description

Installs the current colormap for virtual console *con* on device *info*.

fbgen_update_var

Name

fbgen_update_var — update user defined part of display

Synopsis

```
int fbgen_update_var (int con, struct fb_info * info);
```

Arguments

con

virtual console number

info

frame buffer info structure

Description

Updates the user defined part of the display (’var’ structure) on virtual console *con* for device *info*. This function is called by fbcon.c.

Returns negative errno on error, or zero for success.

fbgen_switch

Name

fbgen_switch — switch to a different virtual console.

Synopsis

```
int fbgen_switch (int con, struct fb_info * info);
```

Arguments

con

virtual console number

info

frame buffer info structure

Description

Switch to virtuall console *con* on device *info*.

Returns zero.

fbgen_blank

Name

fbgen_blank — blank the screen

Synopsis

```
void fbgen_blank (int blank, struct fb_info * info);
```

Arguments

blank

boolean, 0 unblank, 1 blank

info

frame buffer info structure

Description

Blank the screen on device *info*.

20.5. Frame Buffer Video Mode Database

fb_find_mode

Name

fb_find_mode — finds a valid video mode

Synopsis

```
int __init fb_find_mode (struct fb_var_screeninfo * var, struct
fb_info * info, const char * mode_option, const struct
fb_videomode * db, unsigned int dbsize, const struct
fb_videomode * default_mode, unsigned int default_bpp);
```

Arguments

var

frame buffer user defined part of display

info

frame buffer info structure

mode_option

string video mode to find

db

video mode database

dbsize

size of *db*

default_mode

default video mode to fall back to

default_bpp

default color depth in bits per pixel

Description

Finds a suitable video mode, starting with the specified mode in *mode_option* with fallback to *default_mode*. If *default_mode* fails, all modes in the video mode database will be tried.

Valid mode specifiers for *mode_option*:

<xres>x<yres>[-<bpp>][@<refresh>] or <name>[-<bpp>][@<refresh>]

with <xres>, <yres>, <bpp> and <refresh> decimal numbers and <name> a string.

NOTE

The passed struct *var* is _not_ cleared! This allows you to supply values for e.g. the grayscale and accel_flags fields.

Returns zero for failure, 1 if using specified *mode_option*, 2 if using specified *mode_option* with an ignored refresh rate, 3 if default mode is used, 4 if fall back to any valid mode.

__fb_try_mode

Name

__fb_try_mode — test a video mode

Synopsis

```
int __fb_try_mode (struct fb_var_screeninfo * var, struct
fb_info * info, const struct fb_videomode * mode, unsigned int
bpp);
```

Arguments

var

frame buffer user defined part of display

info

frame buffer info structure

mode

frame buffer video mode structure

bpp

color depth in bits per pixel

Description

Tries a video mode to test it’s validity for device *info*.

Returns 1 on success.

20.6. Frame Buffer Macintosh Video Mode Database

console_getmode

Name

`console_getmode` — get current mode

Synopsis

```
int console_getmode (struct vc_mode * mode);
```

Arguments

`mode`

virtual console mode structure

Description

Populates *`mode`* with the current mode held in the global display_info structure.

Note, this function is only for XPMAC compatibility.

Returns zero.

console_setmode

Name

`console_setmode` — sets current console mode

Synopsis

```
int console_setmode (struct vc_mode * mode, int doit);
```

Arguments

`mode`

virtual console mode structure

`doit`

boolean, 0 test mode, 1 test and activate mode

Description

Sets *`mode`* for all virtual consoles if *`doit`* is non-zero, otherwise, test a mode for validity.

Note, this function is only for XPMAC compatibility.

Returns negative errno on error, or zero for success.

console_setcmap

Name

`console_setcmap` — sets palette color map for console

Synopsis

```
int console_setcmap (int n_entries, unsigned char * red,
unsigned char * green, unsigned char * blue);
```

Arguments

`n_entries`

number of entries in the palette (max 16)

`red`

value for red component of palette

`green`

value for green component of palette

`blue`

value for blue component of palette

Description

Sets global palette_cmap structure and activates the palette on the current console.

Note, this function is only for XPMAC compatibility.

Returns negative errno on error, or zero for success.

console_powermode

Name

`console_powermode` — sets monitor power mode

Synopsis

```
int console_powermode (int mode);
```

Arguments

`mode`

power state to set

Description

Sets power state as dictated by *`mode`*.

Note that this function is only for XPMAC compatibility and doesn't do much.

Returns 0 for `VC_POWERMODE_INQUIRY`, -EINVAL for VESA power settings, or -ENIXIO on failure.

mac_vmode_to_var

Name

`mac_vmode_to_var` — converts vmode/cmode pair to var structure

Synopsis

```
int mac_vmode_to_var (int vmode, int cmode, struct
fb_var_screeninfo * var);
```

Arguments

`vmode`

MacOS video mode

`cmode`

MacOS color mode

`var`

frame buffer video mode structure

Description

Converts a MacOS vmode/cmode pair to a frame buffer video mode structure.

Returns negative errno on error, or zero for success.

mac_var_to_vmode

Name

mac_var_to_vmode — convert var structure to MacOS vmode/cmode pair

Synopsis

```
int mac_var_to_vmode (const struct fb_var_screeninfo * var, int
* vmode, int * cmode);
```

Arguments

var

frame buffer video mode structure

vmode

MacOS video mode

cmode

MacOS color mode

Description

Converts a frame buffer video mode structure to a MacOS vmode/cmode pair.

Returns negative errno on error, or zero for success.

mac_map_monitor_sense

Name

`mac_map_monitor_sense` — Convert monitor sense to vmode

Synopsis

```
int mac_map_monitor_sense (int sense);
```

Arguments

`sense`

Macintosh monitor sense number

Description

Converts a Macintosh monitor sense number to a MacOS vmode number.

Returns MacOS vmode video mode number.

mac_find_mode

Name

`mac_find_mode` — find a video mode

Synopsis

```
int __init mac_find_mode (struct fb_var_screeninfo * var, struct
fb_info * info, const char * mode_option, unsigned int
default_bpp);
```

Arguments

var

frame buffer user defined part of display

info

frame buffer info structure

mode_option

video mode name (see mac_modedb[])

default_bpp

default color depth in bits per pixel

Description

Finds a suitable video mode. Tries to set mode specified by *mode_option*. If the name of the wanted mode begins with 'mac', the Mac video mode database will be used, otherwise it will fall back to the standard video mode database.

Note

Function marked as __init and can only be used during system boot.

Returns error code from fb_find_mode (see fb_find_mode function).

20.7. Frame Buffer Fonts

fbcon_find_font

Name

`fbcon_find_font` — find a font

Synopsis

```
struct fbcon_font_desc * fbcon_find_font (char * name);
```

Arguments

`name`

string name of a font

Description

Find a specified font with string name *`name`*.

Returns `NULL` if no font found, or a pointer to the specified font.

fbcon_get_default_font

Name

`fbcon_get_default_font` — get default font

Synopsis

```
struct fbcon_font_desc * fbcon_get_default_font (int xres, int
yres);
```

Arguments

`xres`

screen size of X

`yres`

screen size of Y

Description

Get the default font for a specified screen size. Dimensions are in pixels.

Returns `NULL` if no font is found, or a pointer to the chosen font.