# The Linux Kernel API

#### The Linux Kernel API

This documentation is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, Inc., 59 Temple Place, Suite 330, Boston, MA 02111-1307 USA For more details see the file COPYING in the source distribution of Linux.

# **Table of Contents**

1. Driver Basics	
1.1. Driver Entry and Exit points	
module_init	
module exit	
1.2. Atomic and pointer manipulation	2
atomic_read	
atomic_set	3
atomic_add	
atomic_sub	5
atomic_sub_and_test	6
atomic_inc	7
atomic_dec	8
atomic_dec_and_test	9
atomic_inc_and_test	10
atomic_add_negative	11
get_unaligned	12
put_unaligned	12
1.3. Delaying, scheduling, and timer routines	13
schedule_timeout	14
2. Data Types	16
2.1. Doubly Linked Lists	
· -	16
2.1. Doubly Linked Lists	16 16
2.1. Doubly Linked Listslist_add	16 16 16
2.1. Doubly Linked Lists list_add list_add_tail	16 16 16 17
2.1. Doubly Linked Lists list_add list_add_tail list_del	16 16 16 17
2.1. Doubly Linked Lists list_add list_add_tail list_del list_del_init	16 16 17 18
2.1. Doubly Linked Lists list_add list_add_tail list_del list_del_init list_empty	16 16 17 18 19
2.1. Doubly Linked Lists list_add list_add_tail list_del list_del_init list_empty. list_splice	
2.1. Doubly Linked Lists  list_add list_add_tail list_del list_del_init list_empty list_splice list_entry	
2.1. Doubly Linked Lists  list_add  list_add_tail  list_del  list_del_init  list_empty  list_splice  list_entry  list_for_each	
2.1. Doubly Linked Lists  list_add list_add_tail list_del list_del_init list_empty list_splice list_entry list_for_each list_for_each_safe	
2.1. Doubly Linked Lists  list_add list_add_tail list_del list_del_init list_empty list_splice list_entry list_for_each list_for_each_safe  3. Basic C Library Functions	
2.1. Doubly Linked Lists  list_add list_add_tail list_del list_del_init list_empty list_splice list_entry list_for_each list_for_each_safe  3. Basic C Library Functions  3.1. String Conversions	
2.1. Doubly Linked Lists list_add list_add_tail list_del list_del_init list_empty list_splice list_entry list_for_each list_for_each_safe  3. Basic C Library Functions  3.1. String Conversions simple_strtol	
2.1. Doubly Linked Lists list_add list_add_tail list_del list_del_init list_empty list_splice list_entry list_for_each list_for_each_safe  3. Basic C Library Functions  3.1. String Conversions simple_strtol simple_strtoll	

	snprintf	28
	vsprintf	29
	sprintf	30
3.2.	String Manipulation	31
	strcpy	31
	strncpy	31
	strcat	32
	strncat	33
	stremp	34
	strncmp	35
	strchr	36
	strrchr	36
	strlen	37
	strnlen	38
	strpbrk	38
	strtok	39
	memset	40
	bcopy	41
	memcpy	42
	memmove	43
	memcmp	43
	memscan	44
	strstr	45
	memchr	46
3.3.	Bit Operations	47
	set_bit	47
	set_bit	48
	clear_bit	49
	change_bit	49
	change_bit	50
	test_and_set_bit	51
	test_and_set_bit	52
	test_and_clear_bit	53
	test_and_clear_bit	54
	test_and_change_bit	55
	test_bit	56
	find_first_zero_bit	56
	find_next_zero_bit	57
	ffz	58
	ffs	59

hweight32	60
4. Memory Management in Linux	61
4.1. The Slab Cache	61
kmem_cache_create	61
kmem_cache_shrink	62
kmem_cache_destroy	63
kmem_cache_alloc	64
kmalloc	65
kmem_cache_free	66
kfree	67
5. The proc filesystem	69
5.1. sysctl interface	
register_sysctl_table	
unregister_sysctl_table	
proc_dostring	
proc_dointvec	73
proc_dointvec_minmax	74
proc_doulongvec_minmax	75
proc_doulongvec_ms_jiffies_minmax	76
proc_dointvec_jiffies	77
6. The Linux VFS	80
6.1. The Directory Cache	80
d_invalidate	
d_find_alias	80
prune_dcache	81
shrink_dcache_sb	82
have_submounts	83
shrink_dcache_parent	84
d_alloc	84
d_instantiate	85
d_alloc_root	86
d_lookup	87
d_validate	88
d_delete	89
d_rehash	
d_move	
d_path	
is_subdir	
find_inode_number	94

d_drop	94
d_add	95
dget	96
d_unhashed	97
6.2. Inode Handling	98
mark_inode_dirty	98
write_inode_now	99
clear_inode	100
invalidate_inodes	100
get_empty_inode	101
iunique	102
insert_inode_hash	103
remove_inode_hash	104
iput	105
bmap	105
update_atime	106
make_bad_inode	107
is_bad_inode	108
6.3. Registration and Superblocks	109
register_filesystem	109
unregister_filesystem	110
get_super	111
6.4. File Locks	111
posix_lock_file	112
get_lease	113
lease_get_mtime	114
posix_block_lock	114
posix_unblock_lock	115
lock_may_read	116
lock_may_write	117
fcntl_getlease	118
fcntl_setlease	119
sys_flock	120
get_locks_status	121
7. Linux Networking	123
7.1. Socket Buffer Functions	123
skb_queue_empty	
skb_get	
kfree_skb	
skb_cloned	

skb_shared	126
skb_share_check	127
skb_unshare	128
skb_peek	129
skb_peek_tail	129
skb_queue_len	130
skb_queue_head	131
skb_queue_head	132
skb_queue_tail	133
skb_queue_tail	134
skb_dequeue	135
skb_dequeue	136
skb_insert	137
skb_append	138
skb_unlink	138
skb_dequeue_tail	139
skb_dequeue_tail	140
skb_put	141
skb_push	142
skb_pull	143
skb_headroom	144
skb_tailroom	144
skb_reserve	145
skb_trim	146
skb_orphan	147
skb_queue_purge	148
skb_queue_purge	148
dev_alloc_skb	149
dev_alloc_skb	150
skb_cow	151
skb_over_panic	152
skb_under_panic	153
alloc_skb	154
kfree_skb	155
skb_clone	156
skb_copy	157
pskb_copy	158
pskb_expand_head	159
skb_copy_expand	160
pskb_pull_tail	161

7.2. S	Socket Filter	162
:	sk_run_filter	162
:	sk_chk_filter	163
8. Networ	k device support	165
	Oriver Support	
	init_etherdev	
	alloc_etherdev	
	init_fddidev	167
;	alloc_fddidev	168
į	init_hippi_dev	168
;	alloc_hippi_dev	169
i	init_trdev	170
;	alloc_trdev	171
1	init_fcdev	172
;	alloc_fcdev	173
(	dev_add_pack	174
(	dev_remove_pack	175
-	dev_get_by_name	175
	dev_get_by_name	176
	dev_get	177
-	dev_get_by_index	178
	dev_get_by_index	179
(	dev_alloc_name	179
	dev_alloc	
1	netdev_state_change	181
(	dev_load	182
(	dev_open	183
(	dev_close	184
1	register_netdevice_notifier	185
1	unregister_netdevice_notifier	185
(	dev_queue_xmit	186
1	netif_rx	187
1	net_call_rx_atomic	188
1	register_gifconf	189
1	netdev_set_master	190
	dev_set_promiscuity	
	dev_set_allmulti	192
	dev_ioctl	192
	dev_new_index	193
1	netdev_finish_unregister	194

	unregister_netdevice	195
	8.2. 8390 Based Network Cards	196
	ei_open	196
	ei_close	197
	ei_tx_timeout	198
	ei_interrupt	198
	ethdev_init	199
	NS8390_init	200
	8.3. Synchronous PPP	201
	sppp_input	201
	sppp_close	202
	sppp_open	203
	sppp_reopen	204
	sppp_change_mtu	205
	sppp_do_ioctl	206
	sppp_attach	207
	sppp_detach	207
9. N	Module Support	209
	9.1. Module Loading	209
	request_module	
	call_usermodehelper	
	9.2. Inter Module support	
	inter_module_register	
	inter_module_unregister	
	inter_module_get	212
	inter_module_get_request	213
	inter_module_put	
<b>10.</b> ]	Hardware Interfaces	216
	10.1. Interrupt Handling	216
	disable_irq_nosync	
	disable_irq	
	enable_irq	
	probe_irq_mask	
	10.2. MTRR Handling	
	mtrr_add	
	mtrr_del	
	10.3. PCI Support Library	
	pci_find_slot	
	pci_find_subsys	
	~ ~ ~ , ~	

pci_find_device	224
pci_find_class	225
pci_find_capability	226
pci_find_parent_resource	227
pci_set_power_state	228
pci_save_state	229
pci_restore_state	230
pci_enable_device	231
pci_disable_device	231
pci_enable_wake	232
pci_release_regions	233
pci_request_regions	234
pci_match_device	235
pci_register_driver	236
pci_unregister_driver	237
pci_insert_device	237
pci_remove_device	238
pci_dev_driver	239
pci_set_master	240
pci_setup_device	241
pci_pool_create	241
pci_pool_destroy	243
pci_pool_alloc	244
pci_pool_free	245
10.4. MCA Architecture	246
10.4.1. MCA Device Functions	246
mca_find_adapter	246
mca_find_unused_adapter	247
mca_read_stored_pos	248
mca_read_pos	249
mca_write_pos	249
mca_set_adapter_name	251
mca_set_adapter_procfn	251
mca_is_adapter_used	253
mca_mark_as_used	253
mca_mark_as_unused	254
mca_get_adapter_name	255
mca_isadapter	256
mca_isenabled	256
10.4.2. MCA Bus DMA	257

mca_	_enable_dma	258
mca_	_disable_dma	258
mca_	_set_dma_addr	259
mca_	_get_dma_addr	260
mca_	_set_dma_count	261
mca_	_get_dma_residue	262
mca_	_set_dma_io	262
mca_	_set_dma_mode	263
11. The Device File	e System	265
devfs_register		265
•	er	
•	llink	
devfs_mk_dir		268
devfs_find_han	ndle	269
devfs_get_flags	S	270
	_min	
devfs_get_hand	dle_from_inode	272
devfs_generate	_path	273
devfs_get_ops .		274
devfs_set_file_	size	275
devfs_get_info		275
devfs_set_info		276
devfs_get_pare	ent	277
devfs_get_first_	_child	278
devfs_get_next	_sibling	279
devfs_auto_unr	register	279
devfs_get_unre	egister_slave	280
devfs_get_nam	ıe	281
devfs_register_	chrdev	282
	blkdev	
devfs_unregiste	er_chrdev	284
devfs_unregiste	er_blkdev	285
12. Power Manager	ment	286
_		
	_all	
-		
-		
-		

13. Block Devices	293
blk_cleanup_queue	293
blk_queue_headactive	294
blk_queue_make_request	295
blk_init_queue	296
generic_make_request	297
submit_bh	298
ll_rw_block	299
end_that_request_first	301
14. Miscellaneous Devices	303
misc_register	303
misc_deregister	303
15. Video4Linux	305
video_unregister_device	305
16. Sound Devices	306
register_sound_special	306
register_sound_mixer	
register_sound_midi	307
register_sound_dsp	308
register_sound_synth	309
unregister_sound_special	310
unregister_sound_mixer	311
unregister_sound_midi	312
unregister_sound_dsp	313
unregister_sound_synth	314
17. USB Devices	315
usb_register	315
usb_scan_devices	315
usb_deregister	316
usb_alloc_bus	317
usb_free_bus	318
usb_register_bus	319
usb_deregister_bus	319
usb_match_id	320
usb_alloc_urb	322
usb_free_urb	
usb_control_msg	324
usb_bulk_msg	325

18. 16x50 UART Driver	328
register_serial	328
unregister_serial	328
19. Z85230 Support Library	330
z8530_interrupt	
z8530_sync_open	
z8530_sync_close	
z8530_sync_dma_open	332
z8530_sync_dma_close	333
z8530_sync_txdma_open	334
z8530_sync_txdma_close	335
z8530_describe	336
z8530_init	337
z8530_shutdown	338
z8530_channel_load	339
z8530_null_rx	339
z8530_queue_xmit	340
z8530_get_stats	341
20. Frame Buffer Library	343
20.1. Frame Buffer Memory	343
register_framebuffer	343
unregister_framebuffer	
20.2. Frame Buffer Console	345
fbcon_redraw_clear	345
fbcon_redraw_bmove	346
20.3. Frame Buffer Colormap	347
fb_alloc_cmap	348
fb_copy_cmap	348
fb_get_cmap	350
fb_set_cmap	351
fb_default_cmap	352
fb_invert_cmaps	
20.4. Frame Buffer Generic Functions	353
fbgen_get_fix	
fbgen_get_var	
fbgen_set_var	
fbgen_get_cmap	
fbgen_set_cmap	
fbgen_pan_display	358

fbgen_do_set_var	359
fbgen_set_disp	360
fbgen_install_cmap	361
fbgen_update_var	362
fbgen_switch	
fbgen_blank	364
20.5. Frame Buffer Video Mode Database	
fb_find_mode	365
fb_try_mode	366
20.6. Frame Buffer Macintosh Video Mode Database	368
console_getmode	368
console_setmode	
console_setcmap	369
console_powermode	
mac_vmode_to_var	371
mac_var_to_vmode	
mac_map_monitor_sense	373
mac_find_mode	
20.7. Frame Buffer Fonts	
fbcon_find_font	376
fbcon_get_default_font	
<del>_</del>	

# **Chapter 1. Driver Basics**

# 1.1. Driver Entry and Exit points

## module\_init

#### **Name**

module\_init — driver initialization entry point

## **Synopsis**

```
module_init ( x);
```

### **Arguments**

X

function to be run at kernel boot time or module insertion

### **Description**

module\_init will add the driver initialization routine in the "\_\_initcall.int" code segment if the driver is checked as "y" or static, or else it will wrap the driver initialization routine with init\_module which is used by insmod and modprobe when the driver is used as a module.

## module\_exit

#### Name

module\_exit — driver exit entry point

## **Synopsis**

```
module_exit ( x);
```

## **Arguments**

 $\boldsymbol{X}$ 

function to be run when driver is removed

### **Description**

module\_exit will wrap the driver clean-up code with cleanup\_module when used with rmmod when the driver is a module. If the driver is statically compiled into the kernel, module\_exit has no effect.

## 1.2. Atomic and pointer manipulation

## atomic\_read

#### **Name**

atomic\_read — read atomic variable

## **Synopsis**

```
atomic_read ( v);
```

### **Arguments**

v pointer of type atomic\_t

## **Description**

Atomically reads the value of v. Note that the guaranteed useful range of an atomic\_t is only 24 bits.

## atomic\_set

#### Name

```
atomic_set — set atomic variable
```

## **Synopsis**

```
atomic_set ( v , i);
```

## **Arguments**

```
v pointer of type atomic_t i required value
```

## **Description**

Atomically sets the value of v to i. Note that the guaranteed useful range of an atomic\_t is only 24 bits.

## atomic\_add

#### Name

atomic\_add — add integer to atomic variable

## **Synopsis**

```
void atomic_add (int i, atomic_t * v);
```

## **Arguments**

```
i integer value to add

v

pointer of type atomic_t
```

## **Description**

Atomically adds i to v. Note that the guaranteed useful range of an atomic\_t is only 24 bits.

## atomic\_sub

#### Name

atomic\_sub — subtract the atomic variable

## **Synopsis**

```
void atomic_sub (int i, atomic_t * v);
```

## **Arguments**

```
i integer value to subtract

v

pointer of type atomic_t
```

## **Description**

Atomically subtracts  $\dot{\imath}$  from v. Note that the guaranteed useful range of an atomic\_t is only 24 bits.

## atomic\_sub\_and\_test

#### Name

atomic\_sub\_and\_test — subtract value from variable and test result

## **Synopsis**

```
int atomic_sub_and_test (int i, atomic_t * v);
```

### **Arguments**

```
i integer value to subtract

v

pointer of type atomic_t
```

## **Description**

Atomically subtracts i from v and returns true if the result is zero, or false for all other cases. Note that the guaranteed useful range of an atomic\_t is only 24 bits.

## atomic\_inc

#### Name

atomic\_inc — increment atomic variable

### **Synopsis**

```
void atomic_inc (atomic_t * v);
```

### **Arguments**

V

pointer of type atomic\_t

### **Description**

Atomically increments v by 1. Note that the guaranteed useful range of an atomic\_t is only 24 bits.

## atomic\_dec

#### Name

atomic\_dec — decrement atomic variable

## **Synopsis**

```
void atomic_dec (atomic_t * v);
```

### **Arguments**

v pointer of type atomic\_t

## **Description**

Atomically decrements v by 1. Note that the guaranteed useful range of an atomic\_t is only 24 bits.

## atomic\_dec\_and\_test

#### Name

atomic\_dec\_and\_test — decrement and test

## **Synopsis**

```
int atomic_dec_and_test (atomic_t * v);
```

## **Arguments**

v pointer of type atomic\_t

## **Description**

Atomically decrements v by 1 and returns true if the result is 0, or false for all other cases. Note that the guaranteed useful range of an atomic\_t is only 24 bits.

## atomic\_inc\_and\_test

#### Name

```
atomic_inc_and_test — increment and test
```

## **Synopsis**

```
int atomic_inc_and_test (atomic_t * v);
```

### **Arguments**

v pointer of type atomic\_t

### **Description**

Atomically increments v by 1 and returns true if the result is zero, or false for all other cases. Note that the guaranteed useful range of an atomic\_t is only 24 bits.

## atomic\_add\_negative

#### **Name**

```
atomic_add_negative — add and test if negative
```

### **Synopsis**

```
int atomic_add_negative (int i, atomic_t * v);
```

### **Arguments**

```
i integer value to addvpointer of type atomic_t
```

## **Description**

Atomically adds i to v and returns true if the result is negative, or false when result is greater than or equal to zero. Note that the guaranteed useful range of an atomic\_t is only 24 bits.

## get\_unaligned

#### **Name**

get\_unaligned — get value from possibly mis-aligned location

## **Synopsis**

```
get_unaligned ( ptr);
```

### **Arguments**

ptr

pointer to value

### **Description**

This macro should be used for accessing values larger in size than single bytes at locations that are expected to be improperly aligned, e.g. retrieving a u16 value from a location not u16-aligned.

Note that unaligned accesses can be very expensive on some architectures.

## put\_unaligned

#### Name

put\_unaligned — put value to a possibly mis-aligned location

## **Synopsis**

```
put_unaligned ( val, ptr);
```

### **Arguments**

```
val

value to place

ptr

pointer to location
```

## **Description**

This macro should be used for placing values larger in size than single bytes at locations that are expected to be improperly aligned, e.g. writing a u16 value to a location not u16-aligned.

Note that unaligned accesses can be very expensive on some architectures.

## 1.3. Delaying, scheduling, and timer routines

## schedule\_timeout

#### **Name**

schedule\_timeout — sleep until timeout

### **Synopsis**

signed long schedule\_timeout (signed long timeout);

### **Arguments**

timeout

timeout value in jiffies

### **Description**

Make the current task sleep until timeout jiffies have elapsed. The routine will return immediately unless the current task state has been set (see set\_current\_state).

You can set the task state as follows -

TASK\_UNINTERRUPTIBLE - at least timeout jiffies are guaranteed to pass before the routine returns. The routine will return 0

TASK\_INTERRUPTIBLE - the routine may return early if a signal is delivered to the current task. In this case the remaining time in jiffies will be returned, or 0 if the timer expired in time

The current task state is guaranteed to be TASK\_RUNNING when this routine returns.

Specifying a timeout value of MAX\_SCHEDULE\_TIMEOUT will schedule the CPU away without a bound on the timeout. In this case the return value will be MAX\_SCHEDULE\_TIMEOUT.

In all cases the return value is guaranteed to be non-negative.

# **Chapter 2. Data Types**

## 2.1. Doubly Linked Lists

## list\_add

#### Name

list\_add — add a new entry

## **Synopsis**

```
void list_add (struct list_head * new, struct list_head * head);
```

## **Arguments**

```
new entry to be added

head

list head to add it after
```

## **Description**

Insert a new entry after the specified head. This is good for implementing stacks.

## list\_add\_tail

#### Name

list\_add\_tail — add a new entry

## **Synopsis**

```
void list_add_tail (struct list_head * new, struct list_head *
head);
```

### **Arguments**

```
new
```

new entry to be added

head

list head to add it before

## **Description**

Insert a new entry before the specified head. This is useful for implementing queues.

## list\_del

#### Name

list\_del — deletes entry from list.

## **Synopsis**

```
void list_del (struct list_head * entry);
```

### **Arguments**

entry

the element to delete from the list.

#### **Note**

list\_empty on entry does not return true after this, the entry is in an undefined state.

## list\_del\_init

### Name

list\_del\_init — deletes entry from list and reinitialize it.

## **Synopsis**

```
void list_del_init (struct list_head * entry);
```

## **Arguments**

entry

the element to delete from the list.

# list\_empty

### Name

list\_empty — tests whether a list is empty

## **Synopsis**

```
int list_empty (struct list_head * head);
```

## **Arguments**

head

the list to test.

## list\_splice

#### Name

list\_splice — join two lists

## **Synopsis**

```
void list_splice (struct list_head * list, struct list_head *
head);
```

## **Arguments**

list

the new list to add.

head

the place to add it in the first list.

## list\_entry

#### Name

list\_entry — get the struct for this entry

## **Synopsis**

```
list_entry ( ptr, type, member);
```

## **Arguments**

```
the &struct list_head pointer.

type
the type of the struct this is embedded in.

member
the name of the list_struct within the struct.
```

## list\_for\_each

#### **Name**

```
list_for_each — iterate over a list
```

## **Synopsis**

```
list_for_each ( pos, head);
```

### **Arguments**

```
the &struct list_head to use as a loop counter.

head
the head for your list.
```

## list\_for\_each\_safe

#### **Name**

list\_for\_each\_safe — iterate over a list safe against removal of list entry

## **Synopsis**

```
list_for_each_safe ( pos,  n,  head);
```

### **Arguments**

```
the &struct list_head to use as a loop counter.

n
another &struct list_head to use as temporary storage

head
the head for your list.
```

# **Chapter 3. Basic C Library Functions**

When writing drivers, you cannot in general use routines which are from the C Library. Some of the functions have been found generally useful and they are listed below. The behaviour of these functions may vary slightly from those defined by ANSI, and these deviations are noted in the text.

## 3.1. String Conversions

## simple\_strtol

#### **Name**

simple\_strtol — convert a string to a signed long

## **Synopsis**

```
long simple_strtol (const char * cp, char ** endp, unsigned int
base);
```

### **Arguments**

СР

The start of the string

endp

A pointer to the end of the parsed string will be placed here

base

The number base to use

# simple\_strtoll

#### **Name**

simple\_strtoll — convert a string to a signed long long

## **Synopsis**

```
long long simple_strtoll (const char * cp, char ** endp,
unsigned int base);
```

## **Arguments**

ср

The start of the string

endp

A pointer to the end of the parsed string will be placed here

base

The number base to use

# simple\_strtoul

#### Name

simple\_strtoul — convert a string to an unsigned long

## **Synopsis**

```
unsigned long simple_strtoul (const char * cp, char ** endp, unsigned int base);
```

## **Arguments**

ср

The start of the string

endp

A pointer to the end of the parsed string will be placed here

base

The number base to use

# simple\_strtoull

#### Name

simple\_strtoull — convert a string to an unsigned long long

## **Synopsis**

```
unsigned long long simple_strtoull (const char * cp, char ** endp, unsigned int base);
```

## **Arguments**

ср

The start of the string

endp

A pointer to the end of the parsed string will be placed here

base

The number base to use

# vsnprintf

### **Name**

vsnprintf — Format a string and place it in a buffer

```
int vsnprintf (char * buf, size_t size, const char * fmt,
va_list args);
```

### **Arguments**

```
The buffer to place the result into

size

The size of the buffer, including the trailing null space

fmt

The format string to use

args
```

Arguments for the format string

## **Description**

Call this function if you are already dealing with a va\_list. You probably want snprintf instead.

# **snprintf**

### **Name**

snprintf — Format a string and place it in a buffer

```
int snprintf (char * buf, size_t size, const char * fmt, ...
...);
```

### **Arguments**

```
The buffer to place the result into

size

The size of the buffer, including the trailing null space

fmt

The format string to use @...: Arguments for the format string

variable arguments
```

# vsprintf

#### Name

vsprintf — Format a string and place it in a buffer

## **Synopsis**

```
int vsprintf (char * buf, const char * fmt, va_list args);
```

## **Arguments**

buf

The buffer to place the result into

fmt

The format string to use

args

Arguments for the format string

## **Description**

Call this function if you are already dealing with a va\_list. You probably want sprintf instead.

# sprintf

#### Name

sprintf — Format a string and place it in a buffer

## **Synopsis**

```
int sprintf (char * buf, const char * fmt, ...);
```

## **Arguments**

buf

The buffer to place the result into

fmt

The format string to use @...: Arguments for the format string

. . .

variable arguments

# 3.2. String Manipulation

## strcpy

### **Name**

strcpy — Copy a NUL terminated string

# **Synopsis**

```
char * strcpy (char * dest, const char * src);
```

## **Arguments**

dest

Where to copy the string to

src

Where to copy the string from

## strncpy

#### **Name**

```
strncpy — Copy a length-limited, NUL-terminated string
```

## **Synopsis**

```
char * strncpy (char * dest, const char * src, size_t count);
```

## **Arguments**

dest

Where to copy the string to

src

Where to copy the string from

count

The maximum number of bytes to copy

## **Description**

Note that unlike userspace strncpy, this does not NUL-pad the buffer. However, the result is not NUL-terminated if the source exceeds *count* bytes.

### strcat

#### Name

strcat — Append one NUL-terminated string to another

## **Synopsis**

```
char * strcat (char * dest, const char * src);
```

## **Arguments**

dest

The string to be appended to

src

The string to append to it

## strncat

### Name

strncat — Append a length-limited, NUL-terminated string to another

```
char * strncat (char * dest, const char * src, size_t count);
```

## **Arguments**

```
The string to be appended to

src

The string to append to it

count
```

The maximum numbers of bytes to copy

## **Description**

Note that in contrast to strncpy, strncat ensures the result is terminated.

## strcmp

#### Name

```
\verb|strcmp| - Compare two strings|
```

```
int strcmp (const char * cs, const char * ct);
```

## **Arguments**

```
One string

Ct

Another string
```

## strncmp

### Name

strncmp — Compare two length-limited strings

# **Synopsis**

```
int strncmp (const char * cs, const char * ct, size_t count);
```

## **Arguments**

```
One string

Ct

Another string

Count

The maximum number of bytes to compare
```

## strchr

#### **Name**

strchr — Find the first occurrence of a character in a string

## **Synopsis**

```
char * strchr (const char * s, int c);
```

## **Arguments**

s

The string to be searched

C

The character to search for

## strrchr

### **Name**

strrchr — Find the last occurrence of a character in a string

## **Synopsis**

```
char * strrchr (const char * s, int c);
```

## **Arguments**

S

The string to be searched

C

The character to search for

## strlen

#### **Name**

strlen — Find the length of a string

## **Synopsis**

```
size_t strlen (const char * s);
```

## **Arguments**

S

The string to be sized

## strnlen

#### Name

strnlen — Find the length of a length-limited string

## **Synopsis**

```
size_t strnlen (const char * s, size_t count);
```

## **Arguments**

S

The string to be sized

count

The maximum number of bytes to search

# strpbrk

#### **Name**

strpbrk — Find the first occurrence of a set of characters

## **Synopsis**

```
char * strpbrk (const char * cs, const char * ct);
```

## **Arguments**

CS

The string to be searched

ct

The characters to search for

## strtok

#### Name

strtok — Split a string into tokens

## **Synopsis**

```
char * strtok (char * s, const char * ct);
```

## **Arguments**

S

The string to be searched

ct

The characters to search for

#### **WARNING**

strtok is deprecated, use strsep instead.

#### memset

#### Name

memset — Fill a region of memory with the given value

## **Synopsis**

```
void * memset (void * s, int c, size_t count);
```

## **Arguments**

S

Pointer to the start of the area.

C

The byte to fill the area with

count

The size of the area.

## **Description**

Do not use memset to access IO space, use memset\_io instead.

## bcopy

#### Name

bcopy — Copy one area of memory to another

## **Synopsis**

```
char * bcopy (const char * src, char * dest, int count);
```

### **Arguments**

```
src
```

Where to copy from

dest

Where to copy to

count

The size of the area.

## **Description**

Note that this is the same as memcpy, with the arguments reversed. memcpy is the standard, bcopy is a legacy BSD function.

You should not use this function to access IO space, use memcpy\_toio or memcpy\_fromio instead.

## memcpy

#### Name

memcpy — Copy one area of memory to another

## **Synopsis**

```
void * memcpy (void * dest, const void * src, size_t count);
```

## **Arguments**

```
dest
Where to copy to

src
Where to copy from

count
The size of the area.
```

## **Description**

You should not use this function to access IO space, use memcpy\_toio or memcpy\_fromio instead.

### memmove

#### Name

memmove — Copy one area of memory to another

## **Synopsis**

```
void * memmove (void * dest, const void * src, size_t count);
```

## **Arguments**

dest

Where to copy to

src

Where to copy from

count

The size of the area.

## **Description**

Unlike memcpy, memmove copes with overlapping areas.

## memcmp

#### **Name**

memcmp — Compare two areas of memory

# **Synopsis**

```
int memcmp (const void * cs, const void * ct, size_t count);
```

## **Arguments**

CS

One area of memory

ct

Another area of memory

count

The size of the area.

### memscan

### **Name**

memscan — Find a character in an area of memory.

## **Synopsis**

```
void * memscan (void * addr, int c, size_t size);
```

### **Arguments**

```
addr
The memory area

C
The byte to search for

size
The size of the area.
```

## **Description**

returns the address of the first occurrence of c, or 1 byte past the area if c is not found

## strstr

#### Name

strstr — Find the first substring in a NUL terminated string

```
char * strstr (const char * s1, const char * s2);
```

## **Arguments**

s1

The string to be searched

s2

The string to search for

## memchr

#### Name

memchr — Find a character in an area of memory.

## **Synopsis**

```
void * memchr (const void * s, int c, size_t n);
```

## **Arguments**

S

The memory area

C

The byte to search for

n

The size of the area.

## **Description**

returns the address of the first occurrence of c, or NULL if c is not found

# 3.3. Bit Operations

# set\_bit

#### Name

set\_bit — Atomically set a bit in memory

## **Synopsis**

```
void set_bit (int nr, volatile void * addr);
```

## **Arguments**

```
nr the bit to set addr the address to start counting from
```

### **Description**

This function is atomic and may not be reordered. See \_\_set\_bit if you do not require the atomic guarantees. Note that nr may be almost arbitrarily large; this function is not restricted to acting on a single-word quantity.

## set bit

#### **Name**

```
__set_bit — Set a bit in memory
```

## **Synopsis**

```
void __set_bit (int nr, volatile void * addr);
```

### **Arguments**

```
nr the bit to set addr the address to start counting from
```

## **Description**

Unlike set\_bit, this function is non-atomic and may be reordered. If it's called on the same region of memory simultaneously, the effect may be that only one operation succeeds.

## clear\_bit

#### **Name**

clear\_bit — Clears a bit in memory

## **Synopsis**

```
void clear_bit (int nr, volatile void * addr);
```

### **Arguments**

nr

Bit to clear

addr

Address to start counting from

## **Description**

clear\_bit is atomic and may not be reordered. However, it does not contain a memory barrier, so if it is used for locking purposes, you should call smp\_mb\_\_before\_clear\_bit and/or smp\_mb\_\_after\_clear\_bit in order to ensure changes are visible on other processors.

## \_\_change\_bit

#### Name

```
__change_bit — Toggle a bit in memory
```

## **Synopsis**

```
void __change_bit (int nr, volatile void * addr);
```

## **Arguments**

```
nr the bit to set addr the address to start counting from
```

## **Description**

Unlike change\_bit, this function is non-atomic and may be reordered. If it's called on the same region of memory simultaneously, the effect may be that only one operation succeeds.

# change\_bit

#### Name

change\_bit — Toggle a bit in memory

## **Synopsis**

```
void change_bit (int nr, volatile void * addr);
```

## **Arguments**

nr

Bit to clear

addr

Address to start counting from

## **Description**

change\_bit is atomic and may not be reordered. Note that nr may be almost arbitrarily large; this function is not restricted to acting on a single-word quantity.

## test\_and\_set\_bit

#### Name

test\_and\_set\_bit — Set a bit and return its old value

## **Synopsis**

```
int test_and_set_bit (int nr, volatile void * addr);
```

## **Arguments**

nr

Bit to set

addr

Address to count from

## **Description**

This operation is atomic and cannot be reordered. It also implies a memory barrier.

## \_\_test\_and\_set\_bit

#### **Name**

\_\_test\_and\_set\_bit — Set a bit and return its old value

## **Synopsis**

```
int __test_and_set_bit (int nr, volatile void * addr);
```

### **Arguments**

```
nr
Bit to set

addr

Address to count from
```

## **Description**

This operation is non-atomic and can be reordered. If two examples of this operation race, one can appear to succeed but actually fail. You must protect multiple accesses with a lock.

## test\_and\_clear\_bit

#### Name

test\_and\_clear\_bit — Clear a bit and return its old value

```
int test_and_clear_bit (int nr, volatile void * addr);
```

### **Arguments**

```
nr
Bit to set

addr

Address to count from
```

## **Description**

This operation is atomic and cannot be reordered. It also implies a memory barrier.

# \_\_test\_and\_clear\_bit

### **Name**

```
__test_and_clear_bit — Clear a bit and return its old value
```

## **Synopsis**

```
int __test_and_clear_bit (int nr, volatile void * addr);
```

## **Arguments**

nr

Bit to set

addr

Address to count from

## **Description**

This operation is non-atomic and can be reordered. If two examples of this operation race, one can appear to succeed but actually fail. You must protect multiple accesses with a lock.

# test\_and\_change\_bit

#### **Name**

test\_and\_change\_bit — Change a bit and return its new value

## **Synopsis**

```
int test_and_change_bit (int nr, volatile void * addr);
```

## **Arguments**

nr

Bit to set

addr

Address to count from

## **Description**

This operation is atomic and cannot be reordered. It also implies a memory barrier.

# test\_bit

## Name

test\_bit — Determine whether a bit is set

## **Synopsis**

```
int test_bit (int nr, const volatile void * addr);
```

## **Arguments**

nr

bit number to test

addr

Address to start counting from

## find\_first\_zero\_bit

#### Name

find\_first\_zero\_bit — find the first zero bit in a memory region

### **Synopsis**

```
int find_first_zero_bit (void * addr, unsigned size);
```

### **Arguments**

addr

The address to start the search at

size

The maximum size to search

### **Description**

Returns the bit-number of the first zero bit, not the number of the byte containing a bit.

## find\_next\_zero\_bit

#### **Name**

find\_next\_zero\_bit — find the first zero bit in a memory region

## **Synopsis**

```
int find_next_zero_bit (void * addr, int size, int offset);
```

## **Arguments**

addr

The address to base the search on

size

The maximum size to search

offset

The bitnumber to start searching at

## ffz

#### **Name**

ffz — find first zero in word.

```
unsigned long ffz (unsigned long word);
```

word

The word to search

## **Description**

Undefined if no zero exists, so code should check against ~0UL first.

# ffs

#### **Name**

ffs — find first bit set

# **Synopsis**

```
int ffs (int x);
```

## **Arguments**

 $\boldsymbol{X}$ 

the word to search

## **Description**

This is defined the same way as the libc and compiler builtin ffs routines, therefore differs in spirit from the above ffz (man ffs).

# hweight32

#### **Name**

hweight 32 — returns the hamming weight of a N-bit word

## **Synopsis**

```
hweight32 (x);
```

## **Arguments**

 $\boldsymbol{X}$ 

the word to weigh

## **Description**

The Hamming Weight of a number is the total number of bits set in it.

# Chapter 4. Memory Management in Linux

#### 4.1. The Slab Cache

## kmem\_cache\_create

#### Name

kmem cache create — Create a cache.

## **Synopsis**

```
kmem_cache_t * kmem_cache_create (const char * name, size_t
size, size_t offset, unsigned long flags, void (*ctor) (void*,
kmem_cache_t *, unsigned long), void (*dtor) (void*,
kmem_cache_t *, unsigned long));
```

#### **Arguments**

name

A string which is used in /proc/slabinfo to identify this cache.

size

The size of objects to be created in this cache.

offset

The offset to use within the page.

flags

SLAB flags

ctor

A constructor for the objects.

dtor

A destructor for the objects.

#### **Description**

Returns a ptr to the cache on success, NULL on failure. Cannot be called within a int, but can be interrupted. The ctor is run when new pages are allocated by the cache and the dtor is run before the pages are handed back. The flags are

SLAB\_POISON - Poison the slab with a known test pattern (a5a5a5a5) to catch references to uninitialised memory.

SLAB\_RED\_ZONE - Insert 'Red' zones around the allocated memory to check for buffer overruns.

SLAB\_NO\_REAP - Don't automatically reap this cache when we're under memory pressure.

SLAB\_HWCACHE\_ALIGN - Align the objects in this cache to a hardware cacheline. This can be beneficial if you're counting cycles as closely as davem.

## kmem\_cache\_shrink

#### **Name**

kmem\_cache\_shrink — Shrink a cache.

## **Synopsis**

```
int kmem_cache_shrink (kmem_cache_t * cachep);
```

#### **Arguments**

cachep

The cache to shrink.

## **Description**

Releases as many slabs as possible for a cache. To help debugging, a zero exit status indicates all slabs were released.

# kmem\_cache\_destroy

#### Name

kmem\_cache\_destroy — delete a cache

## **Synopsis**

```
int kmem_cache_destroy (kmem_cache_t * cachep);
```

cachep

the cache to destroy

## **Description**

Remove a kmem\_cache\_t object from the slab cache. Returns 0 on success.

It is expected this function will be called by a module when it is unloaded. This will remove the cache completely, and avoid a duplicate cache being allocated each time a module is loaded and unloaded, if the module doesn't have persistent in-kernel storage across loads and unloads.

The caller must guarantee that noone will allocate memory from the cache during the kmem cache destroy.

## kmem\_cache\_alloc

#### **Name**

kmem\_cache\_alloc — Allocate an object

## **Synopsis**

```
void * kmem_cache_alloc (kmem_cache_t * cachep, int flags);
```

```
cachep
```

The cache to allocate from.

flags

See kmalloc.

## **Description**

Allocate an object from this cache. The flags are only relevant if the cache has no available objects.

## kmalloc

#### **Name**

kmalloc — allocate memory

## **Synopsis**

```
void * kmalloc (size_t size, int flags);
```

## **Arguments**

size

how many bytes of memory are required.

flags

the type of memory to allocate.

#### **Description**

kmalloc is the normal method of allocating memory in the kernel.

The flags argument may be one of:

GFP\_USER - Allocate memory on behalf of user. May sleep.

GFP\_KERNEL - Allocate normal kernel ram. May sleep.

GFP\_ATOMIC - Allocation will not sleep. Use inside interrupt handlers.

Additionally, the GFP\_DMA flag may be set to indicate the memory must be suitable for DMA. This can mean different things on different platforms. For example, on i386, it means that the memory must come from the first 16MB.

## kmem\_cache\_free

#### Name

kmem\_cache\_free — Deallocate an object

#### **Synopsis**

```
void kmem_cache_free (kmem_cache_t * cachep, void * objp);
```

```
cachep
```

The cache the allocation was from.

objp

The previously allocated object.

## **Description**

Free an object which was previously allocated from this cache.

## kfree

#### **Name**

kfree — free previously allocated memory

## **Synopsis**

```
void kfree (const void * objp);
```

## **Arguments**

```
objp
```

pointer returned by kmalloc.

# **Description**

Don't free memory not originally allocated by kmalloc or you will run into trouble.

# Chapter 5. The proc filesystem

## 5.1. sysctl interface

# register\_sysctl\_table

#### Name

register\_sysctl\_table — register a sysctl heirarchy

## **Synopsis**

```
struct ctl_table_header * register_sysctl_table (ctl_table *
table, int insert_at_head);
```

#### **Arguments**

```
table
```

the top-level table structure

```
insert_at_head
```

whether the entry should be inserted in front or at the end

#### **Description**

Register a sysctl table heirarchy. table should be a filled in ctl\_table array. An entry with a ctl\_name of 0 terminates the table.

The members of the &ctl\_table structure are used as follows:

ctl\_name - This is the numeric sysctl value used by sysctl(2). The number must be unique within that level of sysctl

procname - the name of the sysctl file under /proc/sys. Set to NULL to not enter a sysctl file

data - a pointer to data for use by proc\_handler

maxlen - the maximum size in bytes of the data

mode - the file permissions for the /proc/sys file, and for sysctl(2)

child - a pointer to the child sysctl table if this entry is a directory, or NULL.

proc\_handler - the text handler routine (described below)

strategy - the strategy routine (described below)

de - for internal use by the sysctl routines

extra1, extra2 - extra pointers usable by the proc handler routines

Leaf nodes in the sysctl tree will be represented by a single file under /proc; non-leaf nodes will be represented by directories.

sysctl(2) can automatically manage read and write requests through the sysctl table. The data and maxlen fields of the ctl\_table struct enable minimal validation of the values being written to be performed, and the mode field allows minimal authentication.

More sophisticated management can be enabled by the provision of a strategy routine with the table entry. This will be called before any automatic read or write of the data is performed.

The strategy routine may return

- < 0 Error occurred (error is passed to user process)
- 0 OK proceed with automatic read or write.
- > 0 OK read or write has been done by the strategy routine, so return immediately.

There must be a proc\_handler routine for any terminal nodes mirrored under /proc/sys (non-terminals are handled by a built-in directory handler). Several default handlers are available to cover common cases -

```
proc_dostring, proc_dointvec, proc_dointvec_jiffies,
proc_dointvec_minmax, proc_doulongvec_ms_jiffies_minmax,
proc_doulongvec_minmax
```

It is the handler's job to read the input buffer from user memory and process it. The handler should return 0 on success.

This routine returns NULL on a failure to register, and a pointer to the table header on success.

# unregister\_sysctl\_table

#### Name

unregister\_sysctl\_table — unregister a sysctl table heirarchy

## **Synopsis**

```
void unregister_sysctl_table (struct ctl_table_header * header);
```

## **Arguments**

header

the header returned from register\_sysctl\_table

## **Description**

Unregisters the sysctl table and all children. proc entries may not actually be removed until they are no longer used by anyone.

# proc\_dostring

#### **Name**

```
proc_dostring — read a string sysctl
```

#### **Synopsis**

```
int proc_dostring (ctl_table * table, int write, struct file *
filp, void * buffer, size_t * lenp);
```

## **Arguments**

```
the sysctl table

write

TRUE if this is a write to the sysctl file

filp

the file structure

buffer

the user buffer

lenp

the size of the user buffer
```

## **Description**

Reads/writes a string from/to the user buffer. If the kernel buffer provided is not large enough to hold the string, the string is truncated. The copied string is

NULL-terminated. If the string is being read by the user process, it is copied and a newline '\n' is added. It is truncated if the buffer is not large enough.

Returns 0 on success.

# proc\_dointvec

#### Name

proc\_dointvec — read a vector of integers

## **Synopsis**

```
int proc_dointvec (ctl_table * table, int write, struct file *
filp, void * buffer, size_t * lenp);
```

## **Arguments**

```
table
    the sysctl table

write
    TRUE if this is a write to the sysctl file

filp
    the file structure

buffer
    the user buffer
```

lenp

the size of the user buffer

#### **Description**

Reads/writes up to table->maxlen/sizeof(unsigned int) integer values from/to the user buffer, treated as an ASCII string.

Returns 0 on success.

# proc\_dointvec\_minmax

#### **Name**

proc\_dointvec\_minmax — read a vector of integers with min/max values

#### **Synopsis**

```
int proc_dointvec_minmax (ctl_table * table, int write, struct
file * filp, void * buffer, size_t * lenp);
```

#### **Arguments**

```
table
    the sysctl table

write

TRUE if this is a write to the sysctl file
```

```
the file structure

buffer
the user buffer

lenp
the size of the user buffer
```

#### **Description**

Reads/writes up to table->maxlen/sizeof(unsigned int) integer values from/to the user buffer, treated as an ASCII string.

This routine will ensure the values are within the range specified by table->extra1 (min) and table->extra2 (max).

Returns 0 on success.

# proc\_doulongvec\_minmax

#### **Name**

proc\_doulongvec\_minmax — read a vector of long integers with min/max values

## **Synopsis**

```
int proc_doulongvec_minmax (ctl_table * table, int write, struct
file * filp, void * buffer, size_t * lenp);
```

```
the sysctl table

write

TRUE if this is a write to the sysctl file

filp

the file structure

buffer

the user buffer

lenp

the size of the user buffer
```

#### **Description**

Reads/writes up to table->maxlen/sizeof(unsigned long) unsigned long values from/to the user buffer, treated as an ASCII string.

This routine will ensure the values are within the range specified by table->extra1 (min) and table->extra2 (max).

Returns 0 on success.

# proc\_doulongvec\_ms\_jiffies\_minmax

#### Name

proc\_doulongvec\_ms\_jiffies\_minmax — read a vector of millisecond
values with min/max values

#### **Synopsis**

```
int proc_doulongvec_ms_jiffies_minmax (ctl_table * table, int
write, struct file * filp, void * buffer, size_t * lenp);
```

#### **Arguments**

```
the sysctl table

write

TRUE if this is a write to the sysctl file

filp

the file structure

buffer

the user buffer

lenp

the size of the user buffer
```

#### **Description**

Reads/writes up to table->maxlen/sizeof(unsigned long) unsigned long values from/to the user buffer, treated as an ASCII string. The values are treated as milliseconds, and converted to jiffies when they are stored.

This routine will ensure the values are within the range specified by table->extra1 (min) and table->extra2 (max).

Returns 0 on success.

# proc\_dointvec\_jiffies

#### **Name**

proc\_dointvec\_jiffies — read a vector of integers as seconds

#### **Synopsis**

```
int proc_dointvec_jiffies (ctl_table * table, int write, struct
file * filp, void * buffer, size_t * lenp);
```

#### **Arguments**

```
the sysctl table

write

TRUE if this is a write to the sysctl file

filp

the file structure

buffer

the user buffer

lenp

the size of the user buffer
```

## **Description**

Reads/writes up to table->maxlen/sizeof(unsigned int) integer values from/to the user buffer, treated as an ASCII string. The values read are assumed to be in seconds, and

are converted into jiffies.

Returns 0 on success.

# **Chapter 6. The Linux VFS**

# 6.1. The Directory Cache

# d\_invalidate

#### **Name**

d\_invalidate — invalidate a dentry

## **Synopsis**

```
int d_invalidate (struct dentry * dentry);
```

#### **Arguments**

```
dentry

dentry to invalidate
```

## **Description**

Try to invalidate the dentry if it turns out to be possible. If there are other dentries that can be reached through this one we can't delete it and we return -EBUSY. On success we return 0.

no dcache lock.

# d\_find\_alias

#### **Name**

d\_find\_alias — grab a hashed alias of inode

#### **Synopsis**

```
struct dentry * d_find_alias (struct inode * inode);
```

## **Arguments**

inode

inode in question

## **Description**

If inode has a hashed alias - acquire the reference to alias and return it. Otherwise return NULL. Notice that if inode is a directory there can be only one alias and it can be unhashed only if it has no children.

# prune\_dcache

#### **Name**

prune\_dcache — shrink the dcache

## **Synopsis**

```
void prune_dcache (int count);
```

#### **Arguments**

count

number of entries to try and free

## **Description**

Shrink the dcache. This is done when we need more memory, or simply when we need to unmount something (at which point we need to unuse all dentries).

This function may fail to free any resources if all the dentries are in use.

# shrink\_dcache\_sb

#### Name

shrink\_dcache\_sb — shrink dcache for a superblock

## **Synopsis**

```
void shrink_dcache_sb (struct super_block * sb);
```

sb superblock

# **Description**

Shrink the dcache for the specified super block. This is used to free the dcache before unmounting a file system

# have\_submounts

#### **Name**

have\_submounts — check for mounts over a dentry

## **Synopsis**

```
int have_submounts (struct dentry * parent);
```

## **Arguments**

```
parent dentry to check.
```

#### **Description**

Return true if the parent or its subdirectories contain a mount point

# shrink\_dcache\_parent

#### Name

shrink\_dcache\_parent — prune dcache

## **Synopsis**

```
void shrink_dcache_parent (struct dentry * parent);
```

#### **Arguments**

parent

parent of entries to prune

## **Description**

Prune the dcache to remove unused children of the parent dentry.

# d\_alloc

#### Name

d\_alloc — allocate a dcache entry

## **Synopsis**

```
struct dentry * d_alloc (struct dentry * parent, const struct qstr * name);
```

## **Arguments**

```
parent
parent of entry to allocate

name
qstr of the name
```

## **Description**

Allocates a dentry. It returns NULL if there is insufficient memory available. On a success the dentry is returned. The name passed in is copied and the copy passed in may be reused after this call.

# d\_instantiate

#### **Name**

d\_instantiate — fill in inode information for a dentry

#### **Synopsis**

```
void d_instantiate (struct dentry * entry, struct inode *
inode);
```

#### **Arguments**

```
entry
dentry to complete

inode
inode to attach to this dentry
```

#### **Description**

Fill in inode information in the entry.

This turns negative dentries into productive full members of society.

NOTE! This assumes that the inode count has been incremented (or otherwise set) by the caller to indicate that it is now in use by the dcache.

# d\_alloc\_root

#### Name

d\_alloc\_root — allocate root dentry

## **Synopsis**

```
struct dentry * d_alloc_root (struct inode * root_inode);
```

## **Arguments**

root\_inode

inode to allocate the root for

## **Description**

Allocate a root ("/") dentry for the inode given. The inode is instantiated and returned. NULL is returned if there is insufficient memory or the inode passed is NULL.

# d\_lookup

#### **Name**

d\_lookup — search for a dentry

## **Synopsis**

```
struct dentry * d_lookup (struct dentry * parent, struct qstr *
name);
```

## **Arguments**

```
parent
parent dentry

name
qstr of name we wish to find
```

## **Description**

Searches the children of the parent dentry for the name in question. If the dentry is found its reference count is incremented and the dentry is returned. The caller must use d\_put to free the entry when it has finished using it. NULL is returned on failure.

## d\_validate

#### **Name**

d\_validate — verify dentry provided from insecure source

## **Synopsis**

```
int d_validate (struct dentry * dentry, struct dentry *
dparent);
```

```
dentry
```

The dentry alleged to be valid child of dparent

dparent

The parent dentry (known to be valid)

## **Description**

An insecure source has sent us a dentry, here we verify it and dget it. This is used by ncpfs in its readdir implementation. Zero is returned in the dentry is invalid.

# d\_delete

#### Name

d\_delete — delete a dentry

## **Synopsis**

```
void d_delete (struct dentry * dentry);
```

dentry

The dentry to delete

# **Description**

Turn the dentry into a negative dentry if possible, otherwise remove it from the hash queues so it can be deleted later

# d\_rehash

#### **Name**

d\_rehash — add an entry back to the hash

## **Synopsis**

```
void d_rehash (struct dentry * entry);
```

## **Arguments**

entry

dentry to add to the hash

## **Description**

Adds a dentry to the hash according to its name.

## d\_move

#### **Name**

```
d_move — move a dentry
```

## **Synopsis**

```
void d_move (struct dentry * dentry, struct dentry * target);
```

#### **Arguments**

```
dentry
entry to move
target
new dentry
```

## **Description**

Update the dcache to reflect the move of a file name. Negative dcache entries should not be moved in this way.

# \_\_d\_path

#### Name

```
__d_path — return the path of a dentry
```

## **Synopsis**

```
char * __d_path (struct dentry * dentry, struct vfsmount *
vfsmnt, struct dentry * root, struct vfsmount * rootmnt, char *
buffer, int buflen);
```

## **Arguments**

```
dentry
    dentry to report

vfsmnt
    vfsmnt to which the dentry belongs

root
    root dentry

rootmnt
    vfsmnt to which the root dentry belongs

buffer
    buffer to return value in

buflen
    buffer length
```

#### **Description**

Convert a dentry into an ASCII path name. If the entry has been deleted the string " (deleted)" is appended. Note that this is ambiguous. Returns the buffer.

"buflen" should be PAGE\_SIZE or more. Caller holds the dcache\_lock.

## is\_subdir

#### **Name**

is\_subdir — is new dentry a subdirectory of old\_dentry

## **Synopsis**

```
int is_subdir (struct dentry * new_dentry, struct dentry *
old_dentry);
```

#### **Arguments**

```
new_dentry
new dentry

old_dentry

old dentry
```

## **Description**

Returns 1 if new\_dentry is a subdirectory of the parent (at any depth). Returns 0 otherwise.

## find\_inode\_number

#### **Name**

find\_inode\_number — check for dentry with name

## **Synopsis**

```
ino_t find_inode_number (struct dentry * dir, struct qstr *
name);
```

#### **Arguments**

```
dir
directory to check

name

Name to find.
```

#### **Description**

Check whether a dentry already exists for the given name, and return the inode number if it has an inode. Otherwise 0 is returned.

This routine is used to post-process directory listings for filesystems using synthetic inode numbers, and is necessary to keep getcwd working.

## d\_drop

#### **Name**

```
d_drop — drop a dentry
```

### **Synopsis**

```
void d_drop (struct dentry * dentry);
```

### **Arguments**

```
dentry
dentry to drop
```

## **Description**

d\_drop unhashes the entry from the parent dentry hashes, so that it won't be found through a VFS lookup any more. Note that this is different from deleting the dentry - d\_delete will try to mark the dentry negative if possible, giving a successful \_negative\_ lookup, while d\_drop will just make the cache lookup fail.

d\_drop is used mainly for stuff that wants to invalidate a dentry for some reason (NFS timeouts or autofs deletes).

# $d_add$

#### Name

d\_add — add dentry to hash queues

## **Synopsis**

```
void d_add (struct dentry * entry, struct inode * inode);
```

### **Arguments**

```
entry
dentry to add

inode

The inode to attach to this dentry
```

### **Description**

This adds the entry to the hash queues and initializes *inode*. The entry was actually filled in earlier during d\_alloc.

# dget

#### **Name**

dget — get a reference to a dentry

### **Synopsis**

```
struct dentry * dget (struct dentry * dentry);
```

### **Arguments**

```
dentry
dentry to get a reference to
```

### **Description**

Given a dentry or NULL pointer increment the reference count if appropriate and return the dentry. A dentry will not be destroyed when it has references. dget should never be called for dentries with zero reference counter. For these cases (preferably none, functions in dcache.c are sufficient for normal needs and they take necessary precautions) you should hold dcache\_lock and call dget\_locked instead of dget.

## d\_unhashed

#### **Name**

d\_unhashed — is dentry hashed

## **Synopsis**

```
int d_unhashed (struct dentry * dentry);
```

### **Arguments**

```
dentry
entry to check
```

### **Description**

Returns true if the dentry passed is not currently hashed.

# 6.2. Inode Handling

# \_\_mark\_inode\_dirty

## Name

```
__mark_inode_dirty — internal function
```

```
void __mark_inode_dirty (struct inode * inode, int flags);
```

```
inode
  inode to mark

flags
  what kind of dirty (i.e. I_DIRTY_SYNC) Mark an inode as dirty. Callers should
  use mark_inode_dirty or mark_inode_dirty_sync.
```

# write\_inode\_now

### **Name**

write\_inode\_now — write an inode to disk

## **Synopsis**

```
void write_inode_now (struct inode * inode, int sync);
```

### **Arguments**

```
inode
  inode to write to disk
sync
  whether the write should be synchronous or not
```

### **Description**

This function commits an inode to disk immediately if it is dirty. This is primarily needed by knfsd.

# clear\_inode

#### **Name**

```
clear_inode — clear an inode
```

### **Synopsis**

```
void clear_inode (struct inode * inode);
```

### **Arguments**

inode

inode to clear

### **Description**

This is called by the filesystem to tell us that the inode is no longer useful. We just terminate it with extreme prejudice.

## invalidate\_inodes

#### **Name**

invalidate\_inodes — discard the inodes on a device

### **Synopsis**

```
int invalidate_inodes (struct super_block * sb);
```

### **Arguments**

sb

superblock

### **Description**

Discard all of the inodes for a given superblock. If the discard fails because there are busy inodes then a non zero value is returned. If the discard is successful all the inodes have been discarded.

# get\_empty\_inode

#### **Name**

get\_empty\_inode — obtain an inode

### **Synopsis**

```
struct inode * get_empty_inode ( void);
```

### **Arguments**

void

no arguments

### **Description**

This is called by things like the networking layer etc that want to get an inode without any inode number, or filesystems that allocate new inodes with no pre-existing information.

On a successful return the inode pointer is returned. On a failure a NULL pointer is returned. The returned inode is not on any superblock lists.

# iunique

#### Name

iunique — get a unique inode number

```
ino_t iunique (struct super_block * sb, ino_t max_reserved);
```

```
superblock

max_reserved

highest reserved inode number
```

### **Description**

Obtain an inode number that is unique on the system for a given superblock. This is used by file systems that have no natural permanent inode numbering system. An inode number is returned that is higher than the reserved limit but unique.

#### **BUGS**

With a large number of inodes live on the file system this function currently becomes quite slow.

## insert\_inode\_hash

#### **Name**

```
insert_inode_hash — hash an inode
```

```
void insert_inode_hash (struct inode * inode);
```

inode

unhashed inode

### **Description**

Add an inode to the inode hash for this superblock. If the inode has no superblock it is added to a separate anonymous chain.

# remove\_inode\_hash

#### **Name**

remove\_inode\_hash — remove an inode from the hash

## **Synopsis**

```
void remove_inode_hash (struct inode * inode);
```

### **Arguments**

inode

inode to unhash

## **Description**

Remove an inode from the superblock or anonymous hash.

# iput

#### Name

iput — put an inode

## **Synopsis**

```
void iput (struct inode * inode);
```

### **Arguments**

inode

inode to put

### **Description**

Puts an inode, dropping its usage count. If the inode use count hits zero the inode is also then freed and may be destroyed.

# bmap

#### Name

bmap — find a block number in a file

### **Synopsis**

```
int bmap (struct inode * inode, int block);
```

### **Arguments**

```
inode
inode of file
block
block to find
```

### **Description**

Returns the block number on the device holding the inode that is the disk block number for the block of the file requested. That is, asked for block 4 of inode 1 the function will return the disk block relative to the disk start that holds that block of the file.

## update\_atime

#### **Name**

update\_atime — update the access time

### **Synopsis**

```
void update_atime (struct inode * inode);
```

### **Arguments**

inode

inode accessed

## **Description**

Update the accessed time on an inode and mark it for writeback. This function automatically handles read only file systems and media, as well as the "noatime" flag and inode specific "noatime" markers.

# make\_bad\_inode

#### **Name**

make\_bad\_inode — mark an inode bad due to an I/O error

## **Synopsis**

```
void make_bad_inode (struct inode * inode);
```

### **Arguments**

inode

Inode to mark bad

## **Description**

When an inode cannot be read due to a media or remote network failure this function makes the inode "bad" and causes I/O operations on it to fail from this point on.

# is\_bad\_inode

#### Name

is\_bad\_inode — is an inode errored

```
int is_bad_inode (struct inode * inode);
```

inode

inode to test

# **Description**

Returns true if the inode in question has been marked as bad.

# 6.3. Registration and Superblocks

# register\_filesystem

#### **Name**

register\_filesystem — register a new filesystem

# **Synopsis**

```
int register_filesystem (struct file_system_type * fs);
```

### **Arguments**

fs

the file system structure

### **Description**

Adds the file system passed to the list of file systems the kernel is aware of for mount and other syscalls. Returns 0 on success, or a negative errno code on an error.

The &struct file\_system\_type that is passed is linked into the kernel structures and must not be freed until the file system has been unregistered.

# unregister\_filesystem

#### **Name**

unregister\_filesystem — unregister a file system

### **Synopsis**

```
int unregister_filesystem (struct file_system_type * fs);
```

### **Arguments**

fs

filesystem to unregister

### **Description**

Remove a file system that was previously successfully registered with the kernel. An error is returned if the file system is not found. Zero is returned on a success.

Once this function has returned the &struct file\_system\_type structure may be freed or reused.

# get\_super

#### **Name**

get\_super — get the superblock of a device

## **Synopsis**

```
struct super_block * get_super (kdev_t dev);
```

### **Arguments**

dev

device to get the superblock for

# **Description**

Scans the superblock list and finds the superblock of the file system mounted on the device given. NULL is returned if no match is found.

## 6.4. File Locks

# posix\_lock\_file

#### Name

```
posix_lock_file --
```

### **Synopsis**

```
int posix_lock_file (struct file * filp, struct file_lock *
caller, unsigned int wait);
```

### **Arguments**

```
The file to apply the lock to

caller

The lock to be applied

wait

1 to retry automatically, 0 to return -EAGAIN
```

### **Description**

Add a POSIX style lock to a file. We merge adjacent locks whenever possible. POSIX locks are sorted by owner task, then by starting address

#### Kai Petzke writes

To make freeing a lock much faster, we keep a pointer to the lock before the actual one. But the real gain of the new coding was, that lock\_it and unlock\_it became one function.

### To all purists

Yes, I use a few goto's. Just pass on to the next function.

# \_\_get\_lease

### Name

\_\_get\_lease — revoke all outstanding leases on file

## **Synopsis**

```
int __get_lease (struct inode * inode, unsigned int mode);
```

### **Arguments**

```
inode
     the inode of the file to return
mode
     the open mode (read or write)
```

### **Description**

get\_lease (inlined for speed) has checked there already is a lease on this file. Leases are broken on a call to open or truncate. This function can sleep unless you specified O\_NONBLOCK to your open.

# lease\_get\_mtime

#### **Name**

```
lease_get_mtime —
```

### **Synopsis**

```
time_t lease_get_mtime (struct inode * inode);
```

### **Arguments**

inode

the inode

### **Description**

This is to force NFS clients to flush their caches for files with exclusive leases. The justification is that if someone has an exclusive lease, then they could be modifying it.

# posix\_block\_lock

#### Name

posix\_block\_lock — blocks waiting for a file lock

## **Synopsis**

```
void posix_block_lock (struct file_lock * blocker, struct
file_lock * waiter);
```

### **Arguments**

blocker

the lock which is blocking

waiter

the lock which conflicts and has to wait

### **Description**

lockd needs to block waiting for locks.

# posix\_unblock\_lock

#### Name

posix\_unblock\_lock — stop waiting for a file lock

## **Synopsis**

```
void posix_unblock_lock (struct file_lock * waiter);
```

### **Arguments**

waiter

the lock which was waiting

### **Description**

lockd needs to block waiting for locks.

# lock\_may\_read

#### Name

lock\_may\_read — checks that the region is free of locks

```
int lock_may_read (struct inode * inode, loff_t start, unsigned
long len);
```

```
inode
    the inode that is being read
start
    the first byte to read
len
    the number of bytes to read
```

### **Description**

Emulates Windows locking requirements. Whole-file mandatory locks (share modes) can prohibit a read and byte-range POSIX locks can prohibit a read if they overlap.

N.B. this function is only ever called from knfsd and ownership of locks is never checked.

## lock\_may\_write

#### Name

lock\_may\_write — checks that the region is free of locks

```
int lock_may_write (struct inode * inode, loff_t start, unsigned
long len);
```

```
inode
    the inode that is being written

start
    the first byte to write

len
    the number of bytes to write
```

### **Description**

Emulates Windows locking requirements. Whole-file mandatory locks (share modes) can prohibit a write and byte-range POSIX locks can prohibit a write if they overlap.

N.B. this function is only ever called from knfsd and ownership of locks is never checked.

# fcntl\_getlease

#### **Name**

fcntl\_getlease — Enquire what lease is currently active

```
int fcntl_getlease (struct file * filp);
```

```
filp
the file
```

### **Description**

The value returned by this function will be one of

F\_RDLCK to indicate a read-only (type II) lease is held.

F\_WRLCK to indicate an exclusive lease is held.

#### XXX

sfr & i disagree over whether F\_INPROGRESS should be returned to userspace.

# fcntl\_setlease

#### **Name**

fcntl\_setlease — sets a lease on an open file

```
int fcntl_setlease (unsigned int fd, struct file * filp, long arg);
```

```
open file descriptor

filp
file pointer

arg
type of lease to obtain
```

## **Description**

Call this fcntl to establish a lease on the file. Note that you also need to call F\_SETSIG to receive a signal when the lease is broken.

# sys\_flock

#### **Name**

```
sys_flock — flock system call.
```

```
asmlinkage long sys_flock (unsigned int fd, unsigned int cmd);
```

```
fd
```

the file descriptor to lock.

cmd

the type of lock to apply.

### **Description**

```
Apply a FL_FLOCK style lock to an open file descriptor. The cmd can be one of
```

LOCK\_SH -- a shared lock.

LOCK\_EX -- an exclusive lock.

LOCK\_UN -- remove an existing lock.

LOCK\_MAND -- a 'mandatory' flock. This exists to emulate Windows Share Modes.

LOCK\_MAND can be combined with LOCK\_READ or LOCK\_WRITE to allow other processes read and write access respectively.

## get\_locks\_status

#### **Name**

get\_locks\_status — reports lock usage in /proc/locks

```
int get_locks_status (char * buffer, char ** start, off_t
offset, int length);
```

```
buffer
    address in userspace to write into

start
    ?

offset
    how far we are through the buffer

length
    how much to read
```

# **Chapter 7. Linux Networking**

### 7.1. Socket Buffer Functions

# skb\_queue\_empty

#### **Name**

skb\_queue\_empty — check if a queue is empty

## **Synopsis**

```
int skb_queue_empty (struct sk_buff_head * list);
```

### **Arguments**

```
list
queue head
```

### **Description**

Returns true if the queue is empty, false otherwise.

# skb\_get

#### Name

skb\_get — reference buffer

## **Synopsis**

```
struct sk_buff * skb_get (struct sk_buff * skb);
```

### **Arguments**

skb

buffer to reference

## **Description**

Makes another reference to a socket buffer and returns a pointer to the buffer.

# kfree\_skb

### **Name**

kfree\_skb — free an sk\_buff

## **Synopsis**

```
void kfree_skb (struct sk_buff * skb);
```

## **Arguments**

skb

buffer to free

## **Description**

Drop a reference to the buffer and free it if the usage count has hit zero.

# skb\_cloned

### **Name**

skb\_cloned — is the buffer a clone

```
int skb_cloned (struct sk_buff * skb);
```

skb

buffer to check

# **Description**

Returns true if the buffer was generated with skb\_clone and is one of multiple shared copies of the buffer. Cloned buffers are shared data so must not be written to under normal circumstances.

## skb\_shared

#### **Name**

skb\_shared — is the buffer shared

## **Synopsis**

```
int skb_shared (struct sk_buff * skb);
```

### **Arguments**

skb

buffer to check

### **Description**

Returns true if more than one person has a reference to this buffer.

## skb\_share\_check

### **Name**

skb\_share\_check — check if buffer is shared and if so clone it

### **Synopsis**

```
struct sk_buff * skb_share_check (struct sk_buff * skb, int
pri);
```

### **Arguments**

```
skbbuffer to checkpripriority for memory allocation
```

### **Description**

If the buffer is shared the buffer is cloned and the old copy drops a reference. A new clone with a single reference is returned. If the buffer is not shared the original buffer is returned. When being called from interrupt status or with spinlocks held pri must be GFP\_ATOMIC.

NULL is returned on a memory allocation failure.

## skb\_unshare

#### Name

skb\_unshare — make a copy of a shared buffer

### **Synopsis**

```
struct sk_buff * skb_unshare (struct sk_buff * skb, int pri);
```

### **Arguments**

```
skbbuffer to checkpripriority for memory allocation
```

### **Description**

If the socket buffer is a clone then this function creates a new copy of the data, drops a reference count on the old copy and returns the new copy with the reference count at 1. If the buffer is not a clone the original buffer is returned. When called with a spinlock held or from interrupt state <code>pri</code> must be <code>GFP\_ATOMIC</code>

NULL is returned on a memory allocation failure.

## skb\_peek

#### **Name**

```
skb_peek —
```

### **Synopsis**

```
struct sk_buff * skb_peek (struct sk_buff_head * list_);
```

### **Arguments**

```
list_ list to peek at
```

### **Description**

Peek an &sk\_buff. Unlike most other operations you \_MUST\_ be careful with this one. A peek leaves the buffer on the list and someone else may run off with it. You must hold the appropriate locks or have a private queue to do this.

Returns NULL for an empty list or a pointer to the head element. The reference count is not incremented and the reference is therefore volatile. Use with caution.

## skb\_peek\_tail

#### **Name**

```
skb_peek_tail —
```

## **Synopsis**

```
struct sk_buff * skb_peek_tail (struct sk_buff_head * list_);
```

### **Arguments**

```
list_ list to peek at
```

### **Description**

Peek an &sk\_buff. Unlike most other operations you \_MUST\_ be careful with this one. A peek leaves the buffer on the list and someone else may run off with it. You must hold the appropriate locks or have a private queue to do this.

Returns NULL for an empty list or a pointer to the tail element. The reference count is not incremented and the reference is therefore volatile. Use with caution.

# skb\_queue\_len

#### Name

skb\_queue\_len — get queue length

# **Synopsis**

```
__u32 skb_queue_len (struct sk_buff_head * list_);
```

# **Arguments**

list\_

list to measure

# **Description**

Return the length of an &sk\_buff queue.

# \_\_skb\_queue\_head

#### Name

\_\_skb\_queue\_head — queue a buffer at the list head

```
void __skb_queue_head (struct sk_buff_head * list, struct
sk_buff * newsk);
```

# **Arguments**

```
list to use newsk buffer to queue
```

# **Description**

Queue a buffer at the start of a list. This function takes no locks and you must therefore hold required locks before calling it.

A buffer cannot be placed on two lists at the same time.

# skb\_queue\_head

#### **Name**

skb\_queue\_head — queue a buffer at the list head

```
void skb_queue_head (struct sk_buff_head * list, struct sk_buff
* newsk);
```

# **Arguments**

list

list to use

newsk

buffer to queue

# **Description**

Queue a buffer at the start of the list. This function takes the list lock and can be used safely with other locking &sk\_buff functions safely.

A buffer cannot be placed on two lists at the same time.

# \_\_skb\_queue\_tail

#### **Name**

\_\_skb\_queue\_tail — queue a buffer at the list tail

```
void __skb_queue_tail (struct sk_buff_head * list, struct
sk_buff * newsk);
```

# **Arguments**

```
list to use newsk buffer to queue
```

# **Description**

Queue a buffer at the end of a list. This function takes no locks and you must therefore hold required locks before calling it.

A buffer cannot be placed on two lists at the same time.

# skb\_queue\_tail

#### **Name**

skb\_queue\_tail — queue a buffer at the list tail

```
void skb_queue_tail (struct sk_buff_head * list, struct sk_buff
* newsk);
```

## **Arguments**

```
list list to use newsk buffer to queue
```

# **Description**

Queue a buffer at the tail of the list. This function takes the list lock and can be used safely with other locking &sk\_buff functions safely.

A buffer cannot be placed on two lists at the same time.

# \_\_skb\_dequeue

#### **Name**

```
__skb_dequeue — remove from the head of the queue
```

```
struct sk_buff * __skb_dequeue (struct sk_buff_head * list);
```

list

list to dequeue from

# **Description**

Remove the head of the list. This function does not take any locks so must be used with appropriate locks held only. The head item is returned or NULL if the list is empty.

# skb\_dequeue

#### Name

skb\_dequeue — remove from the head of the queue

# **Synopsis**

```
struct sk_buff * skb_dequeue (struct sk_buff_head * list);
```

# **Arguments**

list

list to dequeue from

Remove the head of the list. The list lock is taken so the function may be used safely with other locking list functions. The head item is returned or NULL if the list is empty.

# skb\_insert

#### **Name**

```
skb_insert — insert a buffer
```

# **Synopsis**

```
void skb_insert (struct sk_buff * old, struct sk_buff * newsk);
```

### **Arguments**

old

buffer to insert before

newsk

buffer to insert

# **Description**

Place a packet before a given packet in a list. The list locks are taken and this function is atomic with respect to other list locked calls A buffer cannot be placed on two lists at the same time.

# skb\_append

#### **Name**

```
skb_append — append a buffer
```

# **Synopsis**

```
void skb_append (struct sk_buff * old, struct sk_buff * newsk);
```

# **Arguments**

old

buffer to insert after

newsk

buffer to insert

# **Description**

Place a packet after a given packet in a list. The list locks are taken and this function is atomic with respect to other list locked calls. A buffer cannot be placed on two lists at the same time.

# skb\_unlink

#### **Name**

skb\_unlink — remove a buffer from a list

## **Synopsis**

```
void skb_unlink (struct sk_buff * skb);
```

# **Arguments**

skb

buffer to remove

# **Description**

Place a packet after a given packet in a list. The list locks are taken and this function is atomic with respect to other list locked calls

Works even without knowing the list it is sitting on, which can be handy at times. It also means that THE LIST MUST EXIST when you unlink. Thus a list must have its contents unlinked before it is destroyed.

# \_\_skb\_dequeue\_tail

#### Name

\_\_skb\_dequeue\_tail — remove from the tail of the queue

## **Synopsis**

```
struct sk_buff * __skb_dequeue_tail (struct sk_buff_head *
list);
```

## **Arguments**

list

list to dequeue from

## **Description**

Remove the tail of the list. This function does not take any locks so must be used with appropriate locks held only. The tail item is returned or NULL if the list is empty.

# skb\_dequeue\_tail

#### Name

skb\_dequeue\_tail — remove from the head of the queue

```
struct sk_buff * skb_dequeue_tail (struct sk_buff_head * list);
```

## **Arguments**

list

list to dequeue from

# **Description**

Remove the head of the list. The list lock is taken so the function may be used safely with other locking list functions. The tail item is returned or NULL if the list is empty.

# skb\_put

### Name

skb\_put — add data to a buffer

```
unsigned char * skb_put (struct sk_buff * skb, unsigned int
len);
```

```
skb
buffer to use

len
amount of data to add
```

# **Description**

This function extends the used data area of the buffer. If this would exceed the total buffer size the kernel will panic. A pointer to the first byte of the extra data is returned.

# skb\_push

#### **Name**

skb\_push — add data to the start of a buffer

# **Synopsis**

```
unsigned char * skb_push (struct sk_buff * skb, unsigned int
len);
```

# **Arguments**

skb

buffer to use

len

amount of data to add

## **Description**

This function extends the used data area of the buffer at the buffer start. If this would exceed the total buffer headroom the kernel will panic. A pointer to the first byte of the extra data is returned.

# skb\_pull

#### Name

skb\_pull — remove data from the start of a buffer

# **Synopsis**

```
unsigned char * skb_pull (struct sk_buff * skb, unsigned int
len);
```

## **Arguments**

skb

buffer to use

len

amount of data to remove

This function removes data from the start of a buffer, returning the memory to the headroom. A pointer to the next data in the buffer is returned. Once the data has been pulled future pushes will overwrite the old data.

# skb\_headroom

#### **Name**

skb\_headroom — bytes at buffer head

# **Synopsis**

```
int skb_headroom (const struct sk_buff * skb);
```

## **Arguments**

skb

buffer to check

## **Description**

Return the number of bytes of free space at the head of an &sk\_buff.

# skb\_tailroom

#### Name

skb\_tailroom — bytes at buffer end

# **Synopsis**

```
int skb_tailroom (const struct sk_buff * skb);
```

# **Arguments**

skb

buffer to check

# **Description**

Return the number of bytes of free space at the tail of an sk\_buff

# skb\_reserve

#### Name

skb\_reserve — adjust headroom

```
void skb_reserve (struct sk_buff * skb, unsigned int len);
```

## **Arguments**

```
buffer to alter

len

bytes to move
```

# **Description**

Increase the headroom of an empty &sk\_buff by reducing the tail room. This is only allowed for an empty buffer.

# skb\_trim

#### Name

```
skb_trim — remove end from a buffer
```

```
void skb_trim (struct sk_buff * skb, unsigned int len);
```

```
skb
buffer to alter

len
new length
```

# **Description**

Cut the length of a buffer down by removing data from the tail. If the buffer is already under the length specified it is not modified.

# skb\_orphan

#### Name

```
skb_orphan — orphan a buffer
```

# **Synopsis**

```
void skb_orphan (struct sk_buff * skb);
```

# **Arguments**

```
skb
```

buffer to orphan

If a buffer currently has an owner then we call the owner's destructor function and make the *skb* unowned. The buffer continues to exist but is no longer charged to its former owner.

# skb\_queue\_purge

#### **Name**

```
skb_queue_purge — empty a list
```

# **Synopsis**

```
void skb_queue_purge (struct sk_buff_head * list);
```

## **Arguments**

list

list to empty

## **Description**

Delete all buffers on an &sk\_buff list. Each buffer is removed from the list and one reference dropped. This function takes the list lock and is atomic with respect to other list locking functions.

# \_\_skb\_queue\_purge

#### Name

```
__skb_queue_purge — empty a list
```

# **Synopsis**

```
void __skb_queue_purge (struct sk_buff_head * list);
```

### **Arguments**

list

list to empty

# **Description**

Delete all buffers on an &sk\_buff list. Each buffer is removed from the list and one reference dropped. This function does not take the list lock and the caller must hold the relevant locks to use it.

# \_\_dev\_alloc\_skb

#### **Name**

\_\_dev\_alloc\_skb — allocate an skbuff for sending

```
struct sk_buff * __dev_alloc_skb (unsigned int length, int
gfp_mask);
```

## **Arguments**

```
length
length to allocate

gfp_mask

get_free_pages mask, passed to alloc_skb
```

## **Description**

Allocate a new &sk\_buff and assign it a usage count of one. The buffer has unspecified headroom built in. Users should allocate the headroom they think they need without accounting for the built in space. The built in space is used for optimisations.

NULL is returned in there is no free memory.

# dev\_alloc\_skb

#### **Name**

dev\_alloc\_skb — allocate an skbuff for sending

```
struct sk_buff * dev_alloc_skb (unsigned int length);
```

## **Arguments**

length

length to allocate

## **Description**

Allocate a new &sk\_buff and assign it a usage count of one. The buffer has unspecified headroom built in. Users should allocate the headroom they think they need without accounting for the built in space. The built in space is used for optimisations.

NULL is returned in there is no free memory. Although this function allocates memory it can be called from an interrupt.

# skb\_cow

#### Name

skb\_cow — copy header of skb when it is required

```
int skb_cow (struct sk_buff * skb, unsigned int headroom);
```

skb

buffer to cow

headroom

needed headroom

# **Description**

If the skb passed lacks sufficient headroom or its data part is shared, data is reallocated. If reallocation fails, an error is returned and original skb is not changed.

The result is skb with writable area skb->head...skb->tail and at least *headroom* of space at head.

# skb\_over\_panic

#### **Name**

skb\_over\_panic — private function

```
void skb_over_panic (struct sk_buff * skb, int sz, void * here);
```

```
skb
buffer
sz
size
here
address
```

# **Description**

Out of line support code for skb\_put. Not user callable.

# skb\_under\_panic

#### **Name**

```
skb_under_panic — private function
```

```
void skb_under_panic (struct sk_buff * skb, int sz, void *
here);
```

```
skb
buffer
sz
size
here
address
```

# **Description**

Out of line support code for skb\_push. Not user callable.

# alloc\_skb

#### **Name**

alloc\_skb — allocate a network buffer

```
struct sk_buff * alloc_skb (unsigned int size, int gfp_mask);
```

```
size to allocate

gfp_mask

allocation mask
```

# **Description**

Allocate a new &sk\_buff. The returned buffer has no headroom and a tail room of size bytes. The object has a reference count of one. The return is the buffer. On a failure the return is NULL.

Buffers may only be allocated from interrupts using a gfp\_mask of GFP\_ATOMIC.

# \_\_kfree\_skb

#### **Name**

```
__kfree_skb — private function
```

```
void __kfree_skb (struct sk_buff * skb);
```

skb

buffer

# **Description**

Free an sk\_buff. Release anything attached to the buffer. Clean the state. This is an internal helper function. Users should always call kfree\_skb

# skb\_clone

#### Name

```
skb_clone — duplicate an sk_buff
```

# **Synopsis**

```
struct sk_buff * skb_clone (struct sk_buff * skb, int gfp_mask);
```

# **Arguments**

```
skb
```

buffer to clone

gfp\_mask

allocation priority

Duplicate an &sk\_buff. The new one is not owned by a socket. Both copies share the same packet data but not structure. The new buffer has a reference count of 1. If the allocation fails the function returns NULL otherwise the new buffer is returned.

If this function is called from an interrupt gfp\_mask must be GFP\_ATOMIC.

# skb\_copy

#### **Name**

skb\_copy — create private copy of an sk\_buff

# **Synopsis**

```
struct sk_buff * skb_copy (const struct sk_buff * skb, int
gfp_mask);
```

# **Arguments**

```
buffer to copy

gfp_mask

allocation priority
```

Make a copy of both an &sk\_buff and its data. This is used when the caller wishes to modify the data and needs a private copy of the data to alter. Returns NULL on failure or the pointer to the buffer on success. The returned buffer has a reference count of 1.

As by-product this function converts non-linear &sk\_buff to linear one, so that &sk\_buff becomes completely private and caller is allowed to modify all the data of returned buffer. This means that this function is not recommended for use in circumstances when only header is going to be modified. Use pskb\_copy instead.

# pskb\_copy

#### **Name**

pskb\_copy — create copy of an sk\_buff with private head.

# **Synopsis**

```
struct sk_buff * pskb_copy (struct sk_buff * skb, int gfp_mask);
```

### **Arguments**

```
skb
buffer to copy

gfp_mask
allocation priority
```

Make a copy of both an &sk\_buff and part of its data, located in header. Fragmented data remain shared. This is used when the caller wishes to modify only header of &sk\_buff and needs private copy of the header to alter. Returns NULL on failure or the pointer to the buffer on success. The returned buffer has a reference count of 1.

# pskb\_expand\_head

#### **Name**

pskb\_expand\_head — reallocate header of sk\_buff

## **Synopsis**

```
int pskb_expand_head (struct sk_buff * skb, int nhead, int
ntail, int gfp_mask);
```

### **Arguments**

```
skb
```

buffer to reallocate

nhead

room to add at head

ntail

room to add at tail

```
gfp_mask
allocation priority
```

Expands (or creates identical copy, if &nhead and &ntail are zero) header of skb. &sk\_buff itself is not changed. &sk\_buff MUST have reference count of 1. Returns zero in the case of success or error, if expansion failed. In the last case, &sk\_buff is not changed.

All the pointers pointing into skb header may change and must be reloaded after call to this function.

# skb\_copy\_expand

#### Name

```
skb_copy_expand — copy and expand sk_buff
```

## **Synopsis**

```
struct sk_buff * skb_copy_expand (const struct sk_buff * skb,
int newheadroom, int newtailroom, int gfp_mask);
```

### **Arguments**

```
skb
```

buffer to copy

```
newheadroom

new free bytes at head

newtailroom

new free bytes at tail

gfp_mask

allocation priority
```

Make a copy of both an &sk\_buff and its data and while doing so allocate additional space.

This is used when the caller wishes to modify the data and needs a private copy of the data to alter as well as more space for new fields. Returns NULL on failure or the pointer to the buffer on success. The returned buffer has a reference count of 1.

You must pass GFP\_ATOMIC as the allocation priority if this function is called from an interrupt.

# \_\_pskb\_pull\_tail

#### **Name**

```
__pskb_pull_tail — advance tail of skb header
```

```
unsigned char * __pskb_pull_tail (struct sk_buff * skb, int
delta);
```

skb

buffer to reallocate

delta

number of bytes to advance tail

## **Description**

The function makes a sense only on a fragmented &sk\_buff, it expands header moving its tail forward and copying necessary data from fragmented part.

&sk\_buff MUST have reference count of 1.

Returns NULL (and &sk\_buff does not change) if pull failed or value of new tail of skb in the case of success.

All the pointers pointing into skb header may change and must be reloaded after call to this function.

## 7.2. Socket Filter

# sk\_run\_filter

#### **Name**

sk\_run\_filter — run a filter on a socket

```
int sk_run_filter (struct sk_buff * skb, struct sock_filter *
filter, int flen);
```

# **Arguments**

```
buffer to run the filter on

filter
filter to apply

flen
length of filter
```

# **Description**

Decode and apply filter instructions to the skb->data. Return length to keep, 0 for none. skb is the data we are filtering, filter is the array of filter instructions, and len is the number of filter blocks in the array.

# sk\_chk\_filter

#### **Name**

sk\_chk\_filter — verify socket filter code

```
int sk_chk_filter (struct sock_filter * filter, int flen);
```

## **Arguments**

```
filter
filter to verify

flen
length of filter
```

# **Description**

Check the user's filter code. If we let some ugly filter code slip through kaboom! The filter must contain no references or jumps that are out of range, no illegal instructions and no backward jumps. It must end with a RET instruction

Returns 0 if the rule set is legal or a negative errno code if not.

# Chapter 8. Network device support

# 8.1. Driver Support

# init\_etherdev

#### **Name**

init\_etherdev — Register ethernet device

# **Synopsis**

```
struct net_device * init_etherdev (struct net_device * dev, int
sizeof_priv);
```

# **Arguments**

dev

An ethernet device structure to be filled in, or NULL if a new struct should be allocated.

```
sizeof_priv
```

Size of additional driver-private structure to be allocated for this ethernet device

### **Description**

Fill in the fields of the device structure with ethernet-generic values.

If no device structure is passed, a new one is constructed, complete with a private data area of size <code>sizeof\_priv</code>. A 32-byte (not bit) alignment is enforced for this private data area.

If an empty string area is passed as dev->name, or a new structure is made, a new name string is constructed.

# alloc\_etherdev

#### Name

alloc\_etherdev — Allocates and sets up an ethernet device

# **Synopsis**

```
struct net_device * alloc_etherdev (int sizeof_priv);
```

### **Arguments**

```
sizeof_priv
```

Size of additional driver-private structure to be allocated for this ethernet device

## **Description**

Fill in the fields of the device structure with ethernet-generic values. Basically does everything except registering the device.

Constructs a new net device, complete with a private data area of size sizeof\_priv. A 32-byte (not bit) alignment is enforced for this private data area.

## init\_fddidev

#### **Name**

```
init_fddidev — Register FDDI device
```

## **Synopsis**

```
struct net_device * init_fddidev (struct net_device * dev, int
sizeof priv);
```

### **Arguments**

dev

A FDDI device structure to be filled in, or NULL if a new struct should be allocated.

```
sizeof_priv
```

Size of additional driver-private structure to be allocated for this ethernet device

#### **Description**

Fill in the fields of the device structure with FDDI-generic values.

If no device structure is passed, a new one is constructed, complete with a private data area of size <code>sizeof\_priv</code>. A 32-byte (not bit) alignment is enforced for this private data area.

If an empty string area is passed as dev->name, or a new structure is made, a new name string is constructed.

# alloc\_fddidev

#### **Name**

```
alloc_fddidev — Register FDDI device
```

## **Synopsis**

```
struct net_device * alloc_fddidev (int sizeof_priv);
```

## **Arguments**

```
sizeof_priv
```

Size of additional driver-private structure to be allocated for this FDDI device

## **Description**

Fill in the fields of the device structure with FDDI-generic values.

Constructs a new net device, complete with a private data area of size <code>sizeof\_priv</code>. A 32-byte (not bit) alignment is enforced for this private data area.

# init\_hippi\_dev

#### **Name**

```
init_hippi_dev — Register HIPPI device
```

## **Synopsis**

```
struct net_device * init_hippi_dev (struct net_device * dev, int
sizeof_priv);
```

## **Arguments**

dev

A HIPPI device structure to be filled in, or NULL if a new struct should be allocated.

```
sizeof_priv
```

Size of additional driver-private structure to be allocated for this ethernet device

## **Description**

Fill in the fields of the device structure with HIPPI-generic values.

If no device structure is passed, a new one is constructed, complete with a private data area of size <code>sizeof\_priv</code>. A 32-byte (not bit) alignment is enforced for this private data area.

If an empty string area is passed as dev->name, or a new structure is made, a new name string is constructed.

# alloc\_hippi\_dev

#### **Name**

alloc\_hippi\_dev — Register HIPPI device

## **Synopsis**

```
struct net_device * alloc_hippi_dev (int sizeof_priv);
```

## **Arguments**

sizeof\_priv

Size of additional driver-private structure to be allocated for this HIPPI device

## **Description**

Fill in the fields of the device structure with HIPPI-generic values.

Constructs a new net device, complete with a private data area of size <code>sizeof\_priv</code>. A 32-byte (not bit) alignment is enforced for this private data area.

# init\_trdev

#### **Name**

init\_trdev — Register token ring device

```
struct net_device * init_trdev (struct net_device * dev, int
sizeof_priv);
```

### **Arguments**

dev

A token ring device structure to be filled in, or NULL if a new struct should be allocated.

```
sizeof priv
```

Size of additional driver-private structure to be allocated for this ethernet device

## **Description**

Fill in the fields of the device structure with token ring-generic values.

If no device structure is passed, a new one is constructed, complete with a private data area of size <code>sizeof\_priv</code>. A 32-byte (not bit) alignment is enforced for this private data area.

If an empty string area is passed as dev->name, or a new structure is made, a new name string is constructed.

# alloc\_trdev

#### **Name**

alloc\_trdev — Register token ring device

```
struct net_device * alloc_trdev (int sizeof_priv);
```

## **Arguments**

```
sizeof_priv
```

Size of additional driver-private structure to be allocated for this token ring device

## **Description**

Fill in the fields of the device structure with token ring-generic values.

Constructs a new net device, complete with a private data area of size <code>sizeof\_priv</code>. A 32-byte (not bit) alignment is enforced for this private data area.

# init\_fcdev

#### Name

init\_fcdev — Register fibre channel device

```
struct net_device * init_fcdev (struct net_device * dev, int
sizeof_priv);
```

dev

A fibre channel device structure to be filled in, or NULL if a new struct should be allocated.

```
sizeof priv
```

Size of additional driver-private structure to be allocated for this ethernet device

## **Description**

Fill in the fields of the device structure with fibre channel-generic values.

If no device structure is passed, a new one is constructed, complete with a private data area of size <code>sizeof\_priv</code>. A 32-byte (not bit) alignment is enforced for this private data area.

If an empty string area is passed as dev->name, or a new structure is made, a new name string is constructed.

# alloc\_fcdev

#### **Name**

alloc\_fcdev — Register fibre channel device

```
struct net_device * alloc_fcdev (int sizeof_priv);
```

```
sizeof_priv
```

Size of additional driver-private structure to be allocated for this fibre channel device

## **Description**

Fill in the fields of the device structure with fibre channel-generic values.

Constructs a new net device, complete with a private data area of size sizeof\_priv. A 32-byte (not bit) alignment is enforced for this private data area.

# dev\_add\_pack

#### **Name**

```
dev_add_pack — add packet handler
```

## **Synopsis**

```
void dev_add_pack (struct packet_type * pt);
```

## **Arguments**

pt

packet type declaration

## **Description**

Add a protocol handler to the networking stack. The passed &packet\_type is linked into kernel lists and may not be freed until it has been removed from the kernel lists.

# dev\_remove\_pack

#### **Name**

dev\_remove\_pack — remove packet handler

## **Synopsis**

```
void dev_remove_pack (struct packet_type * pt);
```

### **Arguments**

pt

packet type declaration

# **Description**

Remove a protocol handler that was previously added to the kernel protocol handlers by dev\_add\_pack. The passed &packet\_type is removed from the kernel lists and can be freed or reused once this function returns.

## \_\_dev\_get\_by\_name

#### **Name**

```
__dev_get_by_name — find a device by its name
```

## **Synopsis**

```
struct net_device * __dev_get_by_name (const char * name);
```

### **Arguments**

name

name to find

## **Description**

Find an interface by name. Must be called under RTNL semaphore or <code>dev\_base\_lock</code>. If the name is found a pointer to the device is returned. If the name is not found then <code>NULL</code> is returned. The reference counters are not incremented so the caller must be careful with locks.

# dev\_get\_by\_name

#### **Name**

dev\_get\_by\_name — find a device by its name

```
struct net_device * dev_get_by_name (const char * name);
```

## **Arguments**

name

name to find

## **Description**

Find an interface by name. This can be called from any context and does its own locking. The returned handle has the usage count incremented and the caller must use dev\_put to release it when it is no longer needed. NULL is returned if no matching device is found.

## dev\_get

#### **Name**

dev\_get — test if a device exists

```
int dev_get (const char * name);
```

name

name to test for

## **Description**

Test if a name exists. Returns true if the name is found. In order to be sure the name is not allocated or removed during the test the caller must hold the rtnl semaphore.

This function primarily exists for back compatibility with older drivers.

# \_\_dev\_get\_by\_index

#### **Name**

\_\_dev\_get\_by\_index — find a device by its ifindex

# **Synopsis**

```
struct net_device * __dev_get_by_index (int ifindex);
```

## **Arguments**

ifindex

index of device

### **Description**

Search for an interface by index. Returns NULL if the device is not found or a pointer to the device. The device has not had its reference counter increased so the caller must be careful about locking. The caller must hold either the RTNL semaphore or  $dev\_base\_lock$ .

# dev\_get\_by\_index

#### **Name**

dev\_get\_by\_index — find a device by its ifindex

## **Synopsis**

```
struct net_device * dev_get_by_index (int ifindex);
```

## **Arguments**

ifindex

index of device

## **Description**

Search for an interface by index. Returns NULL if the device is not found or a pointer to the device. The device returned has had a reference added and the pointer is safe until the user calls dev\_put to indicate they have finished with it.

# dev\_alloc\_name

#### **Name**

dev\_alloc\_name — allocate a name for a device

## **Synopsis**

```
int dev_alloc_name (struct net_device * dev, const char * name);
```

### **Arguments**

```
dev device

name

name format string
```

## **Description**

Passed a format string - eg "ltd" it will try and find a suitable id. Not efficient for many devices, not called a lot. The caller must hold the dev\_base or rtnl lock while allocating the name and adding the device in order to avoid duplicates. Returns the number of the unit assigned or a negative errno code.

## dev\_alloc

#### **Name**

dev\_alloc — allocate a network device and name

## **Synopsis**

```
struct net_device * dev_alloc (const char * name, int * err);
```

### **Arguments**

```
name
name format string

err

error return pointer
```

## **Description**

Passed a format string, eg. "ltd", it will allocate a network device and space for the name. NULL is returned if no memory is available. If the allocation succeeds then the name is assigned and the device pointer returned. NULL is returned if the name allocation failed. The cause of an error is returned as a negative errno code in the variable *err* points to.

The caller must hold the dev\_base or RTNL locks when doing this in order to avoid duplicate name allocations.

# netdev\_state\_change

#### Name

netdev\_state\_change — device changes state

## **Synopsis**

```
void netdev_state_change (struct net_device * dev);
```

## **Arguments**

dev

device to cause notification

# **Description**

Called to indicate a device has changed state. This function calls the notifier chains for netdev\_chain and sends a NEWLINK message to the routing socket.

# dev\_load

#### **Name**

dev\_load — load a network module

```
void dev_load (const char * name);
```

## **Arguments**

name

name of interface

## **Description**

If a network interface is not present and the process has suitable privileges this function loads the module. If module loading is not available in this kernel then it becomes a nop.

# dev\_open

#### Name

dev\_open — prepare an interface for use.

```
int dev_open (struct net_device * dev);
```

dev

device to open

## **Description**

Takes a device from down to up state. The device's private open function is invoked and then the multicast lists are loaded. Finally the device is moved into the up state and a NETDEV\_UP message is sent to the netdev notifier chain.

Calling this function on an active interface is a nop. On a failure a negative errno code is returned.

## dev\_close

#### **Name**

dev\_close — shutdown an interface.

# **Synopsis**

```
int dev_close (struct net_device * dev);
```

### **Arguments**

dev

device to shutdown

### **Description**

This function moves an active device into down state. A NETDEV\_GOING\_DOWN is sent to the netdev notifier chain. The device is then deactivated and finally a NETDEV\_DOWN is sent to the notifier chain.

# register\_netdevice\_notifier

#### **Name**

register\_netdevice\_notifier — register a network notifier block

## **Synopsis**

```
int register_netdevice_notifier (struct notifier_block * nb);
```

### **Arguments**

nb

notifier

## **Description**

Register a notifier to be called when network device events occur. The notifier passed is linked into the kernel structures and must not be reused until it has been unregistered. A negative errno code is returned on a failure.

# unregister\_netdevice\_notifier

#### Name

unregister\_netdevice\_notifier — unregister a network notifier block

## **Synopsis**

```
int unregister_netdevice_notifier (struct notifier_block * nb);
```

### **Arguments**

nb

notifier

## **Description**

Unregister a notifier previously registered by register\_netdevice\_notifier. The notifier is unlinked into the kernel structures and may then be reused. A negative errno code is returned on a failure.

# dev\_queue\_xmit

#### **Name**

dev\_queue\_xmit — transmit a buffer

```
int dev_queue_xmit (struct sk_buff * skb);
```

## **Arguments**

skb

buffer to transmit

## **Description**

Queue a buffer for transmission to a network device. The caller must have set the device and priority and built the buffer before calling this function. The function can be called from an interrupt.

A negative errno code is returned on a failure. A success does not guarantee the frame will be transmitted as it may be dropped due to congestion or traffic shaping.

## netif\_rx

#### **Name**

netif\_rx — post buffer to the network code

```
int netif_rx (struct sk_buff * skb);
```

skb

buffer to post

## **Description**

This function receives a packet from a device driver and queues it for the upper (protocol) levels to process. It always succeeds. The buffer may be dropped during processing for congestion control or by the protocol layers.

#### return values

NET\_RX\_SUCCESS (no congestion) NET\_RX\_CN\_LOW (low congestion) NET\_RX\_CN\_MOD (moderate congestion) NET\_RX\_CN\_HIGH (high congestion) NET\_RX\_DROP (packet was dropped)

# net\_call\_rx\_atomic

#### **Name**

```
net_call_rx_atomic —
```

```
void net_call_rx_atomic (void (*fn) (void));
```

fn

function to call

## **Description**

Make a function call that is atomic with respect to the protocol layers.

# register\_gifconf

#### Name

register\_gifconf — register a SIOCGIF handler

# **Synopsis**

```
int register_gifconf (unsigned int family, gifconf_func_t *
gifconf);
```

## **Arguments**

family

Address family

gifconf

Function handler

## **Description**

Register protocol dependent address dumping routines. The handler that is passed must not be freed or reused until it has been replaced by another handler.

## netdev\_set\_master

#### **Name**

```
netdev_set_master — set up master/slave pair
```

## **Synopsis**

```
int netdev_set_master (struct net_device * slave, struct
net_device * master);
```

## **Arguments**

```
slave device

master

new master device
```

# **Description**

Changes the master device of the slave. Pass NULL to break the bonding. The caller must hold the RTNL semaphore. On a failure a negative errno code is returned. On

success the reference counts are adjusted, RTM\_NEWLINK is sent to the routing socket and the function returns zero.

# dev\_set\_promiscuity

#### Name

dev\_set\_promiscuity — update promiscuity count on a device

## **Synopsis**

```
void dev_set_promiscuity (struct net_device * dev, int inc);
```

## **Arguments**

```
dev
device
inc
modifier
```

## **Description**

Add or remove promsicuity from a device. While the count in the device remains above zero the interface remains promiscuous. Once it hits zero the device reverts back to normal filtering operation. A negative inc value is used to drop promiscuity on the device.

# dev\_set\_allmulti

#### Name

dev\_set\_allmulti — update allmulti count on a device

## **Synopsis**

```
void dev_set_allmulti (struct net_device * dev, int inc);
```

### **Arguments**

```
dev
device
inc
modifier
```

## **Description**

Add or remove reception of all multicast frames to a device. While the count in the device remains above zero the interface remains listening to all interfaces. Once it hits zero the device reverts back to normal filtering operation. A negative *inc* value is used to drop the counter when releasing a resource needing all multicasts.

# dev\_ioctl

#### **Name**

```
dev_ioctl — network device ioctl
```

## **Synopsis**

```
int dev_ioctl (unsigned int cmd, void * arg);
```

## **Arguments**

```
cmd
command to issue

arg

pointer to a struct ifreq in user space
```

## **Description**

Issue ioctl functions to devices. This is normally called by the user space syscall interfaces but can sometimes be useful for other purposes. The return value is the return from the syscall if positive or a negative errno code on error.

## dev\_new\_index

#### **Name**

dev\_new\_index — allocate an ifindex

## **Synopsis**

```
int dev_new_index ( void);
```

## **Arguments**

void

no arguments

## **Description**

Returns a suitable unique value for a new device interface number. The caller must hold the rtnl semaphore or the dev\_base\_lock to be sure it remains unique.

# netdev\_finish\_unregister

#### Name

netdev\_finish\_unregister — complete unregistration

```
int netdev_finish_unregister (struct net_device * dev);
```

## **Arguments**

dev device

## **Description**

Destroy and free a dead device. A value of zero is returned on success.

# unregister\_netdevice

#### Name

unregister\_netdevice — remove device from the kernel

```
int unregister_netdevice (struct net_device * dev);
```

dev device

## **Description**

This function shuts down a device interface and removes it from the kernel tables. On success 0 is returned, on a failure a negative errno code is returned.

Callers must hold the rtnl semaphore. See the comment at the end of Space.c for details about the locking. You may want unregister\_netdev instead of this.

## 8.2. 8390 Based Network Cards

# ei\_open

#### **Name**

ei\_open — Open/initialize the board.

```
int ei_open (struct net_device * dev);
```

dev

network device to initialize

## **Description**

This routine goes all-out, setting everything up anew at each open, even though many of these registers should only need to be set once at boot.

# ei\_close

#### Name

ei\_close — shut down network device

## **Synopsis**

```
int ei_close (struct net_device * dev);
```

## **Arguments**

dev

network device to close

## **Description**

Opposite of ei\_open. Only used when "ifconfig <devname> down" is done.

# ei\_tx\_timeout

#### Name

ei\_tx\_timeout — handle transmit time out condition

## **Synopsis**

```
void ei_tx_timeout (struct net_device * dev);
```

## **Arguments**

dev

network device which has apparently fallen asleep

## **Description**

Called by kernel when device never acknowledges a transmit has completed (or failed) - i.e. never posted a Tx related interrupt.

# ei\_interrupt

#### Name

ei\_interrupt — handle the interrupts from an 8390

## **Synopsis**

```
void ei_interrupt (int irq, void * dev_id, struct pt_regs *
regs);
```

## **Arguments**

```
irq
    interrupt number

dev_id
    a pointer to the net_device

regs
    unused
```

## **Description**

Handle the ether interface interrupts. We pull packets from the 8390 via the card specific functions and fire them at the networking stack. We also handle transmit completions and wake the transmit path if neccessary. We also update the counters and do other housekeeping as needed.

# ethdev\_init

#### Name

ethdev\_init — init rest of 8390 device struct

## **Synopsis**

```
int ethdev_init (struct net_device * dev);
```

## **Arguments**

dev

network device structure to init

# **Description**

Initialize the rest of the 8390 device structure. Do NOT \_\_init this, as it is used by 8390 based modular drivers too.

## **NS8390\_init**

#### **Name**

NS8390\_init — initialize 8390 hardware

```
void NS8390_init (struct net_device * dev, int startp);
```

## **Arguments**

```
network device to initialize
startp
boolean. non-zero value to initiate chip processing
```

## **Description**

Must be called with lock held.

# 8.3. Synchronous PPP

# sppp\_input

#### **Name**

sppp\_input — receive and process a WAN PPP frame

```
void sppp_input (struct net_device * dev, struct sk_buff * skb);
```

## **Arguments**

dev

The device it arrived on

skb

The buffer to process

## **Description**

This can be called directly by cards that do not have timing constraints but is normally called from the network layer after interrupt servicing to process frames queued via netif\_rx.

We process the options in the card. If the frame is destined for the protocol stacks then it requeues the frame for the upper level protocol. If it is a control from it is processed and discarded here.

# sppp\_close

#### Name

sppp\_close — close down a synchronous PPP or Cisco HDLC link

### **Synopsis**

```
int sppp_close (struct net_device * dev);
```

#### **Arguments**

dev

The network device to drop the link of

## **Description**

This drops the logical interface to the channel. It is not done politely as we assume we will also be dropping DTR. Any timeouts are killed.

## sppp\_open

#### Name

sppp\_open — open a synchronous PPP or Cisco HDLC link

```
int sppp_open (struct net_device * dev);
```

dev

Network device to activate

# **Description**

Close down any existing synchronous session and commence from scratch. In the PPP case this means negotiating LCP/IPCP and friends, while for Cisco HDLC we simply need to start sending keepalives

## sppp\_reopen

#### **Name**

sppp\_reopen — notify of physical link loss

### **Synopsis**

```
int sppp_reopen (struct net_device * dev);
```

#### **Arguments**

dev

Device that lost the link

This function informs the synchronous protocol code that the underlying link died (for example a carrier drop on X.21)

We increment the magic numbers to ensure that if the other end failed to notice we will correctly start a new session. It happens do to the nature of telco circuits is that you can lose carrier on one endonly.

Having done this we go back to negotiating. This function may be called from an interrupt context.

## sppp\_change\_mtu

#### Name

```
sppp_change_mtu — Change the link MTU
```

## **Synopsis**

```
int sppp_change_mtu (struct net_device * dev, int new_mtu);
```

#### **Arguments**

```
dev
    Device to change MTU on
new_mtu
New MTU
```

Change the MTU on the link. This can only be called with the link down. It returns an error if the link is up or the mtu is out of range.

# sppp\_do\_ioctl

#### **Name**

```
sppp_do_ioctl — Ioctl handler for ppp/hdlc
```

## **Synopsis**

```
int sppp_do_ioctl (struct net_device * dev, struct ifreq * ifr,
int cmd);
```

### **Arguments**

dev

Device subject to ioctl

ifr

Interface request block from the user

cmd

Command that is being issued

This function handles the ioctls that may be issued by the user to control the settings of a PPP/HDLC link. It does both busy and security checks. This function is intended to be wrapped by callers who wish to add additional ioctl calls of their own.

# sppp\_attach

#### **Name**

sppp\_attach — attach synchronous PPP/HDLC to a device

## **Synopsis**

```
void sppp_attach (struct ppp_device * pd);
```

#### **Arguments**

pd

PPP device to initialise

#### **Description**

This initialises the PPP/HDLC support on an interface. At the time of calling the dev element must point to the network device that this interface is attached to. The interface should not yet be registered.

# sppp\_detach

#### Name

sppp\_detach — release PPP resources from a device

## **Synopsis**

```
void sppp_detach (struct net_device * dev);
```

### **Arguments**

dev

Network device to release

# **Description**

Stop and free up any PPP/HDLC resources used by this interface. This must be called before the device is freed.

# **Chapter 9. Module Support**

## 9.1. Module Loading

## request\_module

#### Name

request\_module — try to load a kernel module

### **Synopsis**

```
int request_module (const char * module_name);
```

#### **Arguments**

module\_name

Name of module

#### **Description**

Load a module using the user mode module loader. The function returns zero on success or a negative errno code on failure. Note that a successful module load does not mean the module did not then unload and exit on an error of its own. Callers must check that the service they requested is now available not blindly invoke it.

If module auto-loading support is disabled then this function becomes a no-operation.

## call\_usermodehelper

#### **Name**

call\_usermodehelper — start a usermode application

### **Synopsis**

```
int call_usermodehelper (char * path, char ** argv, char **
envp);
```

### **Arguments**

```
path

pathname for the application

argv

null-terminated argument list

envp

null-terminated environment list
```

### **Description**

Runs a user-space application. The application is started asynchronously. It runs as a child of keventd. It runs with full root capabilities, keventd silently reaps the child when it exits.

Must be called from process context. Returns zero on success, else a negative error code.

# 9.2. Inter Module support

# inter\_module\_register

#### Name

inter\_module\_register — register a new set of inter module data.

#### **Synopsis**

```
void inter_module_register (const char * im_name, struct module
* owner, const void * userdata);
```

#### **Arguments**

```
im_name
    an arbitrary string to identify the data, must be unique

owner
    module that is registering the data, always use THIS_MODULE

userdata
    pointer to arbitrary userdata to be registered
```

#### **Description**

Check that the im\_name has not already been registered, complain if it has. For new data, add it to the inter\_module\_entry list.

# inter\_module\_unregister

#### **Name**

inter\_module\_unregister — unregister a set of inter module data.

#### **Synopsis**

```
void inter_module_unregister (const char * im_name);
```

#### **Arguments**

im\_name

an arbitrary string to identify the data, must be unique

## **Description**

Check that the im\_name has been registered, complain if it has not. For existing data, remove it from the inter\_module\_entry list.

# inter\_module\_get

#### **Name**

inter\_module\_get — return arbitrary userdata from another module.

### **Synopsis**

```
const void * inter_module_get (const char * im_name);
```

### **Arguments**

im name

an arbitrary string to identify the data, must be unique

### **Description**

If the im\_name has not been registered, return NULL. Try to increment the use count on the owning module, if that fails then return NULL. Otherwise return the userdata.

# inter\_module\_get\_request

#### Name

inter\_module\_get\_request — im get with automatic request\_module.

```
const void * inter_module_get_request (const char * im_name,
const char * modname);
```

```
im_name
    an arbitrary string to identify the data, must be unique
modname
    module that is expected to register im_name
```

### **Description**

If inter\_module\_get fails, do request\_module then retry.

# inter\_module\_put

#### **Name**

inter\_module\_put — release use of data from another module.

## **Synopsis**

```
void inter_module_put (const char * im_name);
```

### **Arguments**

```
im_name
an arbitrary string to identify the data, must be unique
```

If the im\_name has not been registered, complain, otherwise decrement the use count on the owning module.

# **Chapter 10. Hardware Interfaces**

## 10.1. Interrupt Handling

## disable\_irq\_nosync

#### **Name**

disable\_irq\_nosync — disable an irq without waiting

### **Synopsis**

```
void disable_irq_nosync (unsigned int irq);
```

#### **Arguments**

irq

Interrupt to disable

#### **Description**

Disable the selected interrupt line. Disables and Enables are nested. Unlike disable\_irq, this function does not ensure existing instances of the IRQ handler have completed before returning.

This function may be called from IRQ context.

## disable\_irq

#### **Name**

disable\_irq — disable an irq and wait for completion

#### **Synopsis**

```
void disable_irq (unsigned int irq);
```

#### **Arguments**

irq

Interrupt to disable

## **Description**

Disable the selected interrupt line. Enables and Disables are nested. This function waits for any pending IRQ handlers for this interrupt to complete before returning. If you use this function while holding a resource the IRQ handler may need you will deadlock.

This function may be called - with care - from IRQ context.

# enable\_irq

#### **Name**

enable\_irq — enable handling of an irq

### **Synopsis**

```
void enable_irq (unsigned int irq);
```

#### **Arguments**

irq

Interrupt to enable

### **Description**

Undoes the effect of one call to disable\_irq. If this matches the last disable, processing of interrupts on this IRQ line is re-enabled.

This function may be called from IRQ context.

# probe\_irq\_mask

#### Name

probe\_irq\_mask — scan a bitmap of interrupt lines

```
unsigned int probe_irq_mask (unsigned long val);
```

val

mask of interrupts to consider

## **Description**

Scan the ISA bus interrupt lines and return a bitmap of active interrupts. The interrupt probe logic state is then returned to its previous value.

#### **Note**

we need to scan all the irq's even though we will only return ISA irq numbers - just so that we reset them all to a known state.

# 10.2. MTRR Handling

### mtrr\_add

#### **Name**

mtrr\_add — Add a memory type region

# **Synopsis**

int mtrr\_add (unsigned long base, unsigned long size, unsigned
int type, char increment);

```
base
```

Physical base address of region

size

Physical size of region

type

Type of MTRR desired

increment

If this is true do usage counting on the region

## **Description**

Memory type region registers control the caching on newer Intel and non Intel processors. This function allows drivers to request an MTRR is added. The details and hardware specifics of each processor's implementation are hidden from the caller, but nevertheless the caller should expect to need to provide a power of two size on an equivalent power of two boundary.

If the region cannot be added either because all regions are in use or the CPU cannot support it a negative value is returned. On success the register number for this entry is returned, but should be treated as a cookie only.

On a multiprocessor machine the changes are made to all processors. This is required on x86 by the Intel processors.

The available types are

MTRR\_TYPE\_UNCACHABLE - No caching

MTRR TYPE WRBACK - Write data back in bursts whenever

MTRR TYPE WRCOMB - Write data back soon but allow bursts

MTRR\_TYPE\_WRTHROUGH - Cache reads but not writes

#### **BUGS**

Needs a quiet flag for the cases where drivers do not mind failures and do not wish system log messages to be sent.

## mtrr\_del

#### **Name**

mtrr\_del — delete a memory type region

## **Synopsis**

```
int mtrr_del (int reg, unsigned long base, unsigned long size);
```

### **Arguments**

```
reg
Register returned by mtrr_add
base
Physical base address
size
Size of region
```

## **Description**

If register is supplied then base and size are ignored. This is how drivers should call it.

Releases an MTRR region. If the usage count drops to zero the register is freed and the region returns to default state. On success the register is returned, on failure a negative error code.

# 10.3. PCI Support Library

# pci\_find\_slot

#### **Name**

pci\_find\_slot — locate PCI device from a given PCI slot

### **Synopsis**

```
struct pci_dev * pci_find_slot (unsigned int bus, unsigned int
devfn);
```

#### **Arguments**

bus

number of PCI bus on which desired PCI device resides

devfn

encodes number of PCI slot in which the desired PCI device resides and the logical device number within that slot in case of multi-function devices.

Given a PCI bus and slot/function number, the desired PCI device is located in system global list of PCI devices. If the device is found, a pointer to its data structure is returned. If no device is found, NULL is returned.

# pci\_find\_subsys

#### **Name**

pci\_find\_subsys — begin or continue searching for a PCI device by vendor/subvendor/device/subdevice id

## **Synopsis**

```
struct pci_dev * pci_find_subsys (unsigned int vendor, unsigned
int device, unsigned int ss_vendor, unsigned int ss_device,
const struct pci_dev * from);
```

#### **Arguments**

vendor

PCI vendor id to match, or PCI\_ANY\_ID to match all vendor ids

device

PCI device id to match, or PCI\_ANY\_ID to match all device ids

ss\_vendor

PCI subsystem vendor id to match, or PCI\_ANY\_ID to match all vendor ids

ss\_device

PCI subsystem device id to match, or PCI\_ANY\_ID to match all device ids

from

Previous PCI device found in search, or NULL for new search.

#### **Description**

Iterates through the list of known PCI devices. If a PCI device is found with a matching <code>vendor</code>, <code>device</code>, <code>ss\_vendor</code> and <code>ss\_device</code>, a pointer to its device structure is returned. Otherwise, <code>NULL</code> is returned. A new search is initiated by passing <code>NULL</code> to the <code>from</code> argument. Otherwise if <code>from</code> is not <code>NULL</code>, searches continue from next device on the global list.

# pci\_find\_device

#### Name

pci\_find\_device — begin or continue searching for a PCI device by vendor/device id

#### **Synopsis**

struct pci\_dev \* pci\_find\_device (unsigned int vendor, unsigned
int device, const struct pci\_dev \* from);

```
vendor
```

PCI vendor id to match, or PCI\_ANY\_ID to match all vendor ids

device

PCI device id to match, or PCI\_ANY\_ID to match all device ids

from

Previous PCI device found in search, or NULL for new search.

#### **Description**

Iterates through the list of known PCI devices. If a PCI device is found with a matching <code>vendor</code> and <code>device</code>, a pointer to its device structure is returned. Otherwise, <code>NULL</code> is returned. A new search is initiated by passing <code>NULL</code> to the <code>from</code> argument. Otherwise if <code>from</code> is not <code>NULL</code>, searches continue from next device on the global list.

## pci\_find\_class

#### **Name**

pci\_find\_class — begin or continue searching for a PCI device by class

```
struct pci_dev * pci_find_class (unsigned int class, const
struct pci_dev * from);
```

class

search for a PCI device with this class designation

from

Previous PCI device found in search, or NULL for new search.

## **Description**

Iterates through the list of known PCI devices. If a PCI device is found with a matching class, a pointer to its device structure is returned. Otherwise, NULL is returned. A new search is initiated by passing NULL to the from argument. Otherwise if from is not NULL, searches continue from next device on the global list.

# pci\_find\_capability

#### **Name**

pci\_find\_capability — query for devices' capabilities

```
int pci_find_capability (struct pci_dev * dev, int cap);
```

```
dev
PCI device to query

cap

capability code
```

#### **Description**

Tell if a device supports a given PCI capability. Returns the address of the requested capability structure within the device's PCI configuration space or 0 in case the device does not support it. Possible values for *cap*:

```
PCI_CAP_ID_PM Power Management

PCI_CAP_ID_AGP Accelerated Graphics Port

PCI_CAP_ID_VPD Vital Product Data

PCI_CAP_ID_SLOTID Slot Identification

PCI_CAP_ID_MSI Message Signalled Interrupts

PCI_CAP_ID_CHSWP CompactPCI HotSwap
```

## pci\_find\_parent\_resource

#### Name

pci\_find\_parent\_resource — return resource region of parent bus of given region

### **Synopsis**

```
struct resource * pci_find_parent_resource (const struct pci_dev
* dev, struct resource * res);
```

### **Arguments**

dev

PCI device structure contains resources to be searched

res

child resource record for which parent is sought

### **Description**

For given resource region of given device, return the resource region of parent bus the given region is contained in or where it should be allocated from.

## pci\_set\_power\_state

#### **Name**

pci\_set\_power\_state — Set the power state of a PCI device

```
int pci_set_power_state (struct pci_dev * dev, int state);
```

dev

PCI device to be suspended

state

Power state we're entering

### **Description**

Transition a device to a new power state, using the Power Management Capabilities in the device's config space.

#### **RETURN VALUE**

-EINVAL if trying to enter a lower state than we're already in. 0 if we're already in the requested state. -EIO if device does not support PCI PM. 0 if we can successfully change the power state.

## pci\_save\_state

#### **Name**

pci\_save\_state — save the PCI configuration space of a device before suspending

```
int pci_save_state (struct pci_dev * dev, u32 * buffer);
```

dev

- PCI device that we're dealing with

buffer

- buffer to hold config space context

### **Description**

buffer must be large enough to hold the entire PCI 2.2 config space (>= 64 bytes).

# pci\_restore\_state

#### **Name**

pci\_restore\_state — Restore the saved state of a PCI device

### **Synopsis**

```
int pci_restore_state (struct pci_dev * dev, u32 * buffer);
```

### **Arguments**

dev

- PCI device that we're dealing with

buffer

- saved PCI config space

# pci\_enable\_device

#### Name

pci\_enable\_device — Initialize device before it's used by a driver.

## **Synopsis**

```
int pci_enable_device (struct pci_dev * dev);
```

### **Arguments**

dev

PCI device to be initialized

### **Description**

Initialize device before it's used by a driver. Ask low-level code to enable I/O and memory. Wake up the device if it was suspended. Beware, this function can fail.

# pci\_disable\_device

#### Name

pci\_disable\_device — Disable PCI device after use

#### **Synopsis**

```
void pci_disable_device (struct pci_dev * dev);
```

#### **Arguments**

dev

PCI device to be disabled

### **Description**

Signal to the system that the PCI device is not in use by the system anymore. This only involves disabling PCI bus-mastering, if active.

# pci\_enable\_wake

#### Name

pci\_enable\_wake — enable device to generate PME# when suspended

## **Synopsis**

```
int pci_enable_wake (struct pci_dev * dev, u32 state, int
enable);
```

### **Arguments**

dev

- PCI device to operate on

state

- Current state of device.

enable

- Flag to enable or disable generation

### **Description**

Set the bits in the device's PM Capabilities to generate PME# when the system is suspended.

-EIO is returned if device doesn't have PM Capabilities. -EINVAL is returned if device supports it, but can't generate wake events. 0 if operation is successful.

# pci\_release\_regions

#### **Name**

pci\_release\_regions — Release reserved PCI I/O and memory resources

### **Synopsis**

```
void pci_release_regions (struct pci_dev * pdev);
```

#### **Arguments**

pdev

PCI device whose resources were previously reserved by pci\_request\_regions

### **Description**

Releases all PCI I/O and memory resources previously reserved by a successful call to pci\_request\_regions. Call this function only after all use of the PCI regions has ceased.

# pci\_request\_regions

#### Name

pci\_request\_regions — Reserved PCI I/O and memory resources

```
int pci_request_regions (struct pci_dev * pdev, char *
res_name);
```

pdev

PCI device whose resources are to be reserved

res name

Name to be associated with resource.

### **Description**

Mark all PCI regions associated with PCI device *pdev* as being reserved by owner *res\_name*. Do not access any address inside the PCI regions unless this call returns successfully.

Returns 0 on success, or EBUSY on error. A warning message is also printed on failure.

# pci\_match\_device

#### **Name**

pci\_match\_device — Tell if a PCI device structure has a matching PCI device id structure

```
const struct pci_device_id * pci_match_device (const struct
pci_device_id * ids, const struct pci_dev * dev);
```

```
ids array of PCI device id structures to search in dev the PCI device structure to match against
```

### **Description**

Used by a driver to check whether a PCI device present in the system is in its list of supported devices. Returns the matching pci\_device\_id structure or NULL if there is no match.

# pci\_register\_driver

#### **Name**

```
pci_register_driver — register a new pci driver
```

### **Synopsis**

```
int pci_register_driver (struct pci_driver * drv);
```

### **Arguments**

drv

the driver structure to register

Adds the driver structure to the list of registered drivers Returns the number of pci devices which were claimed by the driver during registration. The driver remains registered even if the return value is zero.

# pci\_unregister\_driver

#### **Name**

pci\_unregister\_driver — unregister a pci driver

### **Synopsis**

```
void pci_unregister_driver (struct pci_driver * drv);
```

#### **Arguments**

drv

the driver structure to unregister

#### **Description**

Deletes the driver structure from the list of registered PCI drivers, gives it a chance to clean up by calling its remove function for each device it was responsible for, and marks those devices as driverless.

# pci\_insert\_device

#### Name

pci\_insert\_device — insert a hotplug device

### **Synopsis**

```
void pci_insert_device (struct pci_dev * dev, struct pci_bus *
bus);
```

### **Arguments**

```
dev
the device to insert
bus
where to insert it
```

### **Description**

Add a new device to the device lists and notify userspace (/sbin/hotplug).

# pci\_remove\_device

#### **Name**

pci\_remove\_device — remove a hotplug device

## **Synopsis**

```
void pci_remove_device (struct pci_dev * dev);
```

### **Arguments**

dev

the device to remove

## **Description**

Delete the device structure from the device lists and notify userspace (/sbin/hotplug).

# pci\_dev\_driver

### Name

```
pci_dev_driver — get the pci_driver of a device
```

## **Synopsis**

```
struct pci_driver * pci_dev_driver (const struct pci_dev * dev);
```

## **Arguments**

dev

the device to query

# **Description**

Returns the appropriate pci\_driver structure or NULL if there is no registered driver for the device.

# pci\_set\_master

#### Name

pci\_set\_master — enables bus-mastering for device dev

## **Synopsis**

```
void pci_set_master (struct pci_dev * dev);
```

## **Arguments**

dev

the PCI device to enable

Enables bus-mastering on the device and calls pcibios\_set\_master to do the needed arch specific settings.

# pci\_setup\_device

#### **Name**

pci\_setup\_device — fill in class and map information of a device

## **Synopsis**

```
int pci_setup_device (struct pci_dev * dev);
```

### **Arguments**

dev

the device structure to fill

### **Description**

Initialize the device structure with information about the device's vendor, class, memory and IO-space addresses, IRQ lines etc. Called at initialisation of the PCI subsystem and by CardBus services. Returns 0 on success and -1 if unknown type of device (not normal, bridge or CardBus).

# pci\_pool\_create

#### **Name**

pci\_pool\_create — Creates a pool of pci consistent memory blocks, for dma.

### **Synopsis**

```
struct pci_pool * pci_pool_create (const char * name, struct
pci_dev * pdev, size_t size, size_t align, size_t allocation,
int flags);
```

```
name of pool, for diagnostics

pdev

pci device that will be doing the DMA

size

size of the blocks in this pool.

align

alignment requirement for blocks; must be a power of two

allocation

returned blocks won't cross this boundary (or zero)

flags

SLAB_* flags (not all are supported).
```

Returns a pci allocation pool with the requested characteristics, or null if one can't be created. Given one of these pools, pci\_pool\_alloc may be used to allocate memory. Such memory will all have "consistent" DMA mappings, accessible by the device and its driver without using cache flushing primitives. The actual size of blocks allocated may be larger than requested because of alignment.

If allocation is nonzero, objects returned from pci\_pool\_alloc won't cross that size boundary. This is useful for devices which have addressing restrictions on individual DMA transfers, such as not crossing boundaries of 4KBytes.

# pci\_pool\_destroy

#### **Name**

pci\_pool\_destroy — destroys a pool of pci memory blocks.

### **Synopsis**

```
void pci_pool_destroy (struct pci_pool * pool);
```

### **Arguments**

pool

pci pool that will be destroyed

Caller guarantees that no more memory from the pool is in use, and that nothing will try to use the pool after this call.

# pci\_pool\_alloc

#### **Name**

```
pci_pool_alloc — get a block of consistent memory
```

## **Synopsis**

```
void * pci_pool_alloc (struct pci_pool * pool, int mem_flags,
dma_addr_t * handle);
```

```
pci pool that will produce the block

mem_flags

SLAB_KERNEL or SLAB_ATOMIC

handle

pointer to dma address of block
```

This returns the kernel virtual address of a currently unused block, and reports its dma address through the handle. If such a memory block can't be allocated, null is returned.

# pci\_pool\_free

#### **Name**

```
pci_pool_free — put block back into pci pool
```

## **Synopsis**

```
void pci_pool_free (struct pci_pool * pool, void * vaddr,
dma_addr_t dma);
```

```
the pci pool holding the block

vaddr

virtual address of block

dma

dma address of block
```

Caller promises neither device nor driver will again touch this block unless it is first re-allocated.

### 10.4. MCA Architecture

### 10.4.1. MCA Device Functions

# mca\_find\_adapter

#### **Name**

```
mca_find_adapter — scan for adapters
```

## **Synopsis**

```
int mca_find_adapter (int id, int start);
```

```
id
     MCA identification to search for
start
     starting slot
```

Search the MCA configuration for adapters matching the 16bit ID given. The first time it should be called with start as zero and then further calls made passing the return value of the previous call until MCA\_NOTFOUND is returned.

Disabled adapters are not reported.

## mca\_find\_unused\_adapter

#### **Name**

mca\_find\_unused\_adapter — scan for unused adapters

## **Synopsis**

```
int mca_find_unused_adapter (int id, int start);
```

### **Arguments**

```
id
     MCA identification to search for
start
    starting slot
```

### **Description**

Search the MCA configuration for adapters matching the 16bit ID given. The first time it should be called with start as zero and then further calls made passing the return

value of the previous call until MCA\_NOTFOUND is returned.

Adapters that have been claimed by drivers and those that are disabled are not reported. This function thus allows a driver to scan for further cards when some may already be driven.

## mca\_read\_stored\_pos

#### **Name**

mca\_read\_stored\_pos — read POS register from boot data

### **Synopsis**

```
unsigned char mca_read_stored_pos (int slot, int reg);
```

### **Arguments**

```
slot
slot number to read from
reg
register to read from
```

### **Description**

Fetch a POS value that was stored at boot time by the kernel when it scanned the MCA space. The register value is returned. Missing or invalid registers report 0.

## mca\_read\_pos

#### **Name**

mca\_read\_pos — read POS register from card

## **Synopsis**

```
unsigned char mca_read_pos (int slot, int reg);
```

## **Arguments**

```
slot
slot number to read from
reg
register to read from
```

## **Description**

Fetch a POS value directly from the hardware to obtain the current value. This is much slower than mca\_read\_stored\_pos and may not be invoked from interrupt context. It handles the deep magic required for onboard devices transparently.

## mca\_write\_pos

#### **Name**

```
mca_write_pos — read POS register from card
```

### **Synopsis**

```
void mca_write_pos (int slot, int reg, unsigned char byte);
```

### **Arguments**

```
slot
slot number to read from

reg
register to read from

byte
byte to write to the POS registers
```

### **Description**

Store a POS value directly from the hardware. You should not normally need to use this function and should have a very good knowledge of MCA bus before you do so. Doing this wrongly can damage the hardware.

This function may not be used from interrupt context.

Note that this a technically a Bad Thing, as IBM tech stuff says you should only set POS values through their utilities. However, some devices such as the 3c523 recommend that you write back some data to make sure the configuration is consistent. I'd say that IBM is right, but I like my drivers to work.

This function can't do checks to see if multiple devices end up with the same resources, so you might see magic smoke if someone screws up.

# mca\_set\_adapter\_name

#### Name

mca\_set\_adapter\_name — Set the description of the card

## **Synopsis**

```
void mca_set_adapter_name (int slot, char* name);
```

## **Arguments**

```
slot
slot to name

name
text string for the namen
```

### **Description**

This function sets the name reported via /proc for this adapter slot. This is for user information only. Setting a name deletes any previous name.

## mca\_set\_adapter\_procfn

#### **Name**

mca\_set\_adapter\_procfn — Set the /proc callback

### **Synopsis**

```
void mca_set_adapter_procfn (int slot, MCA_ProcFn procfn, void*
dev);
```

### **Arguments**

```
slot
slot to configure

procfn
callback function to call for /proc

dev
device information passed to the callback
```

### **Description**

This sets up an information callback for /proc/mca/slot?. The function is called with the buffer, slot, and device pointer (or some equally informative context information, or nothing, if you prefer), and is expected to put useful information into the buffer. The adapter name, ID, and POS registers get printed before this is called though, so don't do it again.

This should be called with a NULL *procfn* when a module unregisters, thus preventing kernel crashes and other such nastiness.

# mca\_is\_adapter\_used

#### Name

mca\_is\_adapter\_used — check if claimed by driver

## **Synopsis**

```
int mca_is_adapter_used (int slot);
```

## **Arguments**

slot

slot to check

## **Description**

Returns 1 if the slot has been claimed by a driver

## mca\_mark\_as\_used

#### **Name**

mca\_mark\_as\_used — claim an MCA device

## **Synopsis**

```
int mca_mark_as_used (int slot);
```

## **Arguments**

slot

slot to claim

#### **FIXME**

should we make this threadsafe

Claim an MCA slot for a device driver. If the slot is already taken the function returns 1, if it is not taken it is claimed and 0 is returned.

## mca\_mark\_as\_unused

### Name

mca\_mark\_as\_unused — release an MCA device

## **Synopsis**

```
void mca_mark_as_unused (int slot);
```

## **Arguments**

slot slot to claim

## **Description**

Release the slot for other drives to use.

# mca\_get\_adapter\_name

## Name

mca\_get\_adapter\_name — get the adapter description

## **Synopsis**

```
char * mca_get_adapter_name (int slot);
```

```
slot slot to query
```

Return the adapter description if set. If it has not been set or the slot is out range then return NULL.

# mca\_isadapter

#### **Name**

mca\_isadapter — check if the slot holds an adapter

## **Synopsis**

```
int mca_isadapter (int slot);
```

## **Arguments**

```
slot slot to query
```

## **Description**

Returns zero if the slot does not hold an adapter, non zero if it does.

# mca\_isenabled

#### Name

mca\_isenabled — check if the slot holds an adapter

## **Synopsis**

```
int mca_isenabled (int slot);
```

## **Arguments**

```
slot slot to query
```

## **Description**

Returns a non zero value if the slot holds an enabled adapter and zero for any other case.

#### 10.4.2. MCA Bus DMA

# mca\_enable\_dma

### **Name**

mca\_enable\_dma — channel to enable DMA on

## **Synopsis**

void mca\_enable\_dma (unsigned int dmanr);

## **Arguments**

dmanr

DMA channel

### **Description**

Enable the MCA bus DMA on a channel. This can be called from IRQ context.

# mca\_disable\_dma

#### **Name**

mca\_disable\_dma — channel to disable DMA on

## **Synopsis**

```
void mca_disable_dma (unsigned int dmanr);
```

### **Arguments**

dmanr

DMA channel

## **Description**

Enable the MCA bus DMA on a channel. This can be called from IRQ context.

# mca\_set\_dma\_addr

### **Name**

```
mca_set_dma_addr — load a 24bit DMA address
```

# **Synopsis**

```
void mca_set_dma_addr (unsigned int dmanr, unsigned int a);
```

## **Arguments**

```
dmanr
```

DMA channel

а

24bit bus address

## **Description**

Load the address register in the DMA controller. This has a 24bit limitation (16Mb).

# mca\_get\_dma\_addr

#### **Name**

```
mca_get_dma_addr — load a 24bit DMA address
```

## **Synopsis**

```
unsigned int mca_get_dma_addr (unsigned int dmanr);
```

## **Arguments**

dmanr

DMA channel

Read the address register in the DMA controller. This has a 24bit limitation (16Mb). The return is a bus address.

## mca\_set\_dma\_count

#### **Name**

mca\_set\_dma\_count — load a 16bit transfer count

## **Synopsis**

```
void mca_set_dma_count (unsigned int dmanr, unsigned int count);
```

### **Arguments**

dmanr

DMA channel

count

count

## **Description**

Set the DMA count for this channel. This can be up to 64Kbytes. Setting a count of zero will not do what you expect.

# mca\_get\_dma\_residue

#### Name

mca\_get\_dma\_residue — get the remaining bytes to transfer

## **Synopsis**

```
unsigned int mca_get_dma_residue (unsigned int dmanr);
```

## **Arguments**

dmanr

DMA channel

# **Description**

This function returns the number of bytes left to transfer on this DMA channel.

# mca\_set\_dma\_io

#### **Name**

mca\_set\_dma\_io — set the port for an I/O transfer

## **Synopsis**

```
void mca_set_dma_io (unsigned int dmanr, unsigned int io_addr);
```

### **Arguments**

```
dmanr

DMA channel

io_addr

an I/O port number
```

## **Description**

Unlike the ISA bus DMA controllers the DMA on MCA bus can transfer with an I/O port target.

# mca\_set\_dma\_mode

#### Name

```
mca_set_dma_mode — set the DMA mode
```

## **Synopsis**

```
void mca_set_dma_mode (unsigned int dmanr, unsigned int mode);
```

## **Arguments**

dmanr

DMA channel

mode

mode to set

## **Description**

The DMA controller supports several modes. The mode values you can

#### set are

MCA\_DMA\_MODE\_READ when reading from the DMA device.

MCA\_DMA\_MODE\_WRITE to writing to the DMA device.

 $\mbox{MCA\_DMA\_MODE\_IO}$  to do DMA to or from an I/O port.

MCA\_DMA\_MODE\_16 to do 16bit transfers.

# Chapter 11. The Device File System

## devfs\_register

#### **Name**

devfs\_register — Register a device entry.

## **Synopsis**

```
devfs_handle_t devfs_register (devfs_handle_t dir, const char *
name, unsigned int flags, unsigned int major, unsigned int
minor, umode_t mode, void * ops, void * info);
```

### **Arguments**

dir

The handle to the parent devfs directory entry. If this is NULL the new name is relative to the root of the devfs.

name

The name of the entry.

flags

A set of bitwise-ORed flags (DEVFS\_FL\_\*).

major

The major number. Not needed for regular files.

minor

The minor number. Not needed for regular files.

mode

The default file mode.

ops

The &file\_operations or &block\_device\_operations structure. This must not be externally deallocated.

info

An arbitrary pointer which will be written to the <code>private\_data</code> field of the &file structure passed to the device driver. You can set this to whatever you like, and change it once the file is opened (the next file opened will not see this change).

## **Description**

Returns a handle which may later be used in a call to devfs\_unregister. On failure NULL is returned.

# devfs\_unregister

#### Name

devfs\_unregister — Unregister a device entry.

### **Synopsis**

```
void devfs_unregister (devfs_handle_t de);
```

### **Arguments**

de

A handle previously created by devfs\_register or returned from devfs\_find\_handle. If this is NULL the routine does nothing.

# devfs\_mk\_symlink

### Name

```
devfs_mk_symlink —
```

## **Synopsis**

```
int devfs_mk_symlink (devfs_handle_t dir, const char * name,
unsigned int flags, const char * link, devfs_handle_t * handle,
void * info);
```

### **Arguments**

dir

The handle to the parent devfs directory entry. If this is NULL the new name is relative to the root of the devfs.

name

The name of the entry.

flags

A set of bitwise-ORed flags (DEVFS\_FL\_\*).

link

The destination name.

handle

The handle to the symlink entry is written here. This may be NULL.

info

An arbitrary pointer which will be associated with the entry.

### **Description**

Returns 0 on success, else a negative error code is returned.

# devfs\_mk\_dir

### **Name**

devfs\_mk\_dir — Create a directory in the devfs namespace.

## **Synopsis**

```
devfs_handle_t devfs_mk_dir (devfs_handle_t dir, const char *
name, void * info);
```

### **Arguments**

dir

The handle to the parent devfs directory entry. If this is NULL the new name is relative to the root of the devfs.

name

The name of the entry.

info

An arbitrary pointer which will be associated with the entry.

### **Description**

Use of this function is optional. The devfs\_register function will automatically create intermediate directories as needed. This function is provided for efficiency reasons, as it provides a handle to a directory. Returns a handle which may later be used in a call to devfs\_unregister. On failure NULL is returned.

## devfs\_find\_handle

#### **Name**

devfs\_find\_handle — Find the handle of a devfs entry.

### **Synopsis**

devfs\_handle\_t devfs\_find\_handle (devfs\_handle\_t dir, const char
\* name, unsigned int major, unsigned int minor, char type, int
traverse\_symlinks);

### **Arguments**

dir

The handle to the parent devfs directory entry. If this is NULL the name is relative

to the root of the devfs.

name

The name of the entry.

major

The major number. This is used if name is NULL.

minor

The minor number. This is used if name is NULL.

type

The type of special file to search for. This may be either DEVFS\_SPECIAL\_CHR or DEVFS SPECIAL BLK.

traverse\_symlinks

If TRUE then symlink entries in the devfs namespace are traversed. Symlinks pointing out of the devfs namespace will cause a failure. Symlink traversal consumes stack space.

## **Description**

Returns a handle which may later be used in a call to devfs\_unregister, devfs\_get\_flags, or devfs\_set\_flags. On failure NULL is returned.

# devfs\_get\_flags

#### **Name**

devfs\_get\_flags — Get the flags for a devfs entry.

## **Synopsis**

```
int devfs_get_flags (devfs_handle_t de, unsigned int * flags);
```

### **Arguments**

de

The handle to the device entry.

flags

The flags are written here.

## **Description**

Returns 0 on success, else a negative error code.

# devfs\_get\_maj\_min

### **Name**

devfs\_get\_maj\_min — Get the major and minor numbers for a devfs entry.

## **Synopsis**

```
int devfs_get_maj_min (devfs_handle_t de, unsigned int * major,
unsigned int * minor);
```

### **Arguments**

de

The handle to the device entry.

major

The major number is written here. This may be NULL.

minor

The minor number is written here. This may be NULL.

## **Description**

Returns 0 on success, else a negative error code.

# devfs\_get\_handle\_from\_inode

### Name

devfs\_get\_handle\_from\_inode — Get the devfs handle for a VFS inode.

## **Synopsis**

```
devfs_handle_t devfs_get_handle_from_inode (struct inode *
inode);
```

### **Arguments**

inode

The VFS inode.

## **Description**

Returns the devfs handle on success, else NULL.

# devfs\_generate\_path

### **Name**

devfs\_generate\_path — Generate a pathname for an entry, relative to the devfs root.

# **Synopsis**

```
int devfs_generate_path (devfs_handle_t de, char * path, int
buflen);
```

## **Arguments**

de

The devfs entry.

path

The buffer to write the pathname to. The pathname and '\0' terminator will be written at the end of the buffer.

buflen

The length of the buffer.

### **Description**

Returns the offset in the buffer where the pathname starts on success, else a negative error code.

# devfs\_get\_ops

#### Name

devfs\_get\_ops — Get the device operations for a devfs entry.

## **Synopsis**

```
void * devfs_get_ops (devfs_handle_t de);
```

### **Arguments**

de

The handle to the device entry.

Returns a pointer to the device operations on success, else NULL.

# devfs\_set\_file\_size

#### Name

devfs\_set\_file\_size — Set the file size for a devfs regular file.

## **Synopsis**

```
int devfs_set_file_size (devfs_handle_t de, unsigned long size);
```

### **Arguments**

de

The handle to the device entry.

size

The new file size.

## **Description**

Returns 0 on success, else a negative error code.

# devfs\_get\_info

#### Name

devfs\_get\_info — Get the info pointer written to private\_data of de upon open.

## **Synopsis**

```
void * devfs_get_info (devfs_handle_t de);
```

## **Arguments**

de

The handle to the device entry.

## **Description**

Returns the info pointer.

## devfs\_set\_info

#### Name

devfs\_set\_info — Set the info pointer written to private\_data upon open.

## **Synopsis**

```
int devfs_set_info (devfs_handle_t de, void * info);
```

### **Arguments**

```
The handle to the device entry.

info

pointer to the data
```

### **Description**

Returns 0 on success, else a negative error code.

# devfs\_get\_parent

#### **Name**

devfs\_get\_parent — Get the parent device entry.

## **Synopsis**

```
devfs_handle_t devfs_get_parent (devfs_handle_t de);
```

### **Arguments**

de

The handle to the device entry.

### **Description**

Returns the parent device entry if it exists, else NULL.

# devfs\_get\_first\_child

#### Name

devfs\_get\_first\_child — Get the first leaf node in a directory.

### **Synopsis**

```
devfs_handle_t devfs_get_first_child (devfs_handle_t de);
```

### **Arguments**

de

The handle to the device entry.

### **Description**

Returns the leaf node device entry if it exists, else NULL.

# devfs\_get\_next\_sibling

#### **Name**

devfs\_get\_next\_sibling — Get the next sibling leaf node. for a device entry.

## **Synopsis**

```
devfs_handle_t devfs_get_next_sibling (devfs_handle_t de);
```

### **Arguments**

de

The handle to the device entry.

### **Description**

Returns the leaf node device entry if it exists, else NULL.

# devfs\_auto\_unregister

#### **Name**

devfs\_auto\_unregister — Configure a devfs entry to be automatically

unregistered.

### **Synopsis**

```
void devfs_auto_unregister (devfs_handle_t master,
devfs_handle_t slave);
```

#### **Arguments**

master

The master devfs entry. Only one slave may be registered.

slave

The devfs entry which will be automatically unregistered when the master entry is unregistered. It is illegal to call devfs\_unregister on this entry.

# devfs\_get\_unregister\_slave

#### **Name**

devfs\_get\_unregister\_slave — Get the slave entry which will be automatically unregistered.

### **Synopsis**

```
devfs_handle_t devfs_get_unregister_slave (devfs_handle_t
master);
```

#### **Arguments**

master

The master devfs entry.

### **Description**

Returns the slave which will be unregistered when master is unregistered.

# devfs\_get\_name

#### Name

devfs\_get\_name — Get the name for a device entry in its parent directory.

### **Synopsis**

```
const char * devfs_get_name (devfs_handle_t de, unsigned int *
namelen);
```

#### **Arguments**

de

The handle to the device entry.

namelen

The length of the name is written here. This may be NULL.

Returns the name on success, else NULL.

# devfs\_register\_chrdev

#### **Name**

devfs\_register\_chrdev — Optionally register a conventional character driver.

## **Synopsis**

```
int devfs_register_chrdev (unsigned int major, const char *
name, struct file_operations * fops);
```

## **Arguments**

```
major
```

The major number for the driver.

name

The name of the driver (as seen in /proc/devices).

fops

The &file\_operations structure pointer.

This function will register a character driver provided the "devfs=only" option was not provided at boot time. Returns 0 on success, else a negative error code on failure.

## devfs\_register\_blkdev

#### **Name**

devfs\_register\_blkdev — Optionally register a conventional block driver.

### **Synopsis**

```
int devfs_register_blkdev (unsigned int major, const char *
name, struct block_device_operations * bdops);
```

#### **Arguments**

major

The major number for the driver.

name

The name of the driver (as seen in /proc/devices).

bdops

The &block\_device\_operations structure pointer.

This function will register a block driver provided the "devfs=only" option was not provided at boot time. Returns 0 on success, else a negative error code on failure.

## devfs\_unregister\_chrdev

#### **Name**

devfs\_unregister\_chrdev — Optionally unregister a conventional character driver.

### **Synopsis**

int devfs\_unregister\_chrdev (unsigned int major, const char \*
name);

#### **Arguments**

major

The major number for the driver.

name

The name of the driver (as seen in /proc/devices).

## **Description**

This function will unregister a character driver provided the "devfs=only" option was not provided at boot time. Returns 0 on success, else a negative error code on failure.

# devfs\_unregister\_blkdev

#### **Name**

devfs\_unregister\_blkdev — Optionally unregister a conventional block driver.

## **Synopsis**

```
int devfs_unregister_blkdev (unsigned int major, const char *
name);
```

#### **Arguments**

major

The major number for the driver.

name

The name of the driver (as seen in /proc/devices).

### **Description**

This function will unregister a block driver provided the "devfs=only" option was not provided at boot time. Returns 0 on success, else a negative error code on failure.

# **Chapter 12. Power Management**

# pm\_register

#### Name

pm\_register — register a device with power management

## **Synopsis**

```
struct pm_dev * pm_register (pm_dev_t type, unsigned long id,
pm_callback callback);
```

### **Arguments**

```
device type

id

device ID

callback

callback function
```

### **Description**

Add a device to the list of devices that wish to be notified about power management events. A &pm\_dev structure is returned on success, on failure the return is NULL.

The callback function will be called in process context and it may sleep.

# pm\_unregister

#### **Name**

pm\_unregister — unregister a device with power management

## **Synopsis**

```
void pm_unregister (struct pm_dev * dev);
```

### **Arguments**

dev

device to unregister

## **Description**

Remove a device from the power management notification lists. The dev passed must be a handle previously returned by pm\_register.

## pm\_unregister\_all

#### Name

pm\_unregister\_all — unregister all devices with matching callback

## **Synopsis**

```
void pm_unregister_all (pm_callback callback);
```

#### **Arguments**

callback

callback function pointer

### **Description**

Unregister every device that would call the callback passed. This is primarily meant as a helper function for loadable modules. It enables a module to give up all its managed devices without keeping its own private list.

## pm\_send

#### **Name**

pm\_send — send request to a single device

#### **Synopsis**

```
int pm_send (struct pm_dev * dev, pm_request_t rqst, void *
data);
```

#### **Arguments**

```
dev

device to send to

rqst

power management request

data

data for the callback
```

#### **Description**

Issue a power management request to a given device. The PM\_SUSPEND and PM\_RESUME events are handled specially. The data field must hold the intended next state. No call is made if the state matches.

#### **BUGS**

what stops two power management requests occuring in parallel and conflicting.

#### **WARNING**

Calling pm\_send directly is not generally recommended, in paticular there is no locking against the pm\_dev going away. The caller must maintain all needed locking or have 'inside knowledge' on the safety. Also remember that this function is not locked against pm\_unregister. This means that you must handle SMP races on callback execution and unload yourself.

## pm\_send\_all

#### **Name**

pm\_send\_all — send request to all managed devices

## **Synopsis**

```
int pm_send_all (pm_request_t rqst, void * data);
```

#### **Arguments**

```
power management request

data

data for the callback
```

### **Description**

Issue a power management request to a all devices. The PM\_SUSPEND events are handled specially. Any device is permitted to fail a suspend by returning a non zero (error) value from its callback function. If any device vetoes a suspend request then all other devices that have suspended during the processing of this request are restored to their previous state.

#### **WARNING**

This function takes the pm\_devs\_lock. The lock is not dropped until the callbacks have completed. This prevents races against pm locking functions, races against module unload pm\_unregister code. It does mean however that you must not issue pm\_functions within the callback or you will deadlock and users will hate you.

Zero is returned on success. If a suspend fails then the status from the device that vetoes the suspend is returned.

#### **BUGS**

what stops two power management requests occuring in parallel and conflicting.

## pm\_find

#### **Name**

pm\_find — find a device

## **Synopsis**

```
struct pm_dev * pm_find (pm_dev_t type, struct pm_dev * from);
```

### **Arguments**

```
type of device
```

from

where to start looking

### **Description**

Scan the power management list for devices of a specific type. The return value for a matching device may be passed to further calls to this function to find further matches. A NULL indicates the end of the list.

To search from the beginning pass NULL as the from value.

The caller MUST hold the pm\_devs\_lock lock when calling this function. The instant that the lock is dropped all pointers returned may become invalid.

# **Chapter 13. Block Devices**

## blk\_cleanup\_queue

#### **Name**

blk\_cleanup\_queue — release a request\_queue\_t when it is no longer needed

## **Synopsis**

```
void blk_cleanup_queue (request_queue_t * q);
```

### **Arguments**

q

the request queue to be released

## **Description**

blk\_cleanup\_queue is the pair to blk\_init\_queue. It should be called when a request queue is being released; typically when a block device is being de-registered. Currently, its primary task it to free all the &struct request structures that were allocated to the queue.

#### **Caveat**

Hopefully the low level driver will have finished any outstanding requests first...

## blk\_queue\_headactive

#### **Name**

blk\_queue\_headactive — indicate whether head of request queue may be active

### **Synopsis**

```
void blk_queue_headactive (request_queue_t * q, int active);
```

#### **Arguments**

q

The queue which this applies to.

active

A flag indication where the head of the queue is active.

#### **Description**

The driver for a block device may choose to leave the currently active request on the request queue, removing it only when it has completed. The queue handling routines assume this by default for safety reasons and will not involve the head of the request queue in any merging or reordering of requests when the queue is unplugged (and thus may be working on this particular request).

If a driver removes requests from the queue before processing them, then it may indicate that it does so, there by allowing the head of the queue to be involved in merging and reordering. This is done be calling blk\_queue\_headactive with an active flag of 0.

If a driver processes several requests at once, it must remove them (or at least all but one of them) from the request queue.

When a queue is plugged the head will be assumed to be inactive.

## blk\_queue\_make\_request

#### Name

blk\_queue\_make\_request — define an alternate make\_request function for a device

## **Synopsis**

```
void blk_queue_make_request (request_queue_t * q,
make_request_fn * mfn);
```

#### **Arguments**

```
q the request queue for the device to be affected mfn the alternate make_request function
```

### **Description**

The normal way for &struct buffer\_heads to be passed to a device driver is for them to be collected into requests on a request queue, and then to allow the device driver to select requests off that queue when it is ready. This works well for many block devices. However some block devices (typically virtual devices such as md or lvm) do not benefit from the processing on the request queue, and are served best by having the

requests passed directly to them. This can be achieved by providing a function to blk\_queue\_make\_request.

#### Caveat

The driver that does this \*must\* be able to deal appropriately with buffers in "highmemory", either by calling bh\_kmap to get a kernel mapping, to by calling create\_bounce to create a buffer in normal memory.

# blk\_init\_queue

#### Name

blk\_init\_queue — prepare a request queue for use with a block device

### **Synopsis**

```
void blk_init_queue (request_queue_t * q, request_fn_proc *
rfn);
```

#### **Arguments**

```
The &request_queue_t to be initialised 

rfn
```

The function to be called to process requests that have been placed on the queue.

If a block device wishes to use the standard request handling procedures, which sorts requests and coalesces adjacent requests, then it must call blk\_init\_queue. The function rfn will be called when there are requests on the queue that need to be processed. If the device supports plugging, then rfn may not be called immediately when requests are available on the queue, but may be called at some time later instead. Plugged queues are generally unplugged when a buffer belonging to one of the requests on the queue is needed, or due to memory pressure.

rfn is not required, or even expected, to remove all requests off the queue, but only as many as it can handle at a time. If it does leave requests on the queue, it is responsible for arranging that the requests get dealt with eventually.

A global spin lock \$io\_request\_lock must be held while manipulating the requests on the request queue.

The request on the head of the queue is by default assumed to be potentially active, and it is not considered for re-ordering or merging whenever the given queue is unplugged. This behaviour can be changed with blk\_queue\_headactive.

#### **Note**

blk\_init\_queue must be paired with a blk\_cleanup\_queue call when the block device is deactivated (such as at module unload).

## generic\_make\_request

#### Name

```
generic_make_request —
```

#### **Synopsis**

```
void generic_make_request (int rw, struct buffer_head * bh);
```

#### **Arguments**

rw

READ, WRITE, or READA - what sort of I/O is desired.

bh

The buffer head describing the location in memory and on the device.

#### **Description**

generic\_make\_request is used to make I/O requests of block devices. It is passed a &struct buffer\_head and a &rw value. The READ and WRITE options are (hopefully) obvious in meaning. The READA value means that a read is required, but that the driver is free to fail the request if, for example, it cannot get needed resources immediately.

generic\_make\_request does not return any status. The success/failure status of the request, along with notification of completion, is delivered asynchronously through the bh->b\_end\_io function described (one day) else where.

The caller of generic\_make\_request must make sure that b\_page, b\_addr, b\_size are set to describe the memory buffer, that b\_rdev and b\_rsector are set to describe the device address, and the b\_end\_io and optionally b\_private are set to describe how completion notification should be signaled. BH\_Mapped should also be set (to confirm that b\_dev and b\_blocknr are valid).

generic\_make\_request and the drivers it calls may use b\_reqnext, and may change b\_rdev and b\_rsector. So the values of these fields should NOT be depended on after the call to generic\_make\_request. Because of this, the caller should record the device address information in b\_dev and b\_blocknr.

Apart from those fields mentioned above, no other fields, and in particular, no other flags, are changed by generic\_make\_request or any lower level drivers.

## submit\_bh

#### **Name**

```
submit_bh —
```

## **Synopsis**

```
void submit_bh (int rw, struct buffer_head * bh);
```

#### **Arguments**

```
{\it rw} whether to READ or WRITE, or maybe to READA (read ahead) {\it bh}
```

The &struct buffer\_head which describes the I/O

### **Description**

submit\_bh is very similar in purpose to generic\_make\_request, and uses that function to do most of the work.

The extra functionality provided by submit\_bh is to determine b\_rsector from b\_blocknr and b\_size, and to set b\_rdev from b\_dev. This is appropriate for IO requests that come from the buffer cache and page cache which (currently) always use aligned blocks.

### II\_rw\_block

#### **Name**

11\_rw\_block — level access to block devices

### **Synopsis**

```
void ll_rw_block (int rw, int nr, struct buffer_head * * bhs);
```

#### **Arguments**

```
whether to READ or WRITE or maybe READA (readahead)

nr

number of &struct buffer_heads in the array

bhs

array of pointers to &struct buffer_head
```

#### **Description**

ll\_rw\_block takes an array of pointers to &struct buffer\_heads, and requests an I/O operation on them, either a READ or a WRITE. The third READA option is described in the documentation for generic\_make\_request which ll\_rw\_block calls.

This function provides extra functionality that is not in <code>generic\_make\_request</code> that is relevant to buffers in the buffer cache or page cache. In particular it drops any buffer that it cannot get a lock on (with the BH\_Lock state bit), any buffer that appears to be clean when doing a write request, and any buffer that appears to be up-to-date when doing read request. Further it marks as clean buffers that are processed for writing (the buffer cache wont assume that they are actually clean until the buffer gets unlocked).

ll\_rw\_block sets b\_end\_io to simple completion handler that marks the buffer up-to-date (if approriate), unlocks the buffer and wakes any waiters. As client that needs a more interesting completion routine should call submit\_bh (or generic\_make\_request) directly.

#### **Caveat**

All of the buffers must be for the same device, and must also be

# end\_that\_request\_first

#### **Name**

end\_that\_request\_first — end I/O on one buffer.

### **Synopsis**

```
int end_that_request_first (struct request * req, int uptodate,
char * name);
```

#### **Arguments**

```
the request being processed

uptodate

0 for I/O error
```

name

the name printed for an I/O error

## **Description**

Ends I/O on the first buffer attached to req, and sets it up for the next buffer\_head (if any) in the cluster.

#### Return

0 - we are done with this request, call  ${\tt end\_that\_request\_last}$  1 - still buffers pending for this request

#### **Caveat**

Drivers implementing their own end\_request handling must call blk\_finished\_io appropriately.

# **Chapter 14. Miscellaneous Devices**

## misc\_register

#### **Name**

misc\_register — register a miscellaneous device

## **Synopsis**

```
int misc_register (struct miscdevice * misc);
```

### **Arguments**

misc

device structure

## **Description**

Register a miscellaneous device with the kernel. If the minor number is set to MISC\_DYNAMIC\_MINOR a minor number is assigned and placed in the minor field of the structure. For other cases the minor number requested is used.

The structure passed is linked into the kernel and may not be destroyed until it has been unregistered.

A zero is returned on success and a negative errno code for failure.

# misc\_deregister

#### Name

misc\_deregister — unregister a miscellaneous device

### **Synopsis**

```
int misc_deregister (struct miscdevice * misc);
```

## **Arguments**

misc

device to unregister

### **Description**

Unregister a miscellaneous device that was previously successfully registered with misc\_register. Success is indicated by a zero return, a negative errno code indicates an error.

# Chapter 15. Video4Linux

# video\_unregister\_device

#### **Name**

video\_unregister\_device — unregister a video4linux device

## **Synopsis**

```
void video_unregister_device (struct video_device * vfd);
```

## **Arguments**

vfd

the device to unregister

## **Description**

This unregisters the passed device and deassigns the minor number. Future open calls will be met with errors.

# **Chapter 16. Sound Devices**

# register\_sound\_special

#### Name

register\_sound\_special — register a special sound node

## **Synopsis**

int register\_sound\_special (struct file\_operations \* fops, int
unit);

### **Arguments**

fops

File operations for the driver

unit

Unit number to allocate

### **Description**

Allocate a special sound device by minor number from the sound subsystem. The allocated number is returned on succes. On failure a negative error code is returned.

# register\_sound\_mixer

#### Name

register\_sound\_mixer — register a mixer device

### **Synopsis**

```
int register_sound_mixer (struct file_operations * fops, int
dev);
```

### **Arguments**

fops

File operations for the driver

dev

Unit number to allocate

### **Description**

Allocate a mixer device. Unit is the number of the mixer requested. Pass -1 to request the next free mixer unit. On success the allocated number is returned, on failure a negative error code is returned.

# register\_sound\_midi

#### Name

register\_sound\_midi — register a midi device

### **Synopsis**

```
int register_sound_midi (struct file_operations * fops, int
dev);
```

### **Arguments**

fops

File operations for the driver

dev

Unit number to allocate

### **Description**

Allocate a midi device. Unit is the number of the midi device requested. Pass -1 to request the next free midi unit. On success the allocated number is returned, on failure a negative error code is returned.

## register\_sound\_dsp

#### **Name**

register\_sound\_dsp — register a DSP device

### **Synopsis**

```
int register_sound_dsp (struct file_operations * fops, int dev);
```

### **Arguments**

fops

File operations for the driver

dev

Unit number to allocate

### **Description**

Allocate a DSP device. Unit is the number of the DSP requested. Pass -1 to request the next free DSP unit. On success the allocated number is returned, on failure a negative error code is returned.

This function allocates both the audio and dsp device entries together and will always allocate them as a matching pair - eg dsp3/audio3

# register\_sound\_synth

#### Name

register\_sound\_synth — register a synth device

### **Synopsis**

```
int register_sound_synth (struct file_operations * fops, int
dev);
```

### **Arguments**

fops

File operations for the driver

dev

Unit number to allocate

### **Description**

Allocate a synth device. Unit is the number of the synth device requested. Pass -1 to request the next free synth unit. On success the allocated number is returned, on failure a negative error code is returned.

## unregister\_sound\_special

#### Name

unregister\_sound\_special — unregister a special sound device

### **Synopsis**

```
void unregister_sound_special (int unit);
```

### **Arguments**

unit

unit number to allocate

## **Description**

Release a sound device that was allocated with register\_sound\_special. The unit passed is the return value from the register function.

## unregister\_sound\_mixer

#### Name

unregister\_sound\_mixer — unregister a mixer

### **Synopsis**

```
void unregister_sound_mixer (int unit);
```

### **Arguments**

unit

unit number to allocate

### **Description**

Release a sound device that was allocated with register\_sound\_mixer. The unit passed is the return value from the register function.

# unregister\_sound\_midi

#### Name

unregister\_sound\_midi — unregister a midi device

```
void unregister_sound_midi (int unit);
```

unit

unit number to allocate

## **Description**

Release a sound device that was allocated with register\_sound\_midi. The unit passed is the return value from the register function.

## unregister\_sound\_dsp

#### Name

unregister\_sound\_dsp — unregister a DSP device

### **Synopsis**

```
void unregister_sound_dsp (int unit);
```

### **Arguments**

unit

unit number to allocate

### **Description**

Release a sound device that was allocated with register\_sound\_dsp. The unit passed is the return value from the register function.

Both of the allocated units are released together automatically.

## unregister\_sound\_synth

#### **Name**

unregister\_sound\_synth — unregister a synth device

### **Synopsis**

```
void unregister_sound_synth (int unit);
```

### **Arguments**

unit

unit number to allocate

### **Description**

Release a sound device that was allocated with register\_sound\_synth. The unit passed is the return value from the register function.

# **Chapter 17. USB Devices**

## usb\_register

#### **Name**

usb\_register — register a USB driver

## **Synopsis**

```
int usb_register (struct usb_driver * new_driver);
```

### **Arguments**

new\_driver

USB operations for the driver

## **Description**

Registers a USB driver with the USB core. The list of unattached interfaces will be rescanned whenever a new driver is added, allowing the new driver to attach to any recognized devices. Returns a negative error code on failure and 0 on success.

### usb\_scan\_devices

#### **Name**

usb\_scan\_devices — scans all unclaimed USB interfaces

### **Synopsis**

```
void usb_scan_devices ( void);
```

### **Arguments**

void

no arguments

## **Description**

Goes through all unclaimed USB interfaces, and offers them to all registered USB drivers through the 'probe' function. This will automatically be called after usb\_register is called. It is called by some of the USB subsystems after one of their subdrivers are registered.

## usb\_deregister

#### **Name**

usb\_deregister — unregister a USB driver

### **Synopsis**

```
void usb_deregister (struct usb_driver * driver);
```

### **Arguments**

driver

USB operations of the driver to unregister

### **Description**

Unlinks the specified driver from the internal USB driver list.

## usb\_alloc\_bus

#### Name

usb\_alloc\_bus — creates a new USB host controller structure

```
struct usb_bus * usb_alloc_bus (struct usb_operations * op);
```

op

pointer to a struct usb\_operations that this bus structure should use

### **Description**

Creates a USB host controller bus structure with the specified usb\_operations and initializes all the necessary internal objects. (For use only by USB Host Controller Drivers.)

If no memory is available, NULL is returned.

The caller should call usb\_free\_bus when it is finished with the structure.

## usb\_free\_bus

#### **Name**

usb\_free\_bus — frees the memory used by a bus structure

### **Synopsis**

```
void usb_free_bus (struct usb_bus * bus);
```

#### **Arguments**

bus

pointer to the bus to free

### **Description**

(For use only by USB Host Controller Drivers.)

## usb\_register\_bus

#### Name

usb\_register\_bus — registers the USB host controller with the usb core

## **Synopsis**

```
void usb_register_bus (struct usb_bus * bus);
```

### **Arguments**

bus

pointer to the bus to register

### **Description**

(For use only by USB Host Controller Drivers.)

## usb\_deregister\_bus

#### Name

usb\_deregister\_bus — deregisters the USB host controller

### **Synopsis**

```
void usb_deregister_bus (struct usb_bus * bus);
```

### **Arguments**

bus

pointer to the bus to deregister

### **Description**

(For use only by USB Host Controller Drivers.)

## usb\_match\_id

#### **Name**

usb\_match\_id — find first usb\_device\_id matching device or interface

### **Synopsis**

```
const struct usb_device_id * usb_match_id (struct usb_device *
dev, struct usb_interface * interface, const struct
usb_device_id * id);
```

### **Arguments**

dev

the device whose descriptors are considered when matching

interface

the interface of interest

id

array of usb\_device\_id structures, terminated by zero entry

### **Description**

usb\_match\_id searches an array of usb\_device\_id's and returns the first one matching the device or interface, or null. This is used when binding (or rebinding) a driver to an interface. Most USB device drivers will use this indirectly, through the usb core, but some layered driver frameworks use it directly. These device tables are exported with MODULE\_DEVICE\_TABLE, through modutils and "modules.usbmap", to support the driver loading functionality of USB hotplugging.

#### **What Matches**

The "match\_flags" element in a usb\_device\_id controls which members are used. If the corresponding bit is set, the value in the device\_id must match its corresponding member in the device or interface descriptor, or else the device\_id does not match.

"driver\_info" is normally used only by device drivers, but you can create a wildcard "matches anything" usb\_device\_id as a driver's "modules.usbmap" entry if you provide

an id with only a nonzero "driver\_info" field. If you do this, the USB device driver's probe routine should use additional intelligence to decide whether to bind to the specified interface.

#### What Makes Good usb\_device\_id Tables

The match algorithm is very simple, so that intelligence in driver selection must come from smart driver id records. Unless you have good reasons to use another selection policy, provide match elements only in related groups, and order match specifiers from specific to general. Use the macros provided for that purpose if you can.

The most specific match specifiers use device descriptor data. These are commonly used with product-specific matches; the USB\_DEVICE macro lets you provide vendor and product IDs, and you can also match against ranges of product revisions. These are widely used for devices with application or vendor specific bDeviceClass values.

Matches based on device class/subclass/protocol specifications are slightly more general; use the USB\_DEVICE\_INFO macro, or its siblings. These are used with single-function devices where bDeviceClass doesn't specify that each interface has its own class.

Matches based on interface class/subclass/protocol are the most general; they let drivers bind to any interface on a multiple-function device. Use the USB\_INTERFACE\_INFO macro, or its siblings, to match class-per-interface style devices (as recorded in bDeviceClass).

Within those groups, remember that not all combinations are meaningful. For example, don't give a product version range without vendor and product IDs; or specify a protocol without its associated class and subclass.

### usb\_alloc\_urb

#### **Name**

usb\_alloc\_urb — creates a new urb for a USB driver to use

### **Synopsis**

```
urb_t * usb_alloc_urb (int iso_packets);
```

### **Arguments**

```
iso_packets
number of iso packets for this urb
```

### **Description**

Creates an urb for the USB driver to use and returns a pointer to it. If no memory is available, NULL is returned.

If the driver want to use this urb for interrupt, control, or bulk endpoints, pass '0' as the number of iso packets.

The driver should call usb\_free\_urb when it is finished with the urb.

## usb\_free\_urb

#### Name

usb\_free\_urb — frees the memory used by a urb

```
void usb_free_urb (urb_t* urb);
```

urb

pointer to the urb to free

### **Description**

If an urb is created with a call to usb\_create\_urb it should be cleaned up with a call to usb\_free\_urb when the driver is finished with it.

## usb\_control\_msg

#### **Name**

usb\_control\_msg — Builds a control urb, sends it off and waits for completion

## **Synopsis**

```
int usb_control_msg (struct usb_device * dev, unsigned int pipe,
   __u8 request, __u8 requesttype, __u16 value, __u16 index, void *
data, __u16 size, int timeout);
```

### **Arguments**

dev

pointer to the usb device to send the message to

```
pipe
     endpoint "pipe" to send the message to
request
     USB message request value
requesttype
     USB message request type value
value
     USB message value
index
     USB message index value
data
     pointer to the data to send
size
     length in bytes of the data to send
timeout
     time to wait for the message to complete before timing out (if 0 the wait is
    forever)
```

### **Description**

This function sends a simple control message to a specified endpoint and waits for the message to complete, or timeout.

If successful, it returns 0, othwise a negative error number.

Don't use this function from within an interrupt context, like a bottom half handler. If you need a asyncronous message, or need to send a message from within interrupt context, use usb\_submit\_urb

## usb\_bulk\_msg

#### Name

usb\_bulk\_msg — Builds a bulk urb, sends it off and waits for completion

### **Synopsis**

```
int usb_bulk_msg (struct usb_device * usb_dev, unsigned int
pipe, void * data, int len, int * actual_length, int timeout);
```

### **Arguments**

```
usb_dev
    pointer to the usb device to send the message to

pipe
    endpoint "pipe" to send the message to

data
    pointer to the data to send

len
    length in bytes of the data to send

actual_length
    pointer to a location to put the actual length transferred in bytes

timeout
    time to wait for the message to complete before timing out (if 0 the wait is forever)
```

### **Description**

This function sends a simple bulk message to a specified endpoint and waits for the message to complete, or timeout.

If successful, it returns 0, othwise a negative error number. The number of actual bytes transferred will be plaed in the actual\_timeout paramater.

Don't use this function from within an interrupt context, like a bottom half handler. If you need a asyncronous message, or need to send a message from within interrupt context, use usb\_submit\_urb

# Chapter 18. 16x50 UART Driver

# register\_serial

#### **Name**

register\_serial — configure a 16x50 serial port at runtime

### **Synopsis**

```
int register_serial (struct serial_struct * req);
```

### **Arguments**

request structure

### **Description**

Configure the serial port specified by the request. If the port exists and is in use an error is returned. If the port is not currently in the table it is added.

The port is then probed and if neccessary the IRQ is autodetected If this fails an error is returned.

On success the port is ready to use and the line number is returned.

## unregister\_serial

#### Name

unregister\_serial — deconfigure a 16x50 serial port

## **Synopsis**

```
void unregister_serial (int line);
```

### **Arguments**

line

line to deconfigure

## **Description**

The port specified is deconfigured and its resources are freed. Any user of the port is disconnected as if carrier was dropped. Line is the port number returned by register\_serial.

# Chapter 19. Z85230 Support Library

# z8530\_interrupt

#### **Name**

z8530\_interrupt — Handle an interrupt from a Z8530

### **Synopsis**

```
void z8530_interrupt (int irq, void * dev_id, struct pt_regs *
regs);
```

### **Arguments**

```
Interrupt number

dev_id

The Z8530 device that is interrupting.

regs

unused
```

### **Description**

A Z85[2]30 device has stuck its hand in the air for attention. We scan both the channels on the chip for events and then call the channel specific call backs for each channel that has events. We have to use callback functions because the two channels can be in different modes.

## z8530\_sync\_open

#### **Name**

```
z8530_sync_open — Open a Z8530 channel for PIO
```

## **Synopsis**

```
int z8530\_sync\_open (struct net_device * dev, struct z8530\_channel * c);
```

### **Arguments**

dev

The network interface we are using

C

The Z8530 channel to open in synchronous PIO mode

## **Description**

Switch a Z8530 into synchronous mode without DMA assist. We raise the RTS/DTR and commence network operation.

## z8530\_sync\_close

#### Name

```
z8530_sync_close — Close a PIO Z8530 channel
```

### **Synopsis**

```
int z8530_sync_close (struct net_device * dev, struct
z8530_channel * c);
```

### **Arguments**

```
dev

Network device to close

c
```

Z8530 channel to disassociate and move to idle

### **Description**

Close down a Z8530 interface and switch its interrupt handlers to discard future events.

## z8530\_sync\_dma\_open

#### Name

```
z8530_sync_dma_open — Open a Z8530 for DMA I/O
```

### **Synopsis**

```
int z8530_sync_dma_open (struct net_device * dev, struct
z8530_channel * c);
```

### **Arguments**

dev

The network device to attach

C

The Z8530 channel to configure in sync DMA mode.

### **Description**

Set up a Z85x30 device for synchronous DMA in both directions. Two ISA DMA channels must be available for this to work. We assume ISA DMA driven I/O and PC limits on access.

# z8530\_sync\_dma\_close

#### **Name**

```
z8530_sync_dma_close — Close down DMA I/O
```

```
int z8530\_sync\_dma\_close (struct net_device * dev, struct z8530\_channel * c);
```

dev

Network device to detach

C

Z8530 channel to move into discard mode

### **Description**

Shut down a DMA mode synchronous interface. Halt the DMA, and free the buffers.

## z8530\_sync\_txdma\_open

#### Name

z8530\_sync\_txdma\_open — Open a Z8530 for TX driven DMA

```
int z8530_sync_txdma_open (struct net_device * dev, struct
z8530_channel * c);
```

dev

The network device to attach

C

The Z8530 channel to configure in sync DMA mode.

### **Description**

Set up a Z85x30 device for synchronous DMA tranmission. One ISA DMA channel must be available for this to work. The receive side is run in PIO mode, but then it has the bigger FIFO.

## z8530\_sync\_txdma\_close

#### **Name**

z8530\_sync\_txdma\_close — Close down a TX driven DMA channel

```
int z8530\_sync\_txdma\_close (struct net_device * dev, struct z8530\_channel * c);
```

dev

Network device to detach

C

Z8530 channel to move into discard mode

### **Description**

Shut down a DMA/PIO split mode synchronous interface. Halt the DMA, and free the buffers.

## z8530\_describe

#### Name

z8530\_describe — Uniformly describe a Z8530 port

### **Synopsis**

```
void z8530_describe (struct z8530_dev * dev, char * mapping,
unsigned long io);
```

### **Arguments**

dev

Z8530 device to describe

```
mapping
string holding mapping type (eg "I/O" or "Mem")
io
the port value in question
```

### **Description**

Describe a Z8530 in a standard format. We must pass the I/O as the port offset isnt predictable. The main reason for this function is to try and get a common format of report.

## z8530\_init

#### Name

```
z8530_init — Initialise a Z8530 device
```

### **Synopsis**

```
int z8530_init (struct z8530_dev * dev);
```

### **Arguments**

dev

Z8530 device to initialise.

### **Description**

Configure up a Z8530/Z85C30 or Z85230 chip. We check the device is present, identify the type and then program it to hopefully keep quite and behave. This matters a lot, a Z8530 in the wrong state will sometimes get into stupid modes generating 10Khz interrupt streams and the like.

We set the interrupt handler up to discard any events, in case we get them during reset or setp.

Return 0 for success, or a negative value indicating the problem in errno form.

## z8530\_shutdown

#### **Name**

z8530\_shutdown — Shutdown a Z8530 device

## **Synopsis**

```
int z8530_shutdown (struct z8530_dev * dev);
```

#### **Arguments**

dev

The Z8530 chip to shutdown

### **Description**

We set the interrupt handlers to silence any interrupts. We then reset the chip and wait 100uS to be sure the reset completed. Just in case the caller then tries to do stuff.

## z8530\_channel\_load

#### Name

z8530\_channel\_load — Load channel data

## **Synopsis**

```
int z8530_channel_load (struct z8530_channel * c, u8 * rtable);
```

### **Arguments**

C

Z8530 channel to configure

rtable

table of register, value pairs

#### **FIXME**

ioctl to allow user uploaded tables

Load a Z8530 channel up from the system data. We use +16 to indicate the "prime" registers. The value 255 terminates the table.

# z8530\_null\_rx

#### Name

```
z8530_null_rx — Discard a packet
```

### **Synopsis**

```
void z8530_null_rx (struct z8530_channel * c, struct sk_buff *
skb);
```

### **Arguments**

```
c
The channel the packet arrived on skb
The buffer
```

### **Description**

We point the receive handler at this function when idle. Instead of syncppp processing the frames we get to throw them away.

## z8530\_queue\_xmit

#### **Name**

```
z8530_queue_xmit — Queue a packet
```

### **Synopsis**

```
int z8530_queue_xmit (struct z8530_channel * c, struct sk_buff *
skb);
```

### **Arguments**

C
The channel to use

The packet to kick down the channel

### **Description**

Queue a packet for transmission. Because we have rather hard to hit interrupt latencies for the Z85230 per packet even in DMA mode we do the flip to DMA buffer if needed here not in the IRQ.

# z8530\_get\_stats

#### Name

```
z8530_get_stats — Get network statistics
```

### **Synopsis**

```
struct net_device_stats * z8530_get_stats (struct z8530_channel
* c);
```

### **Arguments**

C

The channel to use

## **Description**

Get the statistics block. We keep the statistics in software as the chip doesn't do it for us.

# **Chapter 20. Frame Buffer Library**

The frame buffer drivers depend heavily on four data structures. These structures are declared in include/linux/fb.h. They are fb\_info, fb\_var\_screeninfo, fb\_fix\_screeninfo and fb\_monospecs. The last three can be made available to and from userland.

fb\_info defines the current state of a particular video card. Inside fb\_info, there exists a fb\_ops structure which is a collection of needed functions to make fbdev and fbcon work. fb\_info is only visible to the kernel.

fb\_var\_screeninfo is used to describe the features of a video card that are user defined. With fb\_var\_screeninfo, things such as depth and the resolution may be defined.

The next structure is fb\_fix\_screeninfo. This defines the properties of a card that are created when a mode is set and can't be changed otherwise. A good example of this is the start of the frame buffer memory. This "locks" the address of the frame buffer memory, so that it cannot be changed or moved.

The last structure is fb\_monospecs. In the old API, there was little importance for fb\_monospecs. This allowed for forbidden things such as setting a mode of 800x600 on a fix frequency monitor. With the new API, fb\_monospecs prevents such things, and if used correctly, can prevent a monitor from being cooked. fb\_monospecs will not be useful until kernels 2.5.x.

## 20.1. Frame Buffer Memory

## register\_framebuffer

#### **Name**

register\_framebuffer — registers a frame buffer device

```
int register_framebuffer (struct fb_info * fb_info);
```

fb\_info

frame buffer info structure

### **Description**

Registers a frame buffer device fb\_info.

Returns negative errno on error, or zero for success.

## unregister\_framebuffer

#### Name

unregister\_framebuffer — releases a frame buffer device

## **Synopsis**

```
int unregister_framebuffer (struct fb_info * fb_info);
```

### **Arguments**

fb\_info

frame buffer info structure

### **Description**

Unregisters a frame buffer device fb\_info.

Returns negative errno on error, or zero for success.

### 20.2. Frame Buffer Console

### fbcon\_redraw\_clear

#### **Name**

fbcon\_redraw\_clear — clear area of the screen

### **Synopsis**

```
void fbcon_redraw_clear (struct vc_data * conp, struct display *
p, int sy, int sx, int height, int width);
```

### **Arguments**

```
stucture pointing to current active virtual console

p
display structure

sy
starting Y coordinate
```

```
SX
```

starting X coordinate

height

height of area to clear

width

width of area to clear

### **Description**

Clears a specified area of the screen. All dimensions are in pixels.

## fbcon\_redraw\_bmove

#### **Name**

fbcon\_redraw\_bmove — copy area of screen to another area

## **Synopsis**

```
void fbcon\_redraw\_bmove (struct display * p, int sy, int sx, int dy, int dx, int h, int w);
```

### **Arguments**

р

display structure

```
origin Y coordinate

sx
origin X coordinate

dy
destination Y coordinate

dx
destination X coordinate

h
height of area to copy

w
width of area to copy
```

## **Description**

Copies an area of the screen to another area of the same screen. All dimensions are in pixels.

Note that this function cannot be used together with ypan or ywrap.

# 20.3. Frame Buffer Colormap

# fb\_alloc\_cmap

#### Name

fb\_alloc\_cmap — allocate a colormap

## **Synopsis**

```
int fb_alloc_cmap (struct fb_cmap * cmap, int len, int transp);
```

#### **Arguments**

```
frame buffer colormap structure

len
length of cmap

transp
boolean, 1 if there is transparency, 0 otherwise
```

#### **Description**

Allocates memory for a colormap *cmap*. *len* is the number of entries in the palette.

Returns -1 errno on error, or zero on success.

# fb\_copy\_cmap

#### **Name**

```
fb_copy_cmap — copy a colormap
```

## **Synopsis**

```
void fb_copy_cmap (struct fb_cmap * from, struct fb_cmap * to,
int fsfromto);
```

## **Arguments**

```
from
    frame buffer colormap structure

to
    frame buffer colormap structure

fsfromto
```

determine copy method

## **Description**

Copy contents of colormap from from to to.

#### 0

memcpy function

1

copy\_from\_user function to copy from userspace

2

copy\_to\_user function to copy to userspace

# fb\_get\_cmap

#### **Name**

fb\_get\_cmap — get a colormap

# **Synopsis**

```
int fb_get_cmap (struct fb_cmap * cmap, int kspc, int
(*getcolreg) (u_int, u_int *, u_int *, u_int *, u_int *, struct
fb_info *), struct fb_info * info);
```

## **Arguments**

cmap

frame buffer colormap

kspc

boolean, 0 copy local, 1 put\_user function

```
getcolreg
```

pointer to a function to get a color register

info

frame buffer info structure

## **Description**

Get a colormap cmap for a screen of device info.

Returns negative errno on error, or zero on success.

# fb\_set\_cmap

#### **Name**

fb\_set\_cmap — set the colormap

## **Synopsis**

```
int fb_set_cmap (struct fb_cmap * cmap, int kspc, int
(*setcolreg) (u_int, u_int, u_int, u_int, u_int, struct fb_info
*), struct fb_info * info);
```

## **Arguments**

cmap

frame buffer colormap structure

```
boolean, 0 copy local, 1 get_user function
setcolreg
```

info

frame buffer info structure

## **Description**

Sets the colormap cmap for a screen of device info.

Returns negative errno on error, or zero on success.

# fb\_default\_cmap

#### **Name**

```
fb_default_cmap — get default colormap
```

## **Synopsis**

```
struct fb_cmap * fb_default_cmap (int len);
```

# **Arguments**

```
len
```

size of palette for a depth

#### **Description**

Gets the default colormap for a specific screen depth. *len* is the size of the palette for a particular screen depth.

Returns pointer to a frame buffer colormap structure.

# fb\_invert\_cmaps

#### Name

fb\_invert\_cmaps — invert all defaults colormaps

## **Synopsis**

```
void fb_invert_cmaps ( void);
```

#### **Arguments**

void

no arguments

## **Description**

Invert all default colormaps.

#### 20.4. Frame Buffer Generic Functions

# fbgen\_get\_fix

#### Name

fbgen\_get\_fix — get fixed part of display

### **Synopsis**

```
int fbgen_get_fix (struct fb_fix_screeninfo * fix, int con,
struct fb_info * info);
```

#### **Arguments**

```
fix

fb_fix_screeninfo structure

con

virtual console number

info

frame buffer info structure
```

#### **Description**

Get the fixed information part of the display and place it into fix for virtual console con on device info.

# fbgen\_get\_var

#### Name

fbgen\_get\_var — get user defined part of display

## **Synopsis**

```
int fbgen_get_var (struct fb_var_screeninfo * var, int con,
struct fb_info * info);
```

## **Arguments**

```
var
fb_var_screeninfo structure

con
    virtual console number

info
    frame buffer info structure
```

## **Description**

Get the user defined part of the display and place it into *var* for virtual console *con* on device *info*.

# fbgen\_set\_var

#### Name

fbgen\_set\_var — set the user defined part of display

#### **Synopsis**

```
int fbgen_set_var (struct fb_var_screeninfo * var, int con,
struct fb_info * info);
```

## **Arguments**

```
fb_var_screeninfo user defined part of the display

con

virtual console number

info

frame buffer info structure
```

## **Description**

Set the user defined part of the display as dictated by var for virtual console con on device info.

# fbgen\_get\_cmap

#### Name

```
fbgen_get_cmap — get the colormap
```

## **Synopsis**

```
int fbgen_get_cmap (struct fb_cmap * cmap, int kspc, int con,
struct fb_info * info);
```

## **Arguments**

```
frame buffer colormap structure

kspc
boolean, 0 copy local, 1 put_user function

con
virtual console number

info
frame buffer info structure
```

## **Description**

Gets the colormap for virtual console *con* and places it into *cmap* for device *info*.

Returns negative errno on error, or zero for success.

# fbgen\_set\_cmap

#### **Name**

```
fbgen_set_cmap — set the colormap
```

# **Synopsis**

```
int fbgen_set_cmap (struct fb_cmap * cmap, int kspc, int con,
struct fb_info * info);
```

## **Arguments**

```
frame buffer colormap structure

kspc

boolean, 0 copy local, 1 get_user function

con

virtual console number

info

frame buffer info structure
```

## **Description**

Sets the colormap *cmap* for virtual console *con* on device *info*.

# fbgen\_pan\_display

#### Name

fbgen\_pan\_display — pan or wrap the display

# **Synopsis**

```
int fbgen_pan_display (struct fb_var_screeninfo * var, int con,
struct fb_info * info);
```

#### **Arguments**

```
frame buffer user defined part of display

con

virtual console number

info

frame buffer info structure
```

# **Description**

Pan or wrap virtual console con for device info.

This call looks only at xoffset, yoffset and the FB\_VMODE\_YWRAP flag in var.

# fbgen\_do\_set\_var

#### Name

fbgen\_do\_set\_var — change the video mode

#### **Synopsis**

```
int fbgen_do_set_var (struct fb_var_screeninfo * var, int
isactive, struct fb_info_gen * info);
```

#### **Arguments**

```
frame buffer user defined part of display

isactive
boolean, 0 inactive, 1 active

info
generic frame buffer info structure
```

## **Description**

Change the video mode settings for device *info*. If *isactive* is non-zero, the changes will be activated immediately.

# fbgen\_set\_disp

#### Name

fbgen\_set\_disp — set generic display

## **Synopsis**

```
void fbgen_set_disp (int con, struct fb_info_gen * info);
```

## **Arguments**

con

virtual console number

info

generic frame buffer info structure

## **Description**

Sets a display on virtual console con for device info.

# fbgen\_install\_cmap

#### **Name**

fbgen\_install\_cmap — install the current colormap

## **Synopsis**

```
void fbgen_install_cmap (int con, struct fb_info_gen * info);
```

#### **Arguments**

```
virtual console number

info

generic frame buffer info structure
```

## **Description**

Installs the current colormap for virtual console *con* on device *info*.

# fbgen\_update\_var

#### **Name**

fbgen\_update\_var — update user defined part of display

## **Synopsis**

```
int fbgen_update_var (int con, struct fb_info * info);
```

#### **Arguments**

```
virtual console number
```

frame buffer info structure

## **Description**

Updates the user defined part of the display ('var' structure) on virtual console *con* for device *info*. This function is called by fbcon.c.

Returns negative errno on error, or zero for success.

# fbgen\_switch

#### **Name**

fbgen\_switch — switch to a different virtual console.

#### **Synopsis**

```
int fbgen_switch (int con, struct fb_info * info);
```

#### **Arguments**

con

virtual console number

info

frame buffer info structure

## **Description**

Switch to virtuall console con on device info.

Returns zero.

# fbgen\_blank

#### Name

fbgen\_blank — blank the screen

## **Synopsis**

```
void fbgen_blank (int blank, struct fb_info * info);
```

## **Arguments**

blank

boolean, 0 unblank, 1 blank

info

frame buffer info structure

#### **Description**

Blank the screen on device info.

## 20.5. Frame Buffer Video Mode Database

# fb\_find\_mode

#### Name

fb\_find\_mode — finds a valid video mode

### **Synopsis**

```
int __init fb_find_mode (struct fb_var_screeninfo * var, struct
fb_info * info, const char * mode_option, const struct
fb_videomode * db, unsigned int dbsize, const struct
fb_videomode * default_mode, unsigned int default_bpp);
```

## **Arguments**

var

frame buffer user defined part of display

info

frame buffer info structure

```
mode_option
string video mode to find

db
video mode database

dbsize
size of db

default_mode
default video mode to fall back to

default_bpp
default color depth in bits per pixel
```

#### **Description**

Finds a suitable video mode, starting with the specified mode in mode\_option with fallback to default\_mode. If default\_mode fails, all modes in the video mode database will be tried.

Valid mode specifiers for mode\_option:

```
<xres>x<yres>[-<bpp>][@<refresh>] or <name>[-<bpp>][@<refresh>]
with <xres>, <yres>, <bpp> and <refresh> decimal numbers and <name> a string.
```

#### **NOTE**

The passed struct var is \_not\_ cleared! This allows you to supply values for e.g. the grayscale and accel\_flags fields.

Returns zero for failure, 1 if using specified mode\_option, 2 if using specified mode\_option with an ignored refresh rate, 3 if default mode is used, 4 if fall back to any valid mode.

# \_\_fb\_try\_mode

#### **Name**

```
__fb_try_mode — test a video mode
```

## **Synopsis**

```
int __fb_try_mode (struct fb_var_screeninfo * var, struct
fb_info * info, const struct fb_videomode * mode, unsigned int
bpp);
```

#### **Arguments**

```
frame buffer user defined part of display

info
frame buffer info structure

mode
frame buffer video mode structure

bpp
```

color depth in bits per pixel

## **Description**

Tries a video mode to test it's validity for device info.

Returns 1 on success.

# 20.6. Frame Buffer Macintosh Video Mode Database

# console\_getmode

#### Name

console\_getmode — get current mode

## **Synopsis**

```
int console_getmode (struct vc_mode * mode);
```

# **Arguments**

mode

virtual console mode structure

#### **Description**

Populates *mode* with the current mode held in the global display\_info structure.

Note, this function is only for XPMAC compatibility.

Returns zero.

# console\_setmode

#### Name

console\_setmode — sets current console mode

## **Synopsis**

```
int console_setmode (struct vc_mode * mode, int doit);
```

## **Arguments**

mode

virtual console mode structure

doit

boolean, 0 test mode, 1 test and activate mode

## **Description**

Sets mode for all virtual consoles if doit is non-zero, otherwise, test a mode for validity.

Note, this function is only for XPMAC compatibility.

## console\_setcmap

#### **Name**

console\_setcmap — sets palette color map for console

#### **Synopsis**

```
int console_setcmap (int n_entries, unsigned char * red,
unsigned char * green, unsigned char * blue);
```

#### **Arguments**

```
n_entries

number of entries in the palette (max 16)

red

value for red component of palette

green

value for green component of palette

blue

value for blue component of palette
```

## **Description**

Sets global palette\_cmap structure and activates the palette on the current console.

Note, this function is only for XPMAC compatibility.

# console\_powermode

#### **Name**

console\_powermode — sets monitor power mode

## **Synopsis**

```
int console_powermode (int mode);
```

## **Arguments**

mode

power state to set

## **Description**

Sets power state as dictated by mode.

Note that this function is only for XPMAC compatibility and doesn't do much.

Returns 0 for VC\_POWERMODE\_INQUIRY, -EINVAL for VESA power settings, or -ENIXIO on failure.

# mac\_vmode\_to\_var

#### Name

mac\_vmode\_to\_var — converts vmode/cmode pair to var structure

## **Synopsis**

```
int mac_vmode_to_var (int vmode, int cmode, struct
fb_var_screeninfo * var);
```

## **Arguments**

vmode

MacOS video mode

cmode

MacOS color mode

var

frame buffer video mode structure

## **Description**

Converts a MacOS vmode/cmode pair to a frame buffer video mode structure.

## mac\_var\_to\_vmode

#### Name

mac\_var\_to\_vmode — convert var structure to MacOS vmode/cmode pair

## **Synopsis**

```
int mac_var_to_vmode (const struct fb_var_screeninfo * var, int
* vmode, int * cmode);
```

## **Arguments**

var

frame buffer video mode structure

vmode

MacOS video mode

cmode

MacOS color mode

#### **Description**

Converts a frame buffer video mode structure to a MacOS vmode/cmode pair.

# mac\_map\_monitor\_sense

#### **Name**

mac\_map\_monitor\_sense — Convert monitor sense to vmode

#### **Synopsis**

```
int mac_map_monitor_sense (int sense);
```

## **Arguments**

sense

Macintosh monitor sense number

## **Description**

Converts a Macintosh monitor sense number to a MacOS vmode number.

Returns MacOS vmode video mode number.

# mac\_find\_mode

#### **Name**

mac\_find\_mode — find a video mode

#### **Synopsis**

```
int __init mac_find_mode (struct fb_var_screeninfo * var, struct
fb_info * info, const char * mode_option, unsigned int
default_bpp);
```

#### **Arguments**

```
frame buffer user defined part of display

info
frame buffer info structure

mode_option
video mode name (see mac_modedb[])

default_bpp
default color depth in bits per pixel
```

#### **Description**

Finds a suitable video mode. Tries to set mode specified by *mode\_option*. If the name of the wanted mode begins with 'mac', the Mac video mode database will be used, otherwise it will fall back to the standard video mode database.

#### **Note**

Function marked as \_\_init and can only be used during system boot.

Returns error code from fb\_find\_mode (see fb\_find\_mode function).

## 20.7. Frame Buffer Fonts

# fbcon\_find\_font

#### Name

```
fbcon_find_font — find a font
```

## **Synopsis**

```
struct fbcon_font_desc * fbcon_find_font (char * name);
```

## **Arguments**

name

string name of a font

## **Description**

Find a specified font with string name name.

Returns NULL if no font found, or a pointer to the specified font.

# fbcon\_get\_default\_font

#### Name

fbcon\_get\_default\_font — get default font

## **Synopsis**

```
struct fbcon_font_desc * fbcon_get_default_font (int xres, int
yres);
```

## **Arguments**

```
xres
screen size of X
yres
screen size of Y
```

## **Description**

Get the default font for a specified screen size. Dimensions are in pixels.

Returns NULL if no font is found, or a pointer to the chosen font.